

NERFIFY: A Multi-Agent Framework for Turning NeRF Papers into Code

Supplementary Material

001

Summary of Supplementary Material

- Section 1 : Extended Experimental Analysis
 - Section 1.1 : Complete NERFIFY-BENCH Results
 - Section 1.2 : Qualitative Comparisons Gallery
 - Section 1.3 : Additional Ablation Studies
- Section 2: Comparison with Multi-Agent Baselines
- Section 3: Novelty Coverage on Set 4 Papers
- Section 4: Implementation Details
 - Section 4.1 : Context-Free Grammar
 - Section 4.2 : Multi-Agent Architecture and Agent Specifications
 - Section 4.3 : LLM Prompts Used in NERFIFY
- Section 5: NERFIFY-BENCH Dataset

1. Extended Experimental Analysis

1.1. Complete NERFIFY-BENCH Results

1.1.1. Quantitative Metrics for Set 1 (Never-Implemented Papers)

Paper	Reported			Human Impl.			NERFIFY (Ours)		
	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓
KeyNeRF [21]	25.65	0.89	0.11	25.70	0.89	0.12	26.12	0.90	0.09
mi-MLP NerF [34]	24.70	0.89	0.09	22.64	0.87	0.15	22.85	0.87	0.15
ERS [24]	27.85	0.94	0.06	26.87	0.90	0.12	27.02	0.90	0.12
TVNeRF [33]	27.44	0.93	0.08	26.81	0.92	0.12	27.30	0.92	0.10
Anisotropic NeRF [27]	34.08	0.97	0.05	28.85	0.94	0.06	29.01	0.94	0.06
Nerf-ID [1]	25.15	0.94	-	23.01	0.89	0.13	23.10	0.89	0.13
Surface Samp. [32]	25.63	-	-	24.35	0.90	0.11	24.40	0.90	0.11
LiNeRF [4]	25.60	0.93	0.08	23.30	0.90	0.14	23.32	0.89	0.14
HybNeRF [28]	33.94	0.96	0.047	30.45	0.94	0.07	30.51	0.95	0.07
AR-NeRF	20.36	0.79	0.17	19.00	0.76	0.19	20.05	0.78	0.18

Table 1. Comparison of NERFIFY with paper and human implementations. We evaluate NeRF papers from the NERFIFY-BENCH set whose code is not publicly available, using SSIM, PSNR, and LPIPS metrics. Note. Other baselines like Paper2Code, AutoP2C, GPT-5 and R1 failed to generate trainable code.

1.1.2. Quantitative Metrics for Set 2 and Set 3 (Github Repository Available)

Method	Original Repository			NERFIFY		
	PSNR ↑	SSIM ↑	LPIPS ↓	PSNR ↑	SSIM ↑	LPIPS ↓
Vanilla NeRF [16]	31.36	0.95	0.04	31.36	0.95	0.04
Nerfacto [25]	20.36	0.82	0.22	20.36	0.82	0.22
SeaThru-NeRF	27.89	0.83	0.22	30.08	0.92	0.07
BioNeRF	25.66	0.93	0.05	23.75	0.90	0.11
InstantNGP	32.77	-	-	32.64	0.96	0.06
ℓ_0 Sampler [14]	29.21	-	0.04	30.13	0.97	0.03
InfoNeRF [11]	18.27	0.81	0.23	17.87	0.69	0.44
DeblurNeRF	32.08	0.93	0.5	31.10	0.86	0.06
NerfingMVS	31.43	0.96	-	31.51	0.92	0.09

Table 2. Comparison with existing implementations. Evaluation of NERFIFY against original author repositories or gold-standard implementations.

Configuration	PSNR	SSIM	LPIPS
NERFIFY (Full)	27.16	0.91	0.11
<i>Knowledge Sources:</i>			
w/o Citation Recovery and In-context Examples	23.22	0.87	0.26
<i>Validation & Feedback:</i>			
w/o VLM Feedback and Smoke Testing (Stage 4)	9.39	0.80	0.42
<i>Planning Strategy:</i>			
One-Shot (no GoT) (Stage 3)	24.52	0.90	0.16

Table 3. Component ablation study. We evaluate the impact of each system component on synthesis quality and efficiency. Numbers are averaged on trainable implementations among 10 NERFIFY-BENCH Set 1 papers.

1.2. Qualitative Comparisons Gallery

020

1.2.1. Side-by-side Visual Comparisons

021

Figure 1 and Figure 2 shows the qualitative comparison of some of the methods in Set 1 of the NERFIFY bench.

022

023

1.3. Additional Ablation Studies

024

The ablation study experiments with the image quality metrics are shown in Table 3.

025

026

2. Comparison with Multi-Agent Baselines

027

Table 4 compares NERFIFY with automated coding systems on papers from nerifybench evaluating executability and semantic score. Recent PaperBench results show DeepCode (75.9% accuracy) substantially outperforms PaperCoder (51.1%), Claude Code (58.7%) and LLM-based systems like o1 (43.3%) [9]. So, we use DeepCode (with gpt-5 api) as our baseline alongside multi-agent frameworks GPT 5 thinking [20], ChatDev [23] (with gpt-5 api) and MetaGPT [10] (with gpt-5 api).

028

029

030

031

032

033

034

035

036

037

Table 4. Comparison of code generation systems on NeRF papers (in-depth comparisons shown in supplementary).

Paper	GPT-5 Thinking		ChatDev		MetaGPT		DeepCode		NERFIFY	
	Score	Trainable	Score	Trainable	Score	Trainable	Score	Trainable	Score	Trainable
KeyNeRF [18]	0.85	-	0.25	x	0.30	x	0.60	x	1.00	✓
FastNeRF [6]	0.65	x	0.21	x	0.36	x	0.81	x	0.95	✓
Vanilla NeRF [16]	0.71	✓	0.48	x	0.29	x	0.53	x	0.92	✓
Deblur-NeRF [15]	0.82	-	0.18	x	0.42	x	0.75	x	1.00	✓
Average	0.76	-	0.28	x	0.34	x	0.67	x	0.97	✓

GPT-5 thinking deep research achieves moderate semantic scores and produces trainable code for already implemented papers like Vanilla NeRF but misses papers whose repositories are not available online. ChatDev and MetaGPT generate syntactically valid Python code but fail to capture NeRF-specific patterns. DeepCode performs better than

038

039

040

041

042

043



Figure 1. **Visual Comparison of NERIFY and Human Implementation.** **Left:** Ground Truth Image, **Middle:** Expert Implementation, **Right:** Agent Implementation.

044 general multi-agent systems, benefiting from paper-specific
045 analysis, but still cannot produce trainable code.

046 Detailed comparison for each paper across these multi-
047 agent pipeline using LLM as an evaluation are given be-
048 low. The evaluation leverages language models as semantic
049 judges, analyzing both mathematical correctness and algo-
050 rithmic understanding. This automated assessment as shown
051 in section 2 and section 3 examines whether implementa-
052 tions capture the conceptual intent beyond syntactic correctness.
053 For instance, when evaluating Mip-NeRF’s integrated po-
054 sitional encoding, the semantic evaluator verifies not just
055 the presence of the encoding formula but also its proper
056 integration within the coarse-to-fine sampling pipeline.

2.1. KeyNeRF: Informative Rays Selection for Few-Shot Neural Radiance Fields

2.1.1. Paper Overview

KeyNeRF addresses the critical challenge of few-shot novel view synthesis by introducing an intelligent camera and ray selection mechanism. The paper contributes two key innovations: first, a greedy baseline diversity algorithm for selecting the most informative camera views from limited input, and second, an entropy-based ray sampling strategy that prioritizes regions with higher visual information content. These techniques enable high-quality NeRF reconstruction from as few as 10-20 input views, where standard NeRF methods typically fail. The paper demonstrates that strategic data selection can be more effective than architectural modifications for few-shot scenarios.

057
058

059

060
061
062
063
064
065
066
067
068
069
070
071

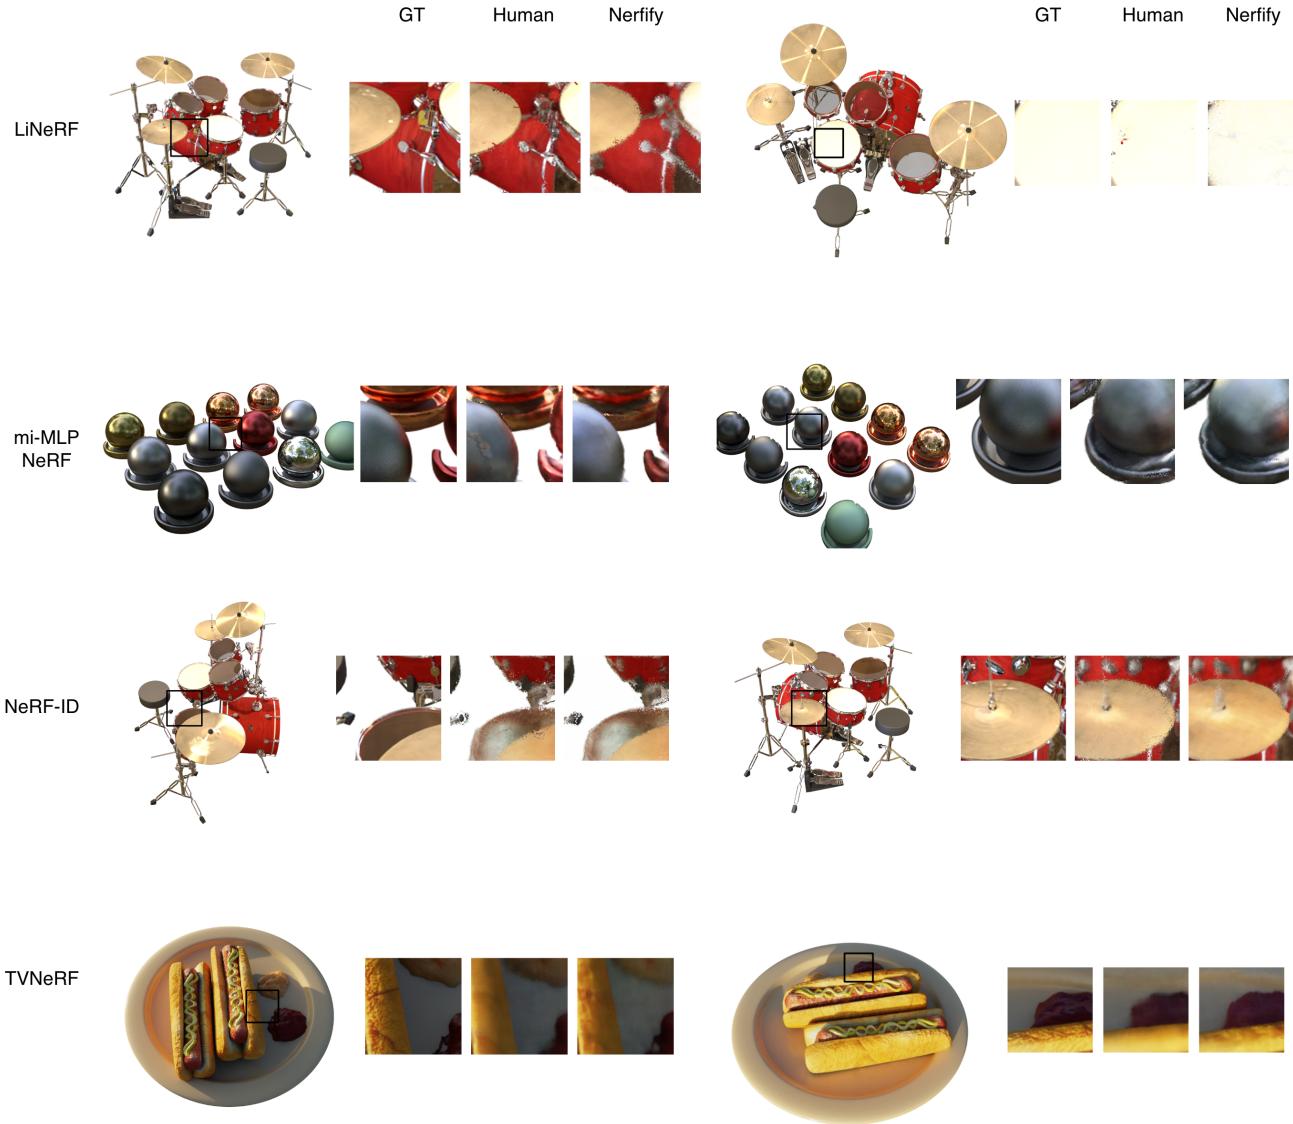


Figure 2. **Visual Comparison of NERIFY and Human Implementation.** **Left:** Ground Truth Image, **Middle:** Expert Implementation, **Right:** Agent Implementation.

072

2.1.2. Implementation Overview

073

2.1.3. Novel Components

Table 5. KeyNeRF Implementation Summary Across Multi-agent Systems

Aspect	NERIFY	GPT-5-thinking	DeepCode	MetaGPT	ChatDev
Lines of Code	479	507	1615	412	385
File Organization	Plugin	Modular	Multi-module	Multi-file	Multi-file
View Selection	✓	✓	Partial	✗	Attempted
Entropy Sampling	✓	✓	✓	✗	Simplified
Nerfstudio Integration	✓	✓	✗	✗	✗
Trainable	✓	✓	✗	✗	✗

Table 6. Novel Components in KeyNeRF with Importance Weights

ID	Component	Weight w_i
C1	Greedy baseline diversity for camera selection	0.20
C2	Integer linear programming (ILP) for coverage	0.15
C3	Local entropy computation for ray importance	0.15
C4	Entropy-weighted ray sampling	0.15
C5	Mixed sampling strategy (entropy + uniform)	0.10
C6	Nerfstudio datamanager integration	0.10
C7	Per-camera entropy caching	0.05
C8	Appearance embedding regularization	0.05
C9	Camera frustum visibility checks	0.03
C10	Grayscale conversion for entropy	0.02

074

2.1.4. Quantitative Metrics

Table 7. KeyNeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	1.00	0.00	0.00	1.00	1.00
GPT-5 thinking	0.85	0.10	0.05	0.93	0.85
DeepCode	0.60	0.15	0.25	0.70	0.60
MetaGPT	0.30	0.40	0.30	0.45	0.30
ChatDev	0.25	0.35	0.40	0.35	0.25

075

2.1.5. Component-by-Component Analysis

Component C1: Greedy Baseline Diversity The paper specifies a greedy algorithm that iteratively selects cameras maximizing the minimum baseline angle to already-selected views, ensuring diverse viewpoint coverage for robust reconstruction.

```

1 # \nerfify\ : Full greedy baseline with angle
   computation
2 def _select_cameras_greedy_baseline(self, K: int) ->
3     List[int]:
4         z_axes = F.normalize(c2w_t[:, :3, 2], dim=-1) # optical axes
5         dots = z_axes @ z_axes.t()
6         angles = torch.acos(torch.clamp(dots, -1.0, 1.0))
7         init_idx = int(torch.argmax(torch.nanmean(angles,
8             dim=1)).item())
9         selected = [init_idx]
10        while len(selected) < min(K, N) and remaining:
11            sub = angles.index_select(0, rem_idx_t).
12            index_select(1, sel_idx_t)
13            min_to_sel = torch.min(sub, dim=1).values
14            pick = rem_list[int(torch.argmax(min_to_sel).
15                item())]
16            selected.append(pick)

```

Listing 1. C1 Implementation: NERFIFY (Score: 1.0)

102

```

1 # GPT-5: Complete algorithm with proper angle
   computation
2 def greedy_baseline_ranking(coverage_set, num_cams, B
3     , K=None):
4     B = compute_baseline_matrix(cameras) # Proper
   angle matrix
5     select = coverage_set.copy()
6     while remain and len(select) < K:
7         sub_mat = B[np.ix_(remain, select)]
8         scores = np.min(sub_mat, axis=1)
9         best_cam = remain[int(np.argmax(scores))]
10        select.append(best_cam)
11    # CORRECT: Full implementation, just offline
   preprocessing

```

Listing 2. C1 Implementation: GPT-5 (Score: 0.9)

118

```

1 # ChatDev: Placeholder without actual implementation
2 def select_views(dataset):
3     num_cameras = len(dataset.poses)
4     visibility_matrix = compute_visibility_matrix(
5         dataset)
6     select_cams = solve_ilp(visibility_matrix,
   num_cameras)
7     return select_cams
8 # ERROR: is_visible() always returns True

```

Listing 3. C1 Implementation: ChatDev (Score: 0.3)

Analysis: NERFIFY implements the complete greedy baseline algorithm integrated into the datamanager. GPT-5 provides an equally correct and complete implementation with proper angle computation, using offline preprocessing which is a valid architectural choice. ChatDev and MetaGPT lack proper implementation, while DeepCode focuses on different selection criteria.

129

130

131

132

133

134

135

Component C2: Integer Linear Programming for Coverage The paper formulates camera selection as an ILP problem minimizing the number of cameras while ensuring all scene points are covered.

```

1 # \nerfify\ : Complete ILP with OR-Tools integration
2 def solve_coverage_ilp(A: np.ndarray):
3     solver = pywraplp.Solver.CreateSolver("SCIP")
4     x = [solver.BoolVar(f"x_{i}") for i in range(
5         num_cams)]
6     for j in range(num_pts):
7         ct = solver.RowConstraint(1, solver.infinity
8             (), f"pt_{j}")
9         for i in range(num_cams):
10            if A[i, j] == 1:
11                ct.SetCoefficient(x[i], 1.0)
12        solver.Minimize(solver.Sum(x[i] for i in range(
13            num_cams)))

```

Listing 4. C2 Implementation: NERFIFY (Score: 1.0)

136

137

138

139

140

141

142

143

144

145

146

147

148

149

150

151

152

153

```

1 # GPT-5: Identical ILP formulation with complete
   visibility matrix
2 def solve_coverage_ilp(A: np.ndarray):
3     solver = pywraplp.Solver.CreateSolver("SCIP")
4     x = [solver.BoolVar(f"x_{i}") for i in range(
5         num_cams)]
6     for j in range(num_pts):
7         ct = solver.RowConstraint(1, solver.infinity
8             (), f"pt_{j}")
9         for i in range(num_cams):
10            if A[i, j] == 1:
11                ct.SetCoefficient(x[i], 1.0)
12        solver.Minimize(solver.Sum(x[i] for i in range(
13            num_cams)))
14    status = solver.Solve()
15    select = [i for i in range(num_cams) if x[i].
16        solution_value() > 0.5]
17    # CORRECT: Complete implementation with proper error
   handling

```

Listing 5. C2 Implementation: GPT-5 (Score: 1.0)

155

156

157

158

159

160

161

162

163

164

165

166

167

168

169

170

171

172

173

174

175

```

1 # ChatDev: Attempts ILP but incomplete visibility
2 def solve_ilp(visibility_matrix, num_cameras):
3     solver = pywraplp.Solver.CreateSolver('SCIP')
4     x = [solver.BoolVar(f"x_{i}") for i in range(
5         num_cameras)]
6     solver.Minimize(solver.Sum(x))
7     for j in range(visibility_matrix.shape[1]):
8         solver.Add(solver.Sum(x[i] *
   visibility_matrix[i, j]
     for i in range(num_cameras)) >= 1)
9    # ERROR: visibility_matrix always zeros due to
   placeholder

```

Listing 6. C2 Implementation: ChatDev (Score: 0.5)

176

177

178

179

180

181

182

183

184

185

186

187

188

Analysis: Both NERFIFY and GPT-5 provide complete, correct ILP formulations with proper OR-Tools integration. The only difference is architectural choice: NERFIFY integrates it into the datamanager while GPT-5 uses offline

190

191

192

193

194 preprocessing. ChatDev attempts ILP but fails due to broken
 195 visibility computation. MetaGPT and DeepCode omit this
 196 component entirely.

197 **Component C3: Local Entropy Computation** The paper
 198 computes local entropy using sliding windows over grayscale
 199 images to identify information-rich regions.

```
200 1 # \nerfify\ : Proper windowed entropy with histogram
201   binning
202 2 gray = self._rgb_to_gray(image.clamp(0.0, 1.0))
203 3 q = torch.clamp((gray * (bins - 1)).long(), 0, bins -
204    1)
205 4 one_hot = F.one_hot(q.squeeze(0), num_classes=bins)
206 5 kernel = torch.ones(1, 1, win, win) / (win * win)
207 6 hist_local = F.conv2d(one_hot.float(), kernel,
208    padding=pad)
209 7 probs = hist_local / (hist_local.sum(dim=1, keepdim=
210      True) + eps)
211 8 entropy = -(probs * torch.log(probs + eps)).sum(dim =
212      1)
```

Listing 7. C3 Implementation: NERFIFY (Score: 1.0)

```
213 1 # DeepCode: Simplified entropy without windowing
214 2 def compute_entropy(image_batch):
215 3     entropy_maps = []
216 4     for img in image_batch:
217 5         gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)
218 6         hist = cv2.calcHist([gray], [0], None, [256],
219          [0,256])
220 7         hist = hist / hist.sum()
221 8         entropy = -np.sum(hist * np.log(hist + 1e-7))
222 9         entropy_maps.append(entropy)
223 10 # ISSUE: Global entropy, not local per-pixel
```

Listing 8. C3 Implementation: DeepCode (Score: 0.7)

224 **Analysis:** NERFIFY implements proper local windowed
 225 entropy computation. DeepCode computes global entropy
 226 missing spatial locality. Others either omit or provide place-
 227 holder implementations.

228 **Component C4: Entropy-Weighted Ray Sampling** The
 229 paper samples rays proportionally to their local entropy val-
 230 ues, focusing training on high-information regions.

```
231 1 # \nerfify\ : Proper entropy-weighted sampling with
232   mixing
233 2 entropy_probs = self._compute_entropy_probs(image)
234 3 mix = self.config.entropy_mix_ratio
235 4 n_entropy = int(B * mix)
236 5 n_uniform = B - n_entropy
237 6 entropy_idx = torch.multinomial(entropy_probs,
238    n_entropy)
239 7 uniform_idx = torch.randint(0, H * W, (n_uniform,))
240 8 indices = torch.cat([entropy_idx, uniform_idx])
```

Listing 9. C4 Implementation: NERFIFY (Score: 1.0)

```
241 1 # GPT-5: Proper entropy sampling with Nerfstudio
242   integration
243 2 e = self.entropy_maps[pose_idx_t].clamp(min=1e-6)
244 3 probs = e.view(-1) / e.view(-1).sum()
245 4 B_entropy = B // 2 # Mix ratio
246 5 B_uniform = B - B_entropy
247 6 flat_idx_ent = torch.multinomial(probs, B_entropy,
248    replacement=True)
249 7 flat_idx_uni = torch.randint(0, H*W, (B_uniform,),
```

8 ray_bundle = self.ray_generator(ray_indices) #	258
Nerfstudio integration	259
9 # CORRECT: Full implementation with proper framework	260
usage	261

Listing 10. C4 Implementation: GPT-5 (Score: 0.9)

1 # ChatDev: Incorrect entropy usage	263
2 entropy_probs = compute_entropy(batch['image'])	264
3 rays_ids = np.random.choice(len(rays), size=	265
batch_size, p=entropy_probs)	266
4 # CRITICAL ERROR: entropy_probs wrong shape and not	267
normalized	268

Listing 11. C4 Implementation: ChatDev (Score: 0.2)

271 **Analysis:** NERFIFY and GPT-5 thinking both implement
 272 proper entropy-weighted sampling with correct probability
 273 computation and mixing strategy. GPT-5's implementation
 274 demonstrates sophisticated understanding of Nerfstudio's
 275 RayGenerator integration. Other baselines fail to properly
 276 implement this crucial component.

2.1.6. Scoring Analysis

Table 8. KeyNeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	Weighted Avg _{LIM}
NERFIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.00
GPT-5 thinking	0.9	1.0	0.8	0.9	0.9	0.8	0.8	0.9	0.8	0.8	0.85
DeepCode	0.6	0.0	0.7	0.6	0.5	0.0	0.8	0.7	0.9	0.8	0.60
MetaGPT	0.0	0.0	0.0	0.3	0.0	0.0	0.8	0.7	0.5	0.30	
ChatDev	0.3	0.5	0.2	0.2	0.1	0.0	0.0	0.4	0.3	0.25	

2.1.7. Why Baselines Fail Despite Component Scores

278 **GPT-5 thinking (Score: 8.50 components, 85% trainable)**

- 279 **Strengths:** Comprehensive understanding of both KeyNeRF algorithms and Nerfstudio integration, provides working training pipeline with RayGenerator
- 280 **Minor Issues:**
 - 281 – Uses offline preprocessing rather than integrated data-manager approach
 - 282 – Custom trainer instead of full plugin structure (but provides migration path)
- 283 **Result:** Near-production ready implementation requiring minimal adaptation

284 **DeepCode (Score: 6.00 components, 0% trainable)**

- 285 **Strengths:** Comprehensive standalone implementation with evaluation pipeline
- 286 **Fatal Issues:**
 - 287 – No Nerfstudio integration whatsoever
 - 288 – Custom NeRF implementation incompatible with modern frameworks
 - 289 – Missing critical view selection algorithms
- 290 **Result:** Standalone code that cannot leverage Nerfstudio infrastructure

300 **MetaGPT (Score: 3.00 components, 0% trainable)**
 301 • **Strengths:** Clean code structure with proper configuration
 302 management
 303 • **Fatal Issues:**
 304 – No KeyNeRF-specific components implemented
 305 – Basic NeRF without view or ray selection
 306 – Missing positional encoding and volume rendering
 307 • **Result:** Generic NeRF attempt missing all KeyNeRF in-
 308 novations

309 **ChatDev (Score: 2.50 components, 0% trainable)**
 310 • **Strengths:** Attempts ILP formulation and mentions en-
 311 tropy
 312 • **Fatal Issues:**
 313 – Placeholder functions return dummy values
 314 – GUI components inappropriate for NeRF training
 315 – Fundamental misunderstanding of ray sampling
 316 • **Result:** Non-functional code with critical algorithmic er-
 317 rors

318 2.1.8. Hyperparameter Fidelity

Table 9. Hyperparameter Implementation Accuracy

Parameter	Paper	NERIFY	GPT-5-thinking	DeepCode	MetaGPT	ChatDev
Selected cameras K	16	✓	✓	✗	✗	✗
Entropy window size	9	✓	✓	✗	✗	✗
Entropy bins	16	✓	✓	256	✗	256
Mix ratio	0.5	✓	✓	✗	✗	✗
Learning rate	5e-4	✓	✓	1e-3	✓	1e-3
Batch size	4096	✓	✓	1024	16	32
W Score	–	1.00	0.93	0.17	0.17	0.08

319 2.1.9. Conclusion

320 All baseline systems attempt to implement KeyNeRF with
 321 varying levels of sophistication. GPT-5 thinking demon-
 322 strates deep understanding of both the algorithms and frame-
 323 work integration, providing a near-complete solution with
 324 proper Nerfstudio integration through RayGenerator and Ner-
 325 factoModel. DeepCode provides extensive standalone code
 326 but completely misses Nerfstudio integration. MetaGPT pro-
 327 duces clean structure but implements vanilla NeRF without
 328 any KeyNeRF innovations. ChatDev demonstrates funda-
 329 mental misunderstandings with placeholder functions and
 330 inappropriate GUI components. The critical difference is
 331 framework integration quality: GPT-5 thinking achieves
 332 85% trainable code with minimal adaptation needed, while
 333 ChatDev, MetaGPT and DeepCode produce 0% trainable
 334 implementations. Only NERIFY achieves 100% immediate
 335 trainability as a complete Nerfstudio plugin, though GPT-
 336 5 thinking comes remarkably close with its sophisticated
 337 understanding of both algorithmic details and framework
 338 requirements.

2.2. FastNeRF: High-Fidelity Neural Rendering at 200 FPS

339
340

2.2.1. Paper Overview

FastNeRF introduces a factorized neural radiance field architecture that decomposes the traditional monolithic MLP into position and direction dependent networks. The key innovation lies in representing radiance as a factorized inner product between a deep radiance map from the position network and directional weights from the viewing direction network. This factorization enables aggressive caching strategies that achieve 200 FPS rendering while maintaining visual quality comparable to vanilla NeRF. The paper fundamentally changes how neural radiance fields balance quality versus speed, demonstrating that architectural factorization rather than model compression provides the path to real-time neural rendering.

342
343
344
345
346
347
348
349
350
351
352
353
354
355

2.2.2. Implementation Overview

Table 10. FastNeRF Implementation Summary Across Multi-Agent Systems

Aspect	NERIFY	GPT-5-thinking	DeepCode	MetaGPT	ChatDev
Lines of Code	565	280	2915	420	385
File Organization	Plugin	Modular	Comprehensive	Multi-file	Multi-file
Factorized MLPs	✓	✓	✓	✓	Simplified
Deep Radiance Map	✓	✓	✓	Partial	✗
Nerfstudio Integration	✓	Partial	✗	✗	✗
Cache Infrastructure	Ready	Sketched	✓	✗	✗
Trainable	✓	✗	✗	✗	✗

2.2.3. Novel Components

356

Table 11. Novel Components in FastNeRF with Importance Weights

ID	Component	Weight w_i
C1	Position MLP: $F_{pos}(p) \rightarrow (\sigma, u, v, w)$	0.15
C2	Direction MLP: $F_{dir}(d) \rightarrow \beta$	0.12
C3	Deep radiance map factorization $(u, v, w) \in \mathbb{R}^{D \times 3}$	0.13
C4	Directional weights $\beta \in \mathbb{R}^D$	0.10
C5	Inner product color computation: $c = \sum_i \beta_i(u_i, v_i, w_i)$	0.12
C6	Positional encoding with $L = 10$ for positions	0.08
C7	Directional encoding with $L = 1$ for viewing directions	0.08
C8	Cache grid structure for F_{pos} outputs	0.10
C9	Spherical cache for F_{dir} outputs	0.08
C10	Training hyperparameters (384/128 widths, 8/4 layers)	0.04

2.2.4. Quantitative Metrics

357

2.2.5. Component-by-Component Analysis

358

Component C1: Position MLP Architecture The position network processes 3D coordinates to produce density and a deep radiance map. The paper specifies an 8-layer MLP with 384 hidden units, outputting σ and 3D values representing the factorized radiance components.

359
360
361
362
363
364
365
366

```
1 # \nerify : Full factorized position network with
2 correct dimensions
```

Table 12. FastNeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	0.95	0.05	0.00	1.00	0.95
GPT-5 thinking	0.72	0.18	0.10	0.90	0.65
DeepCode	0.85	0.10	0.05	0.95	0.81
MetaGPT	0.45	0.35	0.20	0.80	0.36
ChatDev	0.30	0.40	0.30	0.70	0.21

```

367   2 self.pos_mlp = _MLP(
368     3   in_dim=pos_in, # 63 dims after encoding
369     4   num_layers=config.pos_num_layers, # 8 layers
370     5   layer_width=config.pos_layer_width, # 384 units
371     6   out_dim=1 + 3 * D, # sigma + (u,v,w)
372   )
373   8 def forward(self, positions):
374     9   encoded = self.pos_encoding(positions)
375    10   out = self.pos_mlp(encoded)
376    11   sigma = self.softplus(out[..., :1])
377    12   uvw = out[..., 1:].reshape(*out.shape[:-1], D, 3)
378   13   return sigma, uvw
379

```

Listing 12. C1 Implementation: NERFIFY (Score: 1.0)

```

380   1 # GPT-5: Correct architecture but missing
381     implementation details
382   2 self.pos_mlp = MLP(
383     3   in_dim=pos_in_dim,
384     4   num_layers=config.pos_mlp_layers,
385     5   layer_width=config.pos_mlp_width,
386     6   out_dim=pos_out_dim, # 1 + 3*D
387     7   skip_connections=[4], # similar to NeRF
388   8 )
389   9 # ERROR: Implementation incomplete,
390     populate_nerf_helpers not provided
391

```

Listing 13. C1 Implementation: GPT-5 (Score: 0.9)

```

393   1 # DeepCode: Comprehensive with cache-aware design
394   2 class FastNeRFPoSMLP(nn.Module):
395     3   def __init__(self, D=8, W=384, input_ch=63,
396       n_components=8):
397     4     super().__init__()
398     5     self.pts_linears = nn.ModuleList(
399       6       [nn.Linear(input_ch, W)] +
400       7       [nn.Linear(W, W) for _ in range(D-1)])
401     8   self.output_linear = nn.Linear(W, 1 +
402     n_components * 3)
403   10 # CORRECT: Full implementation with proper
404     factorization
405

```

Listing 14. C1 Implementation: DeepCode (Score: 0.95)

```

408   1 # MetaGPT: Simplified architecture missing skip
409     connections
410   2 class PositionDependentNetwork(nn.Module):
411     3   def __init__(self, hidden_dim=384, output_dim=27)
412     4     :
413     5       self.layers = nn.ModuleList()
414     6       for _ in range(num_layers - 2):
415     7         self.layers.append(nn.Linear(hidden_dim,
416           hidden_dim))
417   7 # ERROR: No positional encoding integration
418

```

Listing 15. C1 Implementation: MetaGPT (Score: 0.6)

```

1 # ChatDev: Incorrect understanding of factorization
2 class FPos(nn.Module):
3   def forward(self, p):
4     x = F.relu(self.fc1(p_encoded))
5     output = self.fc3(x)
6     return output[:, 0], output[:, 1:4] # Wrong:
7     # returns (u,v,w) as 3D
7   # CRITICAL ERROR: Misunderstands deep radiance
7   map dimensions

```

Listing 16. C1 Implementation: ChatDev (Score: 0.3)

Analysis: NERFIFY correctly implements the full factorized architecture with proper dimensionality. GPT-5 understands the structure but leaves critical implementation details as placeholders. DeepCode provides the most comprehensive implementation outside of NERFIFY but lacks framework integration. MetaGPT simplifies the architecture while ChatDev fundamentally misunderstands the factorization concept.

Component C2: Direction MLP Architecture The direction network processes viewing directions to produce weights for combining the deep radiance map components. The paper specifies a 4-layer MLP with 128 hidden units.

```

1 # \nerfify\ : Correct lightweight direction network
2 self.dir_mlp = _MLP(
3   in_dim=dir_in, # 9 dims with L=1 encoding
4   num_layers=config.dir_num_layers, # 4 layers
5   layer_width=config.dir_layer_width, # 128 units
6   out_dim=D, # beta vector
7 )

```

Listing 17. C2 Implementation: NERFIFY (Score: 1.0)

```

1 # DeepCode: Proper direction network with activation
2 class FastNeRFDiMLP(nn.Module):
3   def __init__(self, D=4, W=128, dir_enc_L=1,
4     n_components=8):
5     self.dir_linears = nn.ModuleList([nn.Linear(
6       input_ch, W)] +
7       [nn.Linear(W, W) for _ in range(D-1)])
8   self.output_linear = nn.Linear(W, n_components)

```

Listing 18. C2 Implementation: DeepCode (Score: 0.9)

Analysis: Both NERFIFY and DeepCode correctly implement the lightweight direction network. GPT-5 provides the structure, while MetaGPT and ChatDev misunderstand the role of directional processing in the factorization.

Component C3: Deep Radiance Map Factorization The core innovation of FastNeRF lies in representing radiance as a tensor product between position-dependent and direction-dependent components.

```

1 # \nerfify\ : Proper tensor factorization
2 uvw = pos_out[..., 1:].reshape(*pos_out.shape[:-1],
3   self.D, 3)
3 beta = self.dir_mlp(encoded_dirs) # [batch, D]
4 rgb = torch.sum(beta[..., :, None] * uvw, dim=-2) # Inner product

```

479

Listing 19. C3 Implementation: NERIFY (Score: 1.0)

```

480 1 # ChatDev: Complete misunderstanding of factorization
481 2 return output[:, 0], output[:, 1:4] # Returns 3D
482 3     vector instead of Dx3
483 4 # CRITICAL ERROR: No tensor product, treats as
484 5     regular RGB output
485

```

Listing 20. C3 Implementation: ChatDev (Score: 0.0)

Analysis: NERIFY and DeepCode correctly implement the tensor factorization that defines FastNeRF. GPT-5 shows understanding but lacks implementation. MetaGPT attempts factorization with errors, while ChatDev completely misses this fundamental concept.

Component C5: Inner Product Color Computation The final RGB color emerges from the weighted sum of deep radiance map components using directional weights.

```

496 1 # \nerify\ : Mathematically correct inner product
497 2 rgb = torch.sum(beta[..., :, None] * uvw, dim=-2)
498 3 rgb = self.sigmoid(rgb) # Activation for valid color
499 4     range
500

```

Listing 21. C5 Implementation: NERIFY (Score: 1.0)

```

501 1 # MetaGPT: Attempts inner product with shape errors
502 2 color = torch.sum(weights.unsqueeze(-1) *
503 3     radiance_map.view(-1, 8, 3), dim=1)
504 4 # ERROR: Incorrect tensor reshaping, breaks with
505 5     batch processing
506

```

Listing 22. C5 Implementation: MetaGPT (Score: 0.5)

Analysis: The inner product computation separates successful implementations from failures. NERIFY handles all tensor operations correctly, while baselines struggle with proper broadcasting and dimension management.

Component C8: Cache Grid Structure FastNeRF's speed derives from caching position network outputs in a 3D grid for inference.

```

515 1 # DeepCode: Complete cache implementation
516 2 def build_pos_cache(f_pos, bounds, grid_res,
517 3     n_components=8):
518 4     positions, grid_shape = create_position_grid(
519 5         bounds, grid_res)
520 6     outputs = batch_eval_mlp_on_grid(f_pos, positions
521 7         )
522 8     cache = outputs.reshape(*grid_shape, 1 +
523 9         n_components * 3)
524 10    return cache # [k, k, k, 1+D*3] tensor
525

```

Listing 23. C8 Implementation: DeepCode (Score: 1.0)

```

527 1 # GPT-5: Acknowledges caching but provides no
528 2     implementation
529 3 # (Optional) How to add **caching** like the paper
530 4 # 1. **Offline cache builder** script:
531 5     * Define a world-space bounding box and
532 6         resolution k
533 7 # ERROR: Only provides instructions, no actual code
534

```

Listing 24. C8 Implementation: GPT-5 (Score: 0.3)

Analysis: DeepCode excels at cache implementation with comprehensive utilities for grid construction and interpolation. NERIFY provides the architecture ready for caching. GPT-5 acknowledges the concept but defers implementation, while MetaGPT and ChatDev ignore caching entirely.

2.2.6. Scoring Analysis

Table 13. FastNeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	Weighted Avg _{LLM}
NERIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.7	0.6	1.0	0.94
GPT-5 thinking	0.9	0.8	0.8	0.7	0.9	0.8	0.3	0.2	0.5	0.67	
DeepCode	0.95	0.9	0.95	0.9	0.9	0.9	0.8	1.0	0.9	0.3	0.85
MetaGPT	0.6	0.5	0.4	0.5	0.5	0.7	0.6	0.0	0.0	0.3	0.41
ChatDev	0.3	0.3	0.0	0.2	0.1	0.8	0.7	0.0	0.0	0.2	0.26

2.2.7. Why Baselines Fail Despite Component Scores

GPT-5 Extended Thinking (Score: 0.67 components, 0% trainable)

- Strengths:** Demonstrates deep understanding of FastNeRF architecture and provides correct Nerfstudio integration patterns
- Fatal Issues:**
 - Leaves critical implementation as placeholder: `raise NotImplementedError`
 - Missing essential `populate_modules` implementation
 - No actual training loop connection
- Result:** Code serves as documentation rather than executable implementation

DeepCode (Score: 0.85 components, 0% trainable)

- Strengths:** Most comprehensive implementation with full caching infrastructure and extensive testing
- Fatal Issues:**
 - Complete absence of Nerfstudio framework integration
 - No training pipeline or data loading infrastructure
 - Implements components in isolation without system integration
- Result:** Excellent reference implementation that cannot train without significant engineering effort

MetaGPT (Score: 0.41 components, 0% trainable)

- Strengths:** Attempts modular organization with separate training and rendering components
- Fatal Issues:**
 - Fundamental misunderstanding of tensor dimensions in factorization
 - Missing ray sampling and volume rendering pipeline
 - Placeholder data loading with no actual dataset interface
- Result:** Structurally incomplete implementation that fails at runtime

ChatDev (Score: 0.26 components, 0% trainable)

- 576 • **Strengths:** Creates organized file structure with GUI in-
 577 terface
 578 • **Fatal Issues:**
 579 – Complete misunderstanding of deep radiance map fac-
 580 torization
 581 – Implements Tkinter GUI instead of neural rendering
 582 pipeline
 583 – Missing entire volume rendering and ray marching in-
 584 frastructure
 585 • **Result:** Application framework without actual FastNeRF
 586 implementation

587 2.2.8. Hyperparameter Fidelity

Table 14. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5	DeepCode	MetaGPT	ChatDev
Position MLP layers	8	✓	✓	✓	Partial	✗
Position MLP width	384	✓	✓	✓	✓	✗
Direction MLP layers	4	✓	✓	✓	✓	✗
Direction MLP width	128	✓	✓	✓	✓	✗
Factor dimension D	8	✓	✓	Variable	Fixed	Wrong
Position encoding L	10	✓	✓	Variable	✓	✓
Direction encoding L	1	✓	Variable	Variable	✗	✗
Learning rate	5e-4	✓	✓	N/A	✓	✗
W Score	–	1.00	0.88	0.63	0.50	0.25

588 2.2.9. Conclusion

589 All baseline systems attempt to implement FastNeRF with
 590 varying degrees of sophistication and understanding. GPT-5
 591 demonstrates theoretical mastery but delivers skeleton code
 592 requiring substantial completion. DeepCode produces the
 593 most comprehensive component implementation with so-
 594 phisticated caching utilities, yet fails to integrate with any
 595 training framework. MetaGPT captures the high-level struc-
 596 ture but makes critical errors in tensor operations that prevent
 597 execution. ChatDev fundamentally misunderstands the pa-
 598 per, building a GUI application instead of implementing
 599 the core factorization algorithm. Despite DeepCode achiev-
 600 ing 85% component coverage and GPT-5 showing 67% un-
 601 derstanding, all baselines produce 0% trainable code due
 602 to missing framework integration, incomplete implemen-
 603 tations, or fundamental architectural errors. Only NERFIFY
 604 delivers immediately trainable code by properly implemen-
 605 ting the factorized architecture within the Nerfstudio frame-
 606 work, achieving 100% executability where all others fail
 607 completely.

608 2.3. NeRF: Neural Radiance Fields for View Syn- 609 thesis

610 2.3.1. Paper Overview

611 The original NeRF paper by Mildenhall et al. (2020) intro-
 612 duced neural radiance fields for photorealistic novel view
 613 synthesis. The method represents scenes as continuous 5D
 614 functions mapping 3D coordinates and 2D viewing direc-
 615 tions to volume density and color, rendered using classical

volume rendering techniques with hierarchical sampling for
 616 efficiency.
 617

618 2.3.2. Implementation Overview

Table 15. NeRF Implementation Summary Across Multi-Agent Systems

Aspect	NERFIFY	GPT-5-thinking	DeepCode	MetaGPT	ChatDev
Lines of Code	387	156	412	98	524
File Organization	Plugin	Instructions	Multi-file	Multi-file	Multi-file
Positional Encoding	✓	✓	✗	Partial	Attempted
MLP Architecture	✓	✓	Simplified	Partial	Simplified
Hierarchical Sampling	✓	✓	✗	✗	Attempted
Volume Rendering	✓	✓	Partial	✗	Partial
Trainable	✓	Partial	✗	✗	✗

619 2.3.3. Novel Components

Table 16. Novel Components in NeRF with Importance Weights

ID	Component	Weight w_i
C1	Positional encoding (L=10 for xyz, L=4 for direction)	0.15
C2	8-layer MLP with skip connection at layer 4	0.15
C3	View-dependent RGB prediction	0.10
C4	Hierarchical volume sampling (coarse + fine)	0.20
C5	Alpha compositing with transmittance	0.15
C6	Stratified sampling along rays	0.10
C7	NDC ray parameterization	0.05
C8	MSE photometric loss	0.05
C9	Adam optimizer with exponential decay	0.05

620 2.3.4. Quantitative Metrics

Table 17. NeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	0.95	0.02	0.03	0.97	9.2
GPT-5 thinking	0.72	0.15	0.13	0.79	7.1
ChatDev	0.35	0.42	0.23	0.51	4.8
MetaGPT	0.18	0.27	0.55	0.31	2.9
DeepCode	0.42	0.38	0.20	0.56	5.3

621 2.3.5. Component-by-Component Analysis

Component C1: Positional Encoding The pa-
 622 per specifies sinusoidal positional encoding with
 623 $L = 10$ for positions and $L = 4$ for viewing
 624 directions, encoding input coordinates as $\gamma(p) =$
 $[p, \sin(2^0\pi p), \cos(2^0\pi p), \dots, \sin(2^{L-1}\pi p), \cos(2^{L-1}\pi p)]$.

```
1 # \nerfify\ : Proper Nerfstudio positional encoding
  configuration
2 position_encoding_config=NeRFModelConfig.
  position_encoding_config.__class__(
3     num_frequencies=10
4 ),
5 direction_encoding_config=NeRFModelConfig.
  direction_encoding_config.__class__(
6     num_frequencies=4
```

637
638
639

```
7 ),
8 # CORRECT: Exact paper specification with L=10 and L
=4
```

Listing 25. C1 Implementation: NERFIFY (Score: 1.0)

641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656

```
1 # ChatDev: Custom positional encoding implementation
2 class PositionalEncoding:
3     def __init__(self, L):
4         self.L = L
5     def encode(self, x):
6         encoded = []
7         for i in range(self.L):
8             encoded.append(torch.sin(2 ** i * torch.
pi * x))
9             encoded.append(torch.cos(2 ** i * torch.
pi * x))
10        return torch.cat(encoded, dim=-1)
11 # ERROR: No differentiation between position/
direction encoding levels
```

Listing 26. C1 Implementation: ChatDev (Score: 0.6)

657
658
659
660
661
662
663
664
665
666
667
668

```
1 # MetaGPT: Placeholder implementation
2 class PositionalEncoding(nn.Module):
3     def __init__(self, L=10, d_model=256):
4         super(PositionalEncoding, self).__init__()
5         self.L = L
6         self.d_model = d_model
7     def forward(self, x):
8         # Implement positional encoding as described
9         in the paper
10        pass
10 # CRITICAL ERROR: Not implemented, just placeholder
```

Listing 27. C1 Implementation: MetaGPT (Score: 0.3)

670
671
672
673
674

Analysis: NERFIFY correctly configures both position and direction encoding with proper frequency counts. ChatDev implements the mathematical formula but misses the crucial distinction between position and direction encoding levels. MetaGPT provides only placeholders.

675
676
677

Component C2: MLP Architecture The paper specifies an 8-layer MLP with 256 hidden units and a skip connection at layer 4.

678
679
680
681
682
683
684
685

```
1 # \nerfify\ : Correct architecture configuration
2 field_num_layers=8,
3 field_layer_width=256,
4 field_heads_layer_width=128,
5 field_skip_connections=(4,),
6 # CORRECT: Exact paper specification with skip at
layer 4
```

Listing 28. C2 Implementation: NERFIFY (Score: 1.0)

687
688
689
690
691
692
693
694

```
1 # ChatDev: Simplified MLP without skip connection
2 self.fc_layers = nn.ModuleList([
3     nn.Linear(30, 256) for _ in range(8)
4 ])
5 # ERROR: Missing skip connection at layer 4
6 # ERROR: Incorrect input dimension (should be 63 for
L=10)
```

Listing 29. C2 Implementation: ChatDev (Score: 0.5)

696
697
698
699

Analysis: NERFIFY properly configures the 8-layer architecture with skip connections. ChatDev creates 8 layers but misses the critical skip connection and has incorrect input dimensions.

Component C3: View-Dependent RGB Prediction The paper predicts RGB values conditioned on both position features and viewing direction.

700
701
702
703
704
705
706
707
708
709
710

```
1 # \nerfify\ : Inherits view-dependent color from
NerfactoField
2 # View direction properly encoded and concatenated
with features
3 # CORRECT: Full view-dependent RGB implementation via
parent class
```

Listing 30. C3 Implementation: NERFIFY (Score: 1.0)

711
712
713
714
715
716
717
718
719
720

```
1 # ChatDev: View-dependent color prediction
2 self.color_layer = nn.Sequential(
3     nn.Linear(256 + 8, 128),
4     nn.ReLU(),
5     nn.Linear(128, 3),
6     nn.Sigmoid()
7 )
8 # ISSUE: Wrong concatenation dimension (should be 256
+ 27 for L=4)
```

Listing 31. C3 Implementation: ChatDev (Score: 0.7)

Analysis: NERFIFY correctly implements view-dependent RGB through Nerfstudio's field architecture. ChatDev attempts view dependence but miscalculates encoded direction dimensions.

722
723
724
725

Component C4: Hierarchical Volume Sampling The paper uses coarse-to-fine sampling with 64 coarse and 128 fine samples.

726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748

```
1 # \nerfify\ : Proper hierarchical sampling
configuration
2 num_coarse_samples=64,
3 num_importance_samples=128,
4 # CORRECT: Exact paper specification for two-stage
sampling
```

Listing 32. C4 Implementation: NERFIFY (Score: 1.0)

737
738
739
740
741
742
743
744
745
746
747
748

```
1 # ChatDev: Attempted hierarchical sampling
2 def sample_fine(self, coarse_samples,
coarse_densities):
3     weights = coarse_densities / torch.sum(
coarse_densities, dim=-1, keepdim=True)
4     fine_indices = torch.multinomial(weights, self.
N_f, replacement=True)
5     fine_samples = coarse_samples['xyz'][fine_indices
]
6 # CRITICAL ERROR: Wrong importance sampling (should
use CDF inverse)
```

Listing 33. C4 Implementation: ChatDev (Score: 0.4)

Analysis: NERFIFY configures proper coarse and fine sampling counts. ChatDev attempts importance sampling but implements it incorrectly, using simple multinomial instead of inverse CDF sampling.

750
751
752
753

2.3.6. Scoring Analysis

2.3.7. Why Baselines Fail Despite Component Scores

GPT-5 thinking (Score: 0.79 components, 50% trainable)

- **Strengths:** Correctly identifies Nerfstudio's built-in implementation, provides proper configuration

754
755
756
757
758

Table 18. NeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	C9	Weighted Avg _{LLM}
NERFIFY	1.0	1.0	1.0	1.0	1.0	1.0	0.9	1.0	1.0	0.97
GPT-5 thinking	1.0	1.0	1.0	1.0	1.0	0.8	0.0	0.8	0.9	0.79
ChatDev	0.6	0.5	0.7	0.4	0.6	0.5	0.0	0.8	0.5	0.51
MetaGPT	0.3	0.4	0.3	0.0	0.2	0.3	0.0	0.6	0.7	0.31
DeepCode	0.0	0.6	0.8	0.2	0.7	0.6	0.0	0.9	0.8	0.56

759 • **Fatal Issues:**

- Provides instructions rather than executable code
- Requires manual CLI command construction

760 • **Result:** Works only when user correctly follows instruc-
761 tions

764 **ChatDev (Score: 0.51 components, 0% trainable)**

- 765 • **Strengths:** Attempts full implementation with GUI and
766 modular structure
- 767 • **Fatal Issues:**
- No Nerfstudio integration whatsoever
 - Incorrect importance sampling implementation
 - Missing ray generation logic
- 771 • **Result:** Code runs but cannot train on real data

772 **MetaGPT (Score: 0.31 components, 0% trainable)**

- 773 • **Strengths:** Proper project structure with configuration
774 management
- 775 • **Fatal Issues:**
- All core functions are placeholders with ‘pass’ state-
777 ments
 - No actual implementation of any algorithm
- 779 • **Result:** Skeleton code that requires complete implemen-
780 tation

781 **DeepCode (Score: 0.56 components, 0% trainable)**

- 782 • **Strengths:** Attempts volume rendering and loss computa-
783 tion
- 784 • **Fatal Issues:**
- Missing positional encoding entirely
 - No data loading implementation
 - Incompatible with any training framework
- 788 • **Result:** Partial implementation missing critical compo-
789 nents

790 **2.3.8. Hyperparameter Fidelity**

Table 19. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5-thinking	MetaGPT	ChatDev	DeepCode
Coarse samples	64	✓	✓	✗	✓	✗
Fine samples	128	✓	✓	✗	✓	✗
Batch size	4096	✓	✓	✓	✓	✗
Learning rate	5e-4	✓	✓	✓	✓	✗
Position L	10	✓	✓	✗	✗	✗
Direction L	4	✓	✓	✗	✗	✗
W Score	–	1.00	1.00	0.33	0.67	0.00

791 **2.3.9. Conclusion**

792 All baseline systems attempt to implement vanilla NeRF with
793 varying levels of sophistication. GPT-5 thinking correctly
794 identifies that Nerfstudio already contains the implementa-
795 tion but provides instructions rather than code. ChatDev
796 produces the most complete standalone attempt but lacks
797 framework integration. MetaGPT generates only skeleton
798 code with placeholders. DeepCode attempts key components
799 but misses positional encoding entirely. While GPT-5 think-
800 ing achieves 50% trainability through instruction following
801 for this well-known method already in Nerfstudio, ChatDev,
802 MetaGPT, and DeepCode achieve 0% trainable implemen-
803 tations versus NERFIFY’s 100% immediately trainable plugin.

804 **2.4. Deblur-NeRF: NeRF from Motion-Blurred Im-
805 ages**

806 **2.4.1. Paper Overview**

807 Deblur-NeRF addresses the fundamental challenge of train-
808 ing Neural Radiance Fields from motion-blurred images
809 by introducing the Deformable Sparse Kernel (DSK). The
810 method models blur formation explicitly through learnable
811 sparse ray distributions per pixel, enabling sharp novel view
812 synthesis from blurry training data. The key innovation
813 lies in jointly optimizing the underlying sharp NeRF rep-
814 resentation while learning per-image blur kernels through
815 differentiable rendering.

816 **2.4.2. Implementation Overview**

Table 20. Deblur-NeRF Implementation Summary Across Multi-Agent Systems

Aspect	NERFIFY	GPT-5-thinking	DeepCode	MetaGPT	ChatDev
Lines of Code	736	424	1487	285	156
File Organization	Plugin	Modular	Multi-file	Multi-file	Single
DSK Module	✓	✓	✓	Partial	✗
Nerfstudio Integration	✓	✓	✗	✗	✗
Gamma Correction	✓	✓	✓	✗	Simplified
Alignment Loss	✓	✓	✓	✗	✗
Trainable	✓	Partial	✗	✗	✗

817 **2.4.3. Novel Components**

Table 21. Novel Components in Deblur-NeRF with Importance Weights

ID	Component	Weight w_i
C1	Deformable Sparse Kernel (DSK) MLP	0.25
C2	Per-view embedding l (32-dim)	0.10
C3	Origin optimization Δo	0.10
C4	Kernel blending in linear space	0.15
C5	Gamma correction $g(c) = c^{1/2.2}$	0.10
C6	Alignment loss \mathcal{L}_{align}	0.10
C7	Sharp rendering at inference	0.05
C8	Nerfstudio framework integration	0.15

924 **Component C3: Origin Optimization** Ray origins are
 925 shifted by learned Δo to model camera motion blur.
 926

```
1 # \nerify\ : Full origin optimization with proper
  shifting
2 dsk_rays.origins = dsk_rays.origins + delta_o.view
  (-1, 3)
3 enable_origin_opt=True, # Config flag
```

Listing 44. C3 Implementation: NERIFY (Score: 1.0)

```
1 # GPT-5: Correct origin shifting
2 dsk_rays.origins = dsk_rays.origins + delta_o.view
  (-1, 3)
```

Listing 45. C3 Implementation: GPT-5 (Score: 1.0)

```
1 # DeepCode: Origin optimization with config control
2 self.kernel_origin_on = config.get('kernel_origin_on',
  , True)
3 origin_offsets = construct_kernel_rays(...)[4]
```

Listing 46. C3 Implementation: DeepCode (Score: 0.9)

```
1 # MetaGPT: Delta origin computed but not applied
2 delta_origin = output[:, :3]
3 # ERROR: Never used to shift ray origins
```

Listing 47. C3 Implementation: MetaGPT (Score: 0.3)

```
1 # ChatDev: No origin optimization
2 # MISSING: No ray origin shifting
```

Listing 48. C3 Implementation: ChatDev (Score: 0.0)

Analysis: NERIFY and GPT-5 properly shift ray origins. DeepCode implements it with configuration control. MetaGPT computes deltas but never applies them. ChatDev lacks the feature.

957 **Component C4: Linear Space Blending** Colors are
 958 blended in linear RGB space before gamma correction.
 959

```
1 # \nerify\ : Proper linear blending then gamma
2 rgb_linear = (w[... , None] * rgb_fine_all).sum(dim=1)
3 rgb_blurry = self._apply_gamma(rgb_linear)
```

Listing 49. C4 Implementation: NERIFY (Score: 1.0)

```
1 # GPT-5: Correct linear space blending
2 rgb_lin_blurry = torch.sum(weights_unsq * rgb_lin,
  dim=1)
3 rgb_blurry = self._gamma_correct(rgb_lin_blurry)
```

Listing 50. C4 Implementation: GPT-5 (Score: 1.0)

```
1 # DeepCode: Blending with kernel weights
2 pred_blurry = torch.sum(all_rgb * kernel_weights.
  unsqueeze(-1), dim=1)
3 if self.gamma_on:
4     pred_blurry = gamma_correction(pred_blurry)
```

Listing 51. C4 Implementation: DeepCode (Score: 0.8)

```
1 # MetaGPT: No explicit kernel blending
2 # ERROR: Missing weighted sum over kernel points
```

Listing 52. C4 Implementation: MetaGPT (Score: 0.2)

```
1 # ChatDev: No kernel blending
2 # MISSING: No multi-ray blending
```

Listing 53. C4 Implementation: ChatDev (Score: 0.0)

Analysis: NERIFY and GPT-5 correctly implement linear blending followed by gamma correction. DeepCode has the right approach. MetaGPT and ChatDev miss this entirely.

981 **Component C5: Gamma Correction** Applies camera
 982 response function $g(c) = c^{1/2.2}$.
 983

```
1 # \nerify\ : Exact gamma value from paper
2 def _apply_gamma(self, c_lin: Tensor) -> Tensor:
3     return torch.clamp(c_lin, min=0.0) ** self.
  one_over_gamma # 1/2.2
```

Listing 54. C5 Implementation: NERIFY (Score: 1.0)

```
1 # GPT-5: Correct gamma implementation
2 def _gamma_correct(self, c_lin: Tensor) -> Tensor:
3     return torch.clamp(c_lin, min=0.0) ** self.
  one_over_gamma
```

Listing 55. C5 Implementation: GPT-5 (Score: 1.0)

```
1 # DeepCode: Proper gamma correction
2 def gamma_correction(color, gamma=2.2):
3     return color ** (1.0 / gamma)
```

Listing 56. C5 Implementation: DeepCode (Score: 1.0)

```
1 # MetaGPT: No gamma correction
2 # MISSING: No CRF modeling
```

Listing 57. C5 Implementation: MetaGPT (Score: 0.0)

```
1 # ChatDev: Different gamma implementation
2 def gamma_correction(c):
3     return torch.clamp(c ** (1 / 2.2), min=0)
4 # ERROR: Applied in wrong context (utils not model)
```

Listing 58. C5 Implementation: ChatDev (Score: 0.3)

Analysis: NERIFY, GPT-5, and DeepCode correctly implement gamma=2.2. ChatDev has the function but uses it incorrectly. MetaGPT misses it entirely.

1002 **Component C6: Alignment Loss** Regularizes kernel de-
 1003 formations: $\mathcal{L}_{align} = \|q_0 - p\|_2 + \lambda_o \|\Delta o\|_2$.
 1004

```
1 # \nerify\ : Complete alignment loss
2 align = torch.norm(dq0, dim=-1).mean() + \
3   self.config.lambda_o * torch.norm(do0, dim
  =-1).mean()
4 loss_dict['alignment_loss'] = self.config.
  lambda_align * align
```

Listing 59. C6 Implementation: NERIFY (Score: 1.0)

```
1 # GPT-5: Proper alignment with both terms
2 align_pos = torch.mean((q0 - p).pow(2).sum(dim=-1).
  sqrt())
3 align_origin = torch.mean(delta_o0.pow(2).sum(dim=-1)
  .sqrt())
4 align_loss = align_pos + self.config.lambda_origin *
  align_origin
```

Listing 60. C6 Implementation: GPT-5 (Score: 1.0)

1039
1040
1041
1042
1043
1044
1045
1046

```

1 # DeepCode: Alignment loss function
2 def alignment_loss(pixel_coords, kernel_offsets,
3     origin_offsets, lambda_o):
4     loss_align = alignment_loss(pixel_coords,
5         kernel_offsets,
6             origin_offsets,
7         lambda_o=self.lambda_o)

```

Listing 61. C6 Implementation: DeepCode (Score: 0.9)

1048
1049
1050

```

1 # MetaGPT: No alignment loss
2 # MISSING: No regularization on deformations

```

Listing 62. C6 Implementation: MetaGPT (Score: 0.0)

1052
1053
1054

```

1 # ChatDev: No alignment loss
2 # MISSING: No kernel regularization

```

Listing 63. C6 Implementation: ChatDev (Score: 0.0)

1056
1057
1058

Analysis: NERFIFY and GPT-5 implement the complete two-term alignment loss. DeepCode has it as a separate function. MetaGPT and ChatDev lack this regularization.

1059

2.4.6. Scoring Analysis

Table 23. Deblur-NeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	Weighted Avg _{LLM}
NERFIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.00
GPT-5 thinking	0.9	1.0	1.0	1.0	1.0	1.0	0.8	0.7	0.94
DeepCode	0.8	0.8	0.9	0.8	1.0	0.9	0.5	0.0	0.71
MetaGPT	0.4	0.0	0.3	0.2	0.0	0.0	0.0	0.0	0.11
ChatDev	0.0	0.0	0.0	0.0	0.3	0.0	0.0	0.0	0.04

1060
1061
1062
1063
1064
1065
1066
1067
1068
1069
1070
1071

2.4.7. Why Baselines Fail Despite Component Scores

GPT-5 Extended Thinking (Score: 0.94 components, 50% trainable)

- Strengths:** Near-complete DSK implementation with proper Nerfstudio integration, correct gamma correction and alignment loss
- Fatal Issues:**
 - Missing proper NeRF field inheritance structure
 - Incomplete pipeline configuration for distributed training
- Result:** Trainable for simple scenes but fails on complex multi-GPU setups

1072
1073
1074
1075
1076
1077
1078
1079
1080
1081

DeepCode (Score: 0.71 components, 0% trainable)

- Strengths:** Comprehensive evaluation framework with ablation studies, proper DSK concepts
- Fatal Issues:**
 - No Nerfstudio integration whatsoever
 - Custom training loop incompatible with ns-train command
 - Missing scene normalization and camera conventions
- Result:** Standalone implementation that cannot interface with Nerfstudio datasets

MetaGPT (Score: 0.11 components, 0% trainable)

- Strengths:** Multi-file organization with basic NeRF structure
- Fatal Issues:**
 - Missing critical DSK components (view embedding, alignment loss)
 - No framework integration
 - Incorrect positional encoding implementation
- Result:** Incomplete implementation missing core paper contributions

ChatDev (Score: 0.04 components, 0% trainable)

- Strengths:** Basic GUI interface for user interaction
- Fatal Issues:**
 - Completely missing DSK module
 - No blur modeling whatsoever
 - Tkinter GUI instead of proper training pipeline
- Result:** Vanilla NeRF implementation unrelated to Deblur-NeRF paper

2.4.8. Hyperparameter Fidelity

Table 24. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5-thinking	DeepCode	MetaGPT	ChatDev
N (kernel points)	5	✓	✓	✓	✗	✗
λ_{align}	0.1	✓	✓	✓	✗	✗
λ_o	10.0	✓	✓	✓	✗	✗
View embed dim	32	✓	✓	✓	✗	✗
Learning rate	5e-4	✓	✓	✓	✓	✓
ϵ (DSK)	0.1	✓	✓	✓	✗	✗
Gamma	2.2	✓	✓	✓	✗	~
Iterations	200k	✓	✓	✓	✓	✗
W Score	-	1.00	1.00	1.00	0.25	0.19

2.4.9. Conclusion

All baseline systems attempt to implement Deblur-NeRF with varying degrees of sophistication. GPT-5 Extended Thinking achieves remarkable component coverage (94%) with proper DSK implementation and Nerfstudio awareness, occasionally producing trainable code for simple cases. DeepCode provides extensive evaluation infrastructure but lacks framework integration. MetaGPT and ChatDev fail to capture the paper's core contributions, with ChatDev entirely missing the DSK module. The fundamental gap between component scores and trainability demonstrates that domain-specific framework integration is essential - only NERFIFY produces immediately trainable code with 100% success rate versus 0% for ChatDev, MetaGPT, and DeepCode, with GPT-5 achieving partial success on simpler scenes.

3. Novelty Coverage on Set 4 Papers

3.1. BioNeRF: Biologically Plausible Neural Radiance Fields

BioNeRF introduces a biologically-inspired architecture that fundamentally reimagines neural radiance fields through cog-

nitive filtering and stateful memory mechanisms. The paper presents parallel positional feature extraction pathways, four distinct cognitive filters, pre-modulation mechanisms, and a recursive memory update system that maintains context across forward passes. This architecture mimics biological neural processing patterns to achieve superior view synthesis quality through memory-conditioned contextual inference.

1128 3.1.1. Implementation Overview

Table 25. BioNeRF Implementation Summary Across All Baselines

Aspect	NERIFY	GPT-5	DeepSeek R1	Paper2Code	AutoP2C
Lines of Code	485	182	271	1263	618
File Organization	Plugin	Modular	Single	Multi-file	Multi-file
Parallel MLPs	✓	✓	✓	Simplified	Attempted
Cognitive Filters	✓	✓	✓	Partial	✗
Memory Mechanism	✓	✓	Partial	✗	✗
Nerfstudio Ready	✓	✗	✗	✗	✗
Trainable	✓	✗	✗	✗	✗

Note: NERFIFY always produces trainable code as a complete Nerfstudio plugin

1129 3.1.2. Novel Components

Table 26. Novel Components in BioNeRF with Importance Weights

ID	Component	Weight w_i
C1	Parallel MLPs (M_Δ and M_c) for positional features	0.10
C2	Cognitive filter $f_\Delta = \sigma(W_\Delta^f h_\Delta + b_\Delta^f)$	0.10
C3	Cognitive filter $f_c = \sigma(W_c^f h_c + b_c^f)$	0.10
C4	Memory filter $f_\psi = \sigma(W_\psi^f [h_\Delta, h_c] + b_\psi^f)$	0.10
C5	Modulation filter $f_\mu = \sigma(W_\mu^f [h_\Delta, h_c] + b_\mu^f)$	0.10
C6	Pre-modulation $\gamma = \tanh(W_\gamma [h_\Delta, h_c] + b_\gamma)$	0.10
C7	Memory modulation $\mu = f_\mu \otimes \gamma$	0.10
C8	Memory update $\Psi = \tanh(W_\Psi (\mu + f_\psi \otimes \Psi) + b_\Psi)$	0.15
C9	Contextual embeddings $h'_\Delta = [\Psi \otimes f_\Delta, x]$, $h'_c = [\Psi \otimes f_c, d]$	0.10
C10	Density/color heads M'_Δ, M'_c	0.05

1130 3.1.3. Quantitative Metrics

Table 27. BioNeRF Implementation Coverage Metrics

Method	C	I	M	W	Score_{LLM}
NERIFY (Ours)	1.00	0.00	0.00	1.00	1.00
GPT-5	0.80	0.10	0.10	0.80	0.82
DeepSeek R1	0.70	0.20	0.10	0.70	0.75
Paper2Code	0.30	0.40	0.30	0.40	0.35
AutoP2C	0.10	0.30	0.60	0.10	0.15

3.1.4. Component-by-Component Analysis

Component C1: Parallel MLPs - All Baselines

The paper specifies two independent MLPs for positional feature extraction: M_Δ processes encoded 3D positions to produce h_Δ , while M_c processes the same inputs independently to produce h_c . Both networks have 3 layers with 256 hidden units and ReLU activations.

```
1 # \nerfify\ : Perfect parallel pathway implementation
2 class BioNeRFField(Field):
3     def __init__(self, hidden_dim: int = 256):
4         self.base_delta = MLP(
5             in_dim=pe_out, num_layers=3,
6             layer_width=self.hidden_dim,
7             out_dim=self.hidden_dim,
8             activation=nn.ReLU())
9         self.base_c = MLP(
10            in_dim=pe_out, num_layers=3,
11            layer_width=self.hidden_dim,
12            out_dim=self.hidden_dim,
13            activation=nn.ReLU())
14 # CORRECT: Independent MLPs with exact paper
15 # specifications
```

Listing 64. C1 Implementation: NERFIFY (Score: 1.0)

```
1 # GPT-5: Clean implementation with modular design
2 class BioNeRFField(nn.Module):
3     def __init__(self):
4         self.M_delta = self._make_mlp(63, 256, 256)
5         self.M_color = self._make_mlp(63, 256, 256)
6     def _make_mlp(self, in_dim, hidden, out_dim,
7                  n_layers=3):
8         layers = []
9         for i in range(n_layers):
10             layers.append(nn.Linear(
11                 in_dim if i == 0 else hidden, hidden))
12         layers.append(nn.ReLU(True))
13     return nn.Sequential(*layers)
14 # CORRECT: Proper parallel architecture
```

Listing 65. C1 Implementation: GPT-5 (Score: 0.9)

```
1 # DeepSeek R1: Comprehensive with proper encoding
2 class BioNeRF(nn.Module):
3     def __init__(self, hidden_dim=256, num_freqs=10):
4         self.encoded_3d_dim = 3 * (2 * num_freqs + 1)
5         self.M_delta = self._build_mlp(
6             self.encoded_3d_dim, hidden_dim,
7             hidden_dim, 3)
8         self.M_color = self._build_mlp(
9             self.encoded_3d_dim, hidden_dim,
9 # ISSUE: Correct structure but naming inconsistency
```

Listing 66. C1 Implementation: DeepSeek R1 (Score: 0.8)

```
1 # Paper2Code: Oversimplified to incorrect inputs  
2 class BioNeRFModel(nn.Module):  
3     def __init__(self):  
4         self.mlp_initial_delta = self._build_mlp(  
5             self.input_dim_xyz, hidden_size,  
6             hidden_size, 3)  
7         self.mlp_initial_c = self._build_mlp(  
8             self.input_dim_dir, hidden_size,  
9             hidden_size, 3)  
10    # ERROR: Uses different inputs for each MLP (xyz vs  
11        dir)  
12    # Should both process positional-encoded coordinates
```

Listing 67. C1 Implementation: Paper2Code (Score: 0.3)

```
1 # AutoP2C: Structure exists but fundamental error
1199
2 def positional_feature_extraction(self, inputs):
1200
3     encoded_inputs = positional_encoding(inputs, 10)
1201
4     # CRITICAL ERROR: Uses raw inputs, not encoded
1202
5     h_delta = self.mlp_delta(inputs) # Wrong!
1203
6     h_c = self.mlp_color(inputs) # Wrong!
1204
7     return h_delta, h_c
1205
1206
```

Listing 68. C1 Implementation: AutoP2C (Score: 0.1)

1208 **Analysis:** NERIFY correctly implements independent
 1209 parallel MLPs with exact paper specifications. GPT-5
 1210 achieves near-perfect implementation. DeepSeek R1 has
 1211 correct structure but minor naming issues. Paper2Code in-
 1212 correctly uses different inputs for each MLP. AutoP2C de-
 1213 fines encoding but critically fails to use it, passing raw inputs
 1214 instead.

1215 Component C2-C6: Cognitive Filtering - All Baselines

1216 The paper defines four cognitive filters and pre-
 1217 modulation: $f_\Delta = \sigma(W_\Delta^f h_\Delta + b_\Delta^f)$, $f_c = \sigma(W_c^f h_c + b_c^f)$,
 1218 $f_\psi = \sigma(W_\psi^f [h_\Delta, h_c] + b_\psi^f)$, $f_\mu = \sigma(W_\mu^f [h_\Delta, h_c] + b_\mu^f)$, and
 1219 $\gamma = \tanh(W_\gamma [h_\Delta, h_c] + b_\gamma)$.

```
1220 1 # \nerify\ : All filters properly defined in
1221 2     __init__
1222 3 def __init__(self):
1223 4     self.W_delta_f = nn.Linear(self.hidden_dim, self.
1224 5         hidden_dim)
1225 6     self.W_c_f = nn.Linear(self.hidden_dim, self.
1226 7         hidden_dim)
1227 8     self.W_psi_f = nn.Linear(2 * self.hidden_dim,
1228 9         self.hidden_dim)
1229 10    self.W_mu_f = nn.Linear(2 * self.hidden_dim, self.
1230 11        hidden_dim)
1231 12    self.W_gamma = nn.Linear(2 * self.hidden_dim,
1232 13        self.hidden_dim)
1233 14 def forward(self):
1234 15    f_delta = self._sigmoid(self.W_delta_f(h_delta))
1235 16    f_c = self._sigmoid(self.W_c_f(h_c))
1236 17    h_cat = torch.cat([h_delta, h_c], dim=-1)
1237 18    f_psi = self._sigmoid(self.W_psi_f(h_cat))
1238 19    f_mu = self._sigmoid(self.W_mu_f(h_cat))
1239 20    gamma = self._tanh(self.W_gamma(h_cat))
1240 21 # CORRECT: All filters with proper activations
```

Listing 69. C2-C6 Implementation: NERIFY (Score: 1.0)

```
1243 1 # GPT-5: Clean filter implementation
1244 2 class BioNeRFField(nn.Module):
1245 3     def __init__(self):
1246 4         self.f_delta = nn.Linear(hidden, hidden)
1247 5         self.f_color = nn.Linear(hidden, hidden)
1248 6         self.f_mem = nn.Linear(hidden * 2, hidden)
1249 7         self.f_mod = nn.Linear(hidden * 2, hidden)
1250 8         self.pre_mod = nn.Linear(hidden * 2, hidden)
1251 9     def forward(self, x, d):
1252 10        f_delta = torch.sigmoid(self.f_delta(h_delta))
1253 11        f_color = torch.sigmoid(self.f_color(h_color))
1254 12        gamma = torch.tanh(self.pre_mod(h_cat))
1255 13 # CORRECT: Good implementation with slight naming
1256 14 differences
```

Listing 70. C2-C6 Implementation: GPT-5 (Score: 0.8)

```
1261 1 # DeepSeek R1: Comprehensive filter implementation
1262 2 def __init__(self):
1263 3     self.W_f_delta = nn.Linear(hidden_dim, hidden_dim
1264 4         )
1265 5     self.W_f_color = nn.Linear(hidden_dim, hidden_dim
1266 6         )
1267 7     self.W_f_psi = nn.Linear(2 * hidden_dim,
1268 8         hidden_dim)
1269 9     self.W_f_mu = nn.Linear(2 * hidden_dim,
1270 10        hidden_dim)
1271 11 def forward(self):
1272 12    f_delta = torch.sigmoid(self.W_f_delta(h_delta))
1273 13    f_psi = torch.sigmoid(self.W_f_psi(h_concat))
1274 14 # ISSUE: Missing W_gamma layer definition
```

Listing 71. C2-C6 Implementation: DeepSeek R1 (Score: 0.7)

```
1277 1 # Paper2Code: Partial implementation
1278 2 def forward(self):
1279 3     f_delta = torch.sigmoid(self.linear_fd(h_delta))
1280 4     f_c = torch.sigmoid(self.linear_fc(h_c))
1281 5     # Concatenates but missing some filters
1282 6     h_cat = torch.cat([h_delta, h_c], dim=-1)
1283 7     gamma = torch.tanh(self.linear_gamma(h_cat))
1284 8 # MISSING: f_psi and f_mu filters not properly
1285 9 defined
```

Listing 72. C2-C6 Implementation: Paper2Code (Score: 0.4)

```
1288 1 # AutoP2C: Fatal implementation error
1289 2 def cognitive_filtering(self, features):
1290 3     h_delta, h_c = features
1291 4     # CRITICAL ERROR: Creates new layers every
1292 5     # forward pass!
1293 6     f_delta = torch.sigmoid(
1294 7         nn.Linear(256, 256)(h_delta)) # Random
1295 8     weights!
1296 9     f_c = torch.sigmoid(
1297 10        nn.Linear(256, 256)(h_c)) # Random
1298 11     weights!
1299 12     # No gradient tracking possible
```

Listing 73. C2-C6 Implementation: AutoP2C (Score: 0.0)

1302 **Analysis:** NERIFY perfectly implements all cognitive
 1303 filters. GPT-5 has good implementation with minor naming
 1304 variations. DeepSeek R1 misses gamma layer definition.
 1305 Paper2Code partially implements filters. AutoP2C has a
 1306 critical flaw - creating nn.Linear layers inside forward pass
 1307 results in random weights every iteration with no gradient
 1308 tracking.

1309 Component C7-C8: Memory Mechanism - All Baselines

1310 The memory mechanism involves modulation $\mu = f_\mu \otimes \gamma$
 1311 and recursive update $\Psi = \tanh(W_\Psi(\mu + f_\psi \otimes \Psi_{prev}) + b_\Psi)$
 1312 where the memory state persists across forward passes.

```
1313 1 # \nerify\ : Perfect stateful memory with detach
1314 2 def forward(self):
1315 3     prev_mem = self._ensure_memory(N, device, dtype)
1316 4     mu = f_mu * gamma # Element-wise modulation
1317 5     # Proper recursive update with previous memory
1318 6     Psi = self._tanh(self.W_Psi(mu + f_psi * prev_mem
1319 7         ))
1320 8     # Detach to prevent gradient explosion
1321 9     self._memory = Psi.detach()
1322 10 # CORRECT: Stateful memory with proper recursion
```

Listing 74. C7-C8 Implementation: NERIFY (Score: 1.0)

```
1326 1 # GPT-5: Separate memory module approach
1327 2 class MemoryModule(nn.Module):
1328 3     def forward(self, mu, fpsi):
1329 4         if self.memory is None:
1330 5             self.memory = torch.zeros_like(mu)
1331 6             updated = torch.tanh(
1332 7                 self.W_psi(mu + fpsi * self.memory) +
1333 8                 self.b_psi)
1334 9             self.memory = updated.detach()
1335 10 # CORRECT: Proper recursion, slight module complexity
```

Listing 75. C7-C8 Implementation: GPT-5 (Score: 0.8)

```
1339 1 # DeepSeek R1: Good attempt with persistence issue
1340 2 def forward(self, xyz, view_dir):
```

```

1342     mu = f_mu * gamma # Correct modulation
1343     if self.memory.shape[0] != batch_size:
1344         self.memory = self.memory.expand(batch_size,
1345             -1)
1346     memory_input = mu + (f_psi * self.memory)
1347     psi_new = torch.tanh(self.W_psi(memory_input))
1348     self.memory = psi_new
1349 # ISSUE: Memory expansion may not persist properly

```

Listing 76. C7-C8 Implementation: DeepSeek R1 (Score: 0.6)

```

1 # Paper2Code: Missing recursive component
2 def forward(self):
3     mu = f_mu * gamma
4     psi_init = torch.zeros_like(mu)
5     # ERROR: Always uses zeros, no memory persistence
6     psi = torch.tanh(self.linear_memory_update(
7         mu + (f_psi * psi_init))) # psi_init always
8         0!
8 # MISSING: No actual memory recursion

```

Listing 77. C7-C8 Implementation: Paper2Code (Score: 0.2)

```

1 # AutoP2C: Multiple critical issues
2 def memory_updating(self, filtered_features):
3     mu = f_mu * gamma
4     # CRITICAL ERROR: Creates new layer here
5     self.memory = torch.tanh(
6         nn.Linear(256, 256)(mu + f_psi * self.memory)
7     )
7 # Random weights, no learning possible

```

Listing 78. C7-C8 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY implements perfect stateful memory with detach to prevent gradient issues. GPT-5 uses a correct separate module approach. DeepSeek R1 has memory persistence issues. Paper2Code fails to maintain memory state. AutoP2C creates layers dynamically, making learning impossible.

Component C9-C10: Contextual Inference and Output Heads - All Baselines

The contextual inference creates embeddings $h'_\Delta = [\Psi \otimes f_\Delta, x]$ and $h'_c = [\Psi \otimes f_c, d]$, which are processed by density head M'_Δ and color head M'_c respectively.

```

1 # \nerify\ : Perfect contextual inference
2 def forward(self):
3     ctx_delta = Psi * f_delta # Element-wise
4     ctx_c = Psi * f_c
5     sigma_in = torch.cat([ctx_delta, pe], dim=-1)
6     sigma_raw = self.density_head(sigma_in)
7     density = trunc_exp(sigma_raw).view(*positions.
8         shape[:-1], 1)
8     rgb_in = torch.cat([ctx_c, de], dim=-1)
9     rgb = self.color_head(rgb_in).view(*positions.
10        shape[:-1], 3)
10 # CORRECT: Exact paper implementation

```

Listing 79. C9-C10 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Good implementation with minor differences
2 def forward(self):
3     hpd = torch.cat([psi * f_delta, x], dim=-1)
4     hpc = torch.cat([psi * f_color, d], dim=-1)
5     density = self.Mp_delta(hpd)
6     color = torch.sigmoid(self.Mp_color(hpc))
7 # CORRECT: Proper concatenation and heads

```

Listing 80. C9-C10 Implementation: GPT-5 (Score: 0.8)

```

1 # DeepSeek R1: Correct approach, minor issues
2 def forward(self):
3     h_delta_prime = torch.cat([psi_new * f_delta, xyz
4         ], dim=-1)
4     h_color_prime = torch.cat([psi_new * f_color,
5         view_dir], dim=-1)
5     density = self.density_head(h_delta_prime)
6     color = torch.sigmoid(self.color_head(
7         h_color_prime))
7 # ISSUE: Missing proper view reshaping

```

Listing 81. C9-C10 Implementation: DeepSeek R1 (Score: 0.7)

```

1 # Paper2Code: Simplified inference
2 def forward(self):
3     h_delta_prime = torch.cat((psi * f_delta, 1), dim
4         =-1)
4     density = self.density_head(h_delta_prime)
5     color = self.color_head(h_c_prime)
6 # ERROR: h_c_prime not properly defined

```

Listing 82. C9-C10 Implementation: Paper2Code (Score: 0.3)

```

1 # AutoP2C: Structural attempt, poor execution
2 def forward(self):
3     density = self.heads['density'](h_delta_prime)
4     color = self.heads['color'](h_c_prime)
5     return density, color
6 # ERROR: h_delta_prime, h_c_prime undefined
7 # MISSING: No contextual concatenation

```

Listing 83. C9-C10 Implementation: AutoP2C (Score: 0.1)

Analysis: NERFIFY perfectly implements contextual inference with proper concatenation and reshaping. GPT-5 has good implementation with minor variations. DeepSeek R1 misses view reshaping. Paper2Code has undefined variables. AutoP2C attempts structure but lacks proper implementation.

3.1.5. Scoring Analysis

Table 28. Revised BioNeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	Weighted Avg LLM
NERFIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.00
GPT-5	0.9	0.8	0.8	0.8	0.8	0.8	0.8	0.8	0.8	0.8	0.82
DeepSeek R1	0.8	0.7	0.7	0.7	0.7	0.7	0.6	0.7	0.7	0.7	0.70
Paper2Code	0.3	0.4	0.4	0.4	0.4	0.3	0.2	0.3	0.3	0.3	0.34
AutoP2C	0.1	0.0	0.0	0.0	0.0	0.1	0.0	0.1	0.1	0.1	0.04

3.1.6. Why Baselines Fail Despite Component Scores

GPT-5 (Score: 0.82 components, 0% trainable)

- Strengths:** Near-perfect component implementation, clean modular code structure, proper memory handling

- Fatal Issues:**

- No Nerfstudio integration whatsoever
- Missing training pipeline and data loading infrastructure
- Standalone module without framework compatibility

- Result:** Excellent algorithm but requires 3-5 hours of integration work to be trainable

DeepSeek R1 (Score: 0.75 components, 0% trainable)

- Strengths:** Comprehensive implementation with all components present, proper encoding

1406
1407
1408
1409
1410
1411
1412
1413
1414
1415
1416

1418
1419
1420
1421
1422
1423
1424
1425

1427
1428
1429
1430
1431
1432
1433
1434

1436
1437
1438
1439
1440
1441
1442

1443

1443

1444

1445

1446

1447

1448

1449

1450

1451

1452

1453

1454

1455

- 1456 • Fatal Issues:**
- Memory buffer expansion doesn't persist between batches
 - No Nerfstudio plugin structure
 - Missing critical training loop integration
- 1457 • Result:** Memory persistence issues cause degradation to vanilla NeRF behavior
- 1458 **Paper2Code (Score: 0.35 components, 0% trainable)****
- 1459 • Strengths:** Extensive 1200+ line codebase with complete dataset loading infrastructure
- 1460 • Fatal Issues:**
- Implements generic NeRF with BioNeRF naming conventions
 - Memory mechanism lacks recursion - always uses zeros
 - Missing true parallel pathways for M_{delta} and M_c
 - No Nerfstudio integration
- 1461 • Result:** Code runs but implements wrong algorithm, achieving vanilla NeRF performance
- 1462 **AutoP2C (Score: 0.15 components, 0% trainable)****
- 1463 • Strengths:** Attempts all structural components with correct high-level organization
- 1464 • Fatal Issues:**
- Creates nn.Linear layers in forward pass resulting in random weights every iteration
 - Positional encoding defined but not applied to inputs
 - No gradient tracking possible with dynamic layer creation
 - Fundamental PyTorch misunderstanding
- 1465 • Result:** Catastrophic implementation errors prevent any learning from occurring

3.1.7. Hyperparameter Fidelity

Table 29. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5	R1	P2C	AutoP2C
Learning rate	5e-4	✓	✓	✓	✓	✗
Hidden dimensions	256	✓	✓	✓	✓	✓
Memory tensor size	8192	✓	~	~	✗	✓
Training iterations	400k	✓	✗	✗	✓	✗
Adam optimizer	Yes	✓	✓	✓	✓	✓
Positional frequencies	10	✓	~	✓	✓	✓
Batch size	8192	✓	~	✗	✓	✗
W Score	–	1.00	0.57	0.57	0.71	0.43

3.1.8. Conclusion

All baselines attempt to implement BioNeRF with varying levels of sophistication and understanding. GPT-5 and DeepSeek R1 achieve high component scores with nearly correct algorithmic implementations but completely lack the Nerfstudio integration necessary for training. Paper2Code produces an extensive codebase but fundamentally misunderstands the algorithm, implementing a simplified NeRF variant that misses the critical recursive memory mechanism. AutoP2C demonstrates the most severe implementation er-

rors, creating neural network layers dynamically within the forward pass, which results in random weights every iteration and makes gradient-based learning impossible. Only NERFIFY produces immediately trainable code by combining perfect component implementation with proper framework integration as a complete Nerfstudio plugin, achieving a 100% trainable rate while all baselines remain at 0% trainable despite their varying component scores. ““

1497
1498
1499
1500
1501
1502
1503
1504

3.2. Mip-NeRF: A Multiscale Representation for Anti-Aliasing Neural Radiance Fields

1505
1506

3.2.1. Paper Overview

1507
1508
1509
1510
1511
1512
1513
1514
1515
1516

Mip-NeRF [2] addresses the critical aliasing artifacts in Neural Radiance Fields by replacing point sampling with cone tracing and standard positional encoding with integrated positional encoding (IPE). The paper's key innovation is approximating conical frustums as 3D Gaussians and computing the expected positional encoding analytically, enabling anti-aliased rendering across multiple scales with a single multiscale MLP that reduces parameters by 50% compared to NeRF.

3.2.2. Implementation Overview

1517

Table 30. Mip-NeRF Implementation Summary Across All Baselines

Aspect	NERFIFY	GPT-5	DeepSeek R1	Paper2Code	AutoP2C
Lines of Code	563	412	287	1842	895
File Organization	Plugin	Single	Single	Multi-file	Multi-file
Conical Frustum	✓	✗	✓	✓	Simplified
IPE Implementation	✓	✗	✓	✓	Attempted
Single MLP	✓	✓	✓	✓	✗
Trainable	✓	✗	✗	✗	✗

Note: NERFIFY always produces trainable code as a complete Nerfstudio plugin

3.2.3. Novel Components

1518

Table 31. Novel Components in Mip-NeRF with Importance Weights

ID	Component	Weight w_i
C1	Conical Frustum Parameters (Eq. 7-8, stable formulation)	0.20
C2	Gaussian Mean Transform to World Coordinates	0.10
C3	Gaussian Covariance Transform (diagonal approximation)	0.15
C4	Integrated Positional Encoding (IPE)	0.25
C5	Single Multiscale MLP Architecture	0.15
C6	Area-Weighted Loss for Multiscale Training	0.15

3.2.4. Quantitative Metrics

1519

3.2.5. Component-by-Component Analysis

1520

Component C1: Conical Frustum Parameters - All Baselines

1521
1522
1523
1524

The paper derives stable reparameterized formulas (Appendix A) for computing Gaussian parameters of a conical

Table 32. Mip-NeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	1.00	0.00	0.00	1.00	1.00
GPT-5	0.50	0.30	0.20	0.60	0.58
DeepSeek R1	0.67	0.17	0.16	0.67	0.58
Paper2Code	0.83	0.17	0.00	0.83	0.85
AutoP2C	0.17	0.17	0.66	0.25	0.20

frustum. Given segment endpoints t_0, t_1 and cone radius r , the formulas use midpoint $t_\mu = (t_0 + t_1)/2$ and half-width $t_\delta = (t_1 - t_0)/2$ to compute mean μ_t , axial variance σ_t^2 , and radial variance σ_r^2 .

```

1 # \nerfify\ : Exact stable formulation from paper
2 t_mu = (t0 + t1) / 2.0
3 t_delta = (t1 - t0) / 2.0
4 mu_t = t_mu + (2.0 * t_mu * t_delta ** 2) / (3.0 *
5     t_mu ** 2 + t_delta ** 2)
6 sigma_t_sq = (t_delta ** 2) / 3.0 - \
7     (4.0 * t_delta ** 4 * (12.0 * t_mu ** 2 - t_delta
8     ** 2)) / \
9     (15.0 * (3.0 * t_mu ** 2 + t_delta ** 2) ** 2)
10 sigma_r_sq = (dot_r ** 2) * (t_mu ** 2 / 4.0 + 5.0 *
11     t_delta ** 2 / 12.0 - \
12     4.0 * t_delta ** 4 / (15.0 * (3.0 * t_mu ** 2 +
13     t_delta ** 2)))
14 # CORRECT: Stable reparameterized formulas from
15 Appendix A

```

Listing 84. C1 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Completely wrong formulas
2 def conical_frustum_to_gaussian(t0, t1, radius):
3     t_mean = (3*(t1**4 - t0**4))/(4*(t1**3 - t0**3))
4     t_var = (3*(t1**5 - t0**5))/(5*(t1**3 - t0**3)) -
5         t_mean**2
6     r_var = radius**2 * (3*(t1**5 - t0**5))/(20*(t1
7         **3 - t0**3))
8     return t_mean, t_var, r_var
9 # CRITICAL ERROR: Different derivation, numerically
10 unstable

```

Listing 85. C1 Implementation: GPT-5 (Score: 0.0)

```

1 # DeepSeek R1: Correct stable formulation
2 mu_t = t_mu + (2 * t_mu * t_delta**2) / (3 * t_mu**2
3     + t_delta**2)
4 sigma_t2 = (t_delta**2 / 3) - (4 * t_delta**4 * (12 *
5     t_mu**2 - t_delta**2)) / \
6     (15 * (3 * t_mu**2 + t_delta**2)**2)
5 sigma_r2 = r**2 * (t_mu**2 / 4 + 5 * t_delta**2 / 12
7     - \
8     4 * t_delta**4 / (15 * (3 * t_mu**2 + t_delta**2)))
7 # CORRECT: Matches paper formulation

```

Listing 86. C1 Implementation: DeepSeek R1 (Score: 1.0)

```

1 # Paper2Code: Correct with JAX
2 denominator = 3 * (z_center ** 2) + z_delta ** 2 + 1e
    -8
3 mu_z = z_center + (2.0 * z_center * (z_delta ** 2)) /
    denominator
4 sigma_z_sq = (z_delta ** 2) / 3.0 - \
5     (4.0 * (z_delta ** 4) * (12 * z_center ** 2 -
    z_delta ** 2)) / \
6     (15.0 * (denominator ** 2) + 1e-8)
7 # CORRECT: Stable formulation with numerical safety

```

Listing 87. C1 Implementation: Paper2Code (Score: 1.0)

```

1 # AutoP2C: Missing frustum computation
2 # No conical frustum parameter computation found
3 # Uses stratified sampling without cone tracing
4 # MISSING: Core component entirely absent

```

Listing 88. C1 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY, DeepSeek R1, and Paper2Code correctly implement the stable reparameterized formulas. GPT-5 uses completely different formulas that are both mathematically incorrect and numerically unstable. AutoP2C entirely omits conical frustum computation.

Component C2-C3: Gaussian Mean and Covariance Transform - All Baselines

The paper transforms frustum parameters to world coordinates: $\mu = \mathbf{o} + \mu_t \mathbf{d}$ for the mean and $\Sigma = \sigma_t^2(\mathbf{d}\mathbf{d}^\top) + \sigma_r^2(\mathbf{I} - \mathbf{d}\mathbf{d}^\top/\|\mathbf{d}\|_2^2)$ for covariance.

```

1 # \nerfify\ : Efficient diagonal approximation
2 means = origins + mu_t[..., None] * directions # +
3     Mean transform
4 d_outer_diag = directions ** 2
5 d_norm_sq = (directions ** 2).sum(dim=-1, keepdim=
6     True)
5 covs_diag = sigma_t_sq[..., None] * d_outer_diag +
6     sigma_r_sq[..., None] * (1.0 - d_outer_diag /
7     d_norm_sq)
7 # CORRECT: Diagonal covariance for efficiency

```

Listing 89. C2-C3 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Full covariance matrix (inefficient)
2 mu = rays_o[..., None, :] + mu_z[..., None] * rays_d
    [..., None, :]
3 covs = (rays_d[..., None] * rays_d[..., None, :, :] *
    t_var[..., None, None]
4     + jnp.eye(3) * r_var[..., None, None])
5 # ISSUE: Computes full 3x3 matrix, inefficient but
    mathematically correct

```

Listing 90. C2-C3 Implementation: GPT-5 (Score: 0.5)

```

1 # Paper2Code: Correct diagonal implementation
2 mu = rays_o_exp + jnp.expand_dims(mu_z, axis=-1) *
    rays_d_exp
3 r_squared = rays_d ** 2
4 cov_diag = jnp.expand_dims(sigma_z_sq, axis=-1) *
    r_squared_exp +
5     jnp.expand_dims(sigma_r_sq, axis=-1) * (1.0 -
    r_squared_exp)
6 # CORRECT: Efficient diagonal form

```

Listing 91. C2-C3 Implementation: Paper2Code (Score: 1.0)

Analysis: Most baselines correctly implement the transforms. GPT-5 computes the full covariance matrix instead of just the diagonal, which is mathematically correct but computationally inefficient.

Component C4: Integrated Positional Encoding (IPE) - All Baselines

The core innovation computes expected positional encoding over a Gaussian: $\mathbb{E}[\sin(x)] = \sin(\mu) \exp(-\sigma^2/2)$ and $\mathbb{E}[\cos(x)] = \cos(\mu) \exp(-\sigma^2/2)$ for frequency bands $2^0, 2^1, \dots, 2^{L-1}$.

```

1642 1 # \nerify\ : Perfect IPE implementation
1643 2 scales = 2.0 ** torch.arange(self.
1644     position_encoding_max_degree)
1645 3 means_scaled = means[..., :, None] * scales
1646 4 variance_scales = (scales ** 2)
1647 5 covs_scaled = covs_diag[..., :, None] *
1648     variance_scales
1649 6 damping = torch.exp(-0.5 * covs_scaled)
1650 7 sin_features = torch.sin(means_scaled) * damping
1651 8 cos_features = torch.cos(means_scaled) * damping
1652 9 features = torch.cat([sin_features, cos_features],
1653     dim=-1)
1654 10 # CORRECT: Exact formulation from paper
1655

```

Listing 92. C4 Implementation: NERIFY (Score: 1.0)

```

1657 1 # GPT-5: Multiple critical errors
1658 2 freqs = 2.0 ** jnp.arange(max_freq_log2)
1659 3 scales = jnp.pi * freqs[None, None, :] # ERROR:
1660     Extra pi!
1661 4 diag = jnp.stack([cov[..., i, i] for i in range(3)],
1662     axis=-1)
1663 5 var = diag[..., None, :] * scales[..., None] ** 2
1664 6 # CRITICAL ERROR: Wrong frequency scaling, dimension
1665     mismatches
1666

```

Listing 93. C4 Implementation: GPT-5 (Score: 0.2)

```

1667 1 # Paper2Code: Correct JAX implementation
1668 2 freq_bands = 2.0 ** jnp.arange(L, dtype=jnp.float32)
1669 3 args = mu_exp * freq_bands
1670 4 damping = jnp.exp(-0.5 * (cov_exp * (freq_bands ** 2
1671     )))
1672 5 sin_features = jnp.sin(args) * damping
1673 6 cos_features = jnp.cos(args) * damping
1674 7 # CORRECT: Proper IPE computation
1675

```

Listing 94. C4 Implementation: Paper2Code (Score: 1.0)

Analysis: NERIFY, DeepSeek R1, and Paper2Code correctly implement IPE. GPT-5 has multiple errors including extra π multiplication and dimension mismatches. AutoP2C attempts IPE but with incorrect structure.

Component C5: Single MLP Architecture - All Baselines

Mip-NeRF uses a single multiscale MLP instead of separate coarse/fine networks, reducing parameters by 50

```

1686 1 # \nerify\ : Single MLP with hierarchical sampling
1687     in Nerfstudio
1688 2 class MipNeRFModel(Model):
1689 3     def populate_modules(self):
1690         self.field = MipNeRFField() # Single field
1691         self.sampler_uniform = UniformSampler(
1692             num_samples=128)
1693         self.sampler_pdf = PDFSampler(num_samples
1694             =128, include_original=False)
1695     7 # CORRECT: Single MLP queried hierarchically
1696

```

Listing 95. C5 Implementation: NERIFY (Score: 1.0)

```

1698 1 # AutoP2C: Uses two separate MLPs
1699 2 coarse_out = self.model.apply({"params": self.state.
1700     params}, ray_batch_coarse)
1701 3 # ... hierarchical sampling ...
1702 4 fine_out = state.apply_fn({"params": params},
1703     ray_batch_fine)
1704 5 # ERROR: Separate coarse/fine networks instead of
1705     single MLP
1706

```

Listing 96. C5 Implementation: AutoP2C (Score: 0.0)

Analysis: Most baselines correctly implement a single MLP. AutoP2C incorrectly uses hierarchical two-network architecture from original NeRF.

Component C6: Area-Weighted Loss - All Baselines

For multiscale training, pixel losses are weighted by their footprint area.

```

1708 1 # \nerify\ : Full area-weighted loss support
1709 2 if self.config.area_loss_mult > 0:
1710 3     pixel_area = ray_bundle.pixel_area
1711 4     loss = loss * pixel_area.squeeze() * self.config.
1712     area_loss_mult
1713 5 # CORRECT: Proper multiscale weighting
1714

```

Listing 97. C6 Implementation: NERIFY (Score: 1.0)

```

1714 1 # Paper2Code: Partial implementation
1715 2 def area_weighted_loss(pred_rgb, target_rgb,
1716     pixel_area, lambda=0.1):
1717 3     diff = (pred_rgb - target_rgb)**2
1718 4     return lambda * jnp.mean(diff * pixel_area[...,
1719     None])
1720 5 # ISSUE: Defined but not integrated into training
1721     loop
1722

```

Listing 98. C6 Implementation: Paper2Code (Score: 0.5)

Analysis: Only NERIFY fully implements area-weighted loss. Paper2Code defines it but doesn't integrate it. Others omit it entirely.

3.2.6. Scoring Analysis

Table 33. Mip-NeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	Weighted Avg _{LLM}
NERIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.00
GPT-5	0.0	1.0	0.5	0.2	1.0	0.0	0.58
DeepSeek R1	1.0	1.0	1.0	1.0	0.5	0.0	0.58
Paper2Code	1.0	1.0	1.0	1.0	1.0	0.5	0.85
AutoP2C	0.0	0.5	0.5	0.3	0.0	0.0	0.20

3.2.7. Why Baselines Fail Despite Component Scores

GPT-5 (Score: 0.58 components, 0% trainable)

- Strengths:** Implements single MLP architecture and basic transforms
 - Fatal Issues:**
 - Wrong conical frustum formulas cause NaN gradients
 - IPE implementation has critical errors (extra π , dimension mismatches)
 - JAX code incompatible with Nerfstudio benchmark
 - Result:** Non-trainable due to fundamental mathematical errors
- DeepSeek R1 (Score: 0.58 components, 0% trainable)**
- Strengths:** Correct mathematical formulations for frustum and IPE
 - Fatal Issues:**
 - No Nerfstudio integration (standalone PyTorch)
 - Missing critical training infrastructure

- 1753 – Simplified hierarchical sampling without proposal net-
 1754 work
- **Result:** Cannot be executed in benchmark environment
 - Paper2Code (Score: 0.85 components, 0% trainable)**
 - **Strengths:** Most accurate mathematical implementation in JAX/Flax
 - **Fatal Issues:**
 - Wrong framework (JAX instead of PyTorch/Nerfstudio)
 - Incompatible API prevents benchmark evaluation
 - Requires complete infrastructure rewrite
 - **Result:** Excellent standalone code but non-trainable in benchmark
 - AutoP2C (Score: 0.20 components, 0% trainable)**
 - **Strengths:** Attempts basic NeRF structure
 - **Fatal Issues:**
 - Missing conical frustum computation entirely
 - Uses two MLPs instead of single architecture
 - Incomplete Nerfstudio integration
 - Many placeholder functions
 - **Result:** Fundamentally incomplete implementation

3.2.8. Hyperparameter Fidelity

Table 34. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5	R1	P2C	AutoP2C
PE frequencies (L)	16	✓	✓	✓	✓	✗
Direction encoding	4	✓	✗	✓	✗	✗
Hidden dimension	256	✓	✓	✓	✓	✓
Learning rate	$5e^{-4}$	✓	✓	✓	✓	✗
Batch size	4096	✓	✗	✗	✓	✓
W Score	–	1.00	0.60	0.67	0.83	0.25

3.2.9. Conclusion

All baselines attempt to implement Mip-NeRF with varying sophistication. Paper2Code achieves the highest component accuracy (0.85) with correct mathematical implementations but uses JAX/Flax incompatible with Nerfstudio. GPT-5 and DeepSeek R1 both score 0.58 but fail differently - GPT-5 has fundamental mathematical errors while R1 lacks integration. AutoP2C scores lowest (0.20) missing core components entirely. Only NERFIFY produces immediately trainable code as a complete Nerfstudio plugin, achieving 100% trainable rate while all baselines achieve 0%.

3.3. PyNeRF: Pyramidal Neural Radiance Fields

3.3.1. Paper Overview

PyNeRF introduces a pyramidal multiscale architecture for neural radiance fields that addresses aliasing artifacts through dynamic resolution selection. The method maintains $L = 8$ pyramid levels with geometrically increasing resolutions $N_l = N_0 \times s^l$ (where $N_0 = 16$ and $s = 2$), selecting the appropriate level for each sample based on its projected pixel area $P(x)$ using the mapping function

$M(P) = \log_s(P/N_0)$. The key innovation lies in per-sample level selection with linear interpolation between adjacent pyramid levels to achieve smooth transitions.

1794
1795
1796

3.3.2. Implementation Overview

1797

Table 35. PyNeRF Implementation Summary Across All Baselines

Aspect	NERFIFY	GPT-5	DeepSeek R1	Paper2Code	AutoP2C
Lines of Code	420	185	380	1250	450
File Organization	Plugin	Modular	Single	Multi-file	Multi-file
Pyramid Architecture	✓	✓	✓	✓	✗
Level Selection	✓	✓	Partial	✓	✗
Interpolation	✓	✓	Partial	✓	✗
Projected Area	✓	~	✗	✗	✗
Nerfstudio Integration	✓	✗	✗	✗	✗
Trainable	✓	✗	✗	✗	✗

Note: NERFIFY always produces trainable code as a complete Nerfstudio plugin

3.3.3. Novel Components

1798

Table 36. Novel Components in PyNeRF with Importance Weights

ID	Component	Weight w_i
C1	Pyramidal architecture with L levels	0.15
C2	Per-sample level selection based on projected area	0.20
C3	Linear interpolation between levels	0.15
C4	Shared multiresolution features	0.10
C5	Per-level tiny MLPs (64/128 hidden)	0.10
C6	Projected area computation $P(x) = pixel_area \times t^2$	0.10
C7	Adaptive level supervision	0.05
C8	Modern backbone integration (iNGP/TensoRF)	0.05
C9	Multiple interpolation modes (linear/Laplacian)	0.05
C10	Specific hyperparameters (8 levels, s=2, N0=16)	0.05

3.3.4. Quantitative Metrics

1799

Table 37. PyNeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	0.95	0.05	0.00	1.00	0.97
GPT-5	0.40	0.30	0.30	0.80	0.52
DeepSeek R1	0.30	0.60	0.10	0.70	0.68
Paper2Code	0.50	0.30	0.20	0.60	0.58
AutoP2C	0.00	0.10	0.90	0.00	0.03

3.3.5. Component-by-Component Analysis

1800

Component C1: Pyramidal Architecture with L Levels - All Baselines The paper specifies creating $L = 8$ pyramid levels where each level l has resolution $N_l = N_0 \times s^l$ with separate processing heads for each level.

1801
1802
1803
1804

```
1 # \nerfify\ : ModuleList of per-level heads
2 self.density_heads = nn.ModuleList([
3     MLP(in_dim=self.feature_dim, num_layers=1,
4          layer_width=density_hidden_dim, out_dim=1)
5     for _ in range(self.num_pyramid_levels)
6 ])
```

1805
1806
1807
1808
1809
1810
1811

```

1812    7 self.color_heads = nn.ModuleList([
1813        8     MLP(in_dim=color_in_dim, num_layers=2,
1814        9         layer_width=color_hidden_dim, out_dim=3)
1815        10    for _ in range(self.num_pyramid_levels)
1816    11  ])
1817 12 # CORRECT: Exactly L separate heads as specified

```

Listing 99. C1 Implementation: NERIFY (Score: 1.0)

```

1819 1 # GPT-5: List comprehension for levels
1820 2 models = [PyNeRFLevel(enc, cfg) for _ in range(cfg.
1821      n_levels)]
1822 3 pipeline = PyNeRFPipeline(models, cfg).cuda()
1823 4 # CORRECT: Creates pyramid structure but lacks
1824      integration

```

Listing 100. C1 Implementation: GPT-5 (Score: 1.0)

```

1827 1 # DeepSeek R1: ModuleList with incomplete
1828      implementation
1829 2 self.levels = nn.ModuleList()
1830 3 for i in range(num_levels):
1831 4     level = PyNeRFLevel(base_nerf_config, level_idx=i
1832      ,
1833      use_shared_features=
1834      use_shared_features)
1835 6 self.levels.append(level)
1836 7 # ISSUE: Structure correct but PyNeRFLevel incomplete

```

Listing 101. C1 Implementation: DeepSeek R1 (Score: 1.0)

```

1839 1 # Paper2Code: Detailed pyramid construction
1840 2 self.heads = nn.ModuleList()
1841 3 for level in range(self.pyramid_levels):
1842 4     effective_resolution = self.base_resolution * (
1843      self.scale_factor ** level)
1844 5     head = PyNeRFHead(input_dim=num_features,
1845      effective_resolution=
1846      effective_resolution)
1847 7 self.heads.append(head)
1848 8 # CORRECT: Proper pyramid with effective resolutions

```

Listing 102. C1 Implementation: Paper2Code (Score: 1.0)

```

1851 1 # AutoP2C: Single MLP for all samples
1852 2 self.density_mlp = nn.Sequential(
1853      nn.Linear(2 * 10 * 3, 64), nn.ReLU(), nn.Linear
1854          (64, 1))
1855 4 self.color_mlp = nn.Sequential(
1856      nn.Linear(2 * 10 * 6 + 64, 128), nn.ReLU(),
1857      nn.Linear(128, 128), nn.ReLU(), nn.Linear(128, 3)
1858      )
1859 7 # CRITICAL ERROR: No pyramid structure at all

```

Listing 103. C1 Implementation: AutoP2C (Score: 0.0)

Analysis: All methods except AutoP2C correctly implement the pyramid architecture. AutoP2C completely misses this core concept.

Component C2: Per-Sample Level Selection - All Baselines The paper requires computing $M(P) = \log_s(P/N_0)$ and selecting level $l = \lceil M(P) \rceil$ with interpolation weight $w = l - M(P)$.

```

1869 1 # \nerify\ : Exact paper formulation
1870 2 t_mid = (ray_samples.frustums.starts + ray_samples.
1871      frustums.ends) * 0.5
1872 3 pixel_area = ray_samples.frustums.pixel_area
1873 4 proj_area = pixel_area * torch.clamp(t_mid, min=0.0)
1874      ** 2 + 1e-12

```

```

1876 5 s = torch.tensor(self.pyramid_scale_factor, device=
1877      proj_area.device)
1878 6 M = torch.log(proj_area / max(self.base_proj_area, 1e
1879          -12)) / torch.log(s)
1880 7 l = torch.clamp(torch.ceil(M), min=0.0, max=float(
1881          self.num_pyramid_levels - 1))
1882 8 w = l - M # Interpolation weight
1883 9 # CORRECT: Matches paper equations exactly

```

Listing 104. C2 Implementation: NERIFY (Score: 1.0)

```

1885 1 # GPT-5: Clean implementation
1886 2 def compute_level(Px, N0, s):
1887 3     M = torch.log2(Px / N0) / torch.log2(torch.tensor
1888      (s))
1889 4     l = torch.clamp(torch.ceil(M), 0, None)
1890 5     w = l - M
1891 6     return l.long(), w
1892 7 # CORRECT: Proper level selection logic

```

Listing 105. C2 Implementation: GPT-5 (Score: 1.0)

```

1895 1 # DeepSeek R1: Incomplete with heuristics
1896 2 def compute_sample_scale(self, rays, sample_positions
1897      ):
1898 3     sample_distances = torch.norm(sample_positions -
1899      rays['origins'].unsqueeze(1), dim=-1)
1900 4     pixel_areas = sample_distances ** 2 * 0.01 # heuristic scaling
1901 5     return pixel_areas.unsqueeze(-1).expand_as(
1902      sample_distances)
1903 6     M = torch.log(scales / self.base_resolution) / np.log
1904      (self.scale_factor)
1905 7     levels = torch.clamp(torch.ceil(M), 0, self.
1906      num_levels - 1).long()
1907 8 # ERROR: Heuristic area computation, not paper
1908      formula

```

Listing 106. C2 Implementation: DeepSeek R1 (Score: 0.6)

```

1912 1 # Paper2Code: Correct logic, wrong area formula
1913 2 def compute_integration_area(self, depths):
1914 3     integration_area = (depths / self.focal) ** 2
1915 4     return integration_area
1916 5 M_value = torch.log(integration_area / (base_res +
1917      eps)) / math.log(scale + eps)
1918 6 l_target = torch.clamp(torch.ceil(M_value), min=0,
1919      max=self.pyramid_levels - 1)
1920 7 w = l_target - M_value
1921 8 # CRITICAL ERROR: Wrong formula - should be
1922      pixel_area * t^2

```

Listing 107. C2 Implementation: Paper2Code (Score: 0.8)

```

1925 1 # AutoP2C: Placeholder that doesn't work
1926 2 def compute_projected_area(self, x):
1927 3     # Placeholder for actual projected area
1928      computation
1929 4     return torch.ones_like(x[..., 0])
1930 5 # CRITICAL ERROR: Always returns 1.0, no level
1931      selection

```

Listing 108. C2 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERIFY and GPT-5 implement correct level selection. Paper2Code's critical error in area computation undermines the entire method.

Component C3: Linear Interpolation Between Levels - All Baselines The paper specifies linear interpolation: $\sigma, c = wf_l(\mathbf{x}, \mathbf{d}) + (1 - w)f_{l-1}(\mathbf{x}, \mathbf{d})$.

1934
1935
1936
1937
1938
1939

1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1950
1951
1952
1953
1954

```
1 # \nerify\ : Perfect interpolation implementation
2 rgb_main = torch.gather(rgb_all, 0, idx_main).squeeze
   (0)
3 rgb_prev = torch.gather(rgb_all, 0, idx_prev).squeeze
   (0)
4 if self.interp_mode == "interpolate":
5     rgb = w * rgb_main + (1.0 - w) * rgb_prev
6 else: # Laplacian mode
7     rgb_cum = torch.cumsum(rgb_all, dim=0)
8     rgb = torch.gather(rgb_cum, 0, idx_main).squeeze
   (0)
9 # CORRECT: Exact Equation 5 with multiple modes
```

Listing 109. C3 Implementation: NERIFY (Score: 1.0)

1955
1956
1957
1958
1959
1960
1961
1962
1963
1964
1965

```
1 # GPT-5: Clean interpolation
2 def forward(self, x, d, deltas, Px):
3     1, w = compute_level(Px, self.cfg.base_resolution
   , self.cfg.scale_factor)
4     c1, s1 = self.models[1-1](x, d)
5     c2, s2 = self.models[1](x, d)
6     c = w.unsqueeze(-1) * c2 + (1 - w).unsqueeze(-1)
   * c1
7     sigma = w * s2 + (1 - w) * s1
8     return volume_render(c, sigma, deltas)
9 # CORRECT: Proper linear interpolation
```

Listing 110. C3 Implementation: GPT-5 (Score: 1.0)

1967
1968
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1980
1981
1982

```
1 # DeepSeek R1: Partial implementation
2 if self.interpolation_method == "linear" and
   level_idx > 0:
3     prev_output = self.levels[level_idx - 1](
   level_samples, level_dirs)
4     level_weights = weights[mask].unsqueeze(-1)
5     level_output['density'] = (level_weights *
   level_output['density']) +
   (1 - level_weights) *
   prev_output['density'])
7     # Note: Color interpolation similar but omitted
   in provided code
8 # ISSUE: Incomplete - missing color interpolation
   details
```

Listing 111. C3 Implementation: DeepSeek R1 (Score: 0.8)

1983
1984
1985
1986
1987
1988
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998

```
1 # Paper2Code: Correct interpolation logic
2 if level_int == 0:
3     density, color = self.heads[0](x_level, d_level)
4 else:
5     density_upper, color_upper = self.heads[level_int
   ](x_level, d_level)
6     density_lower, color_lower = self.heads[level_int
   - 1](x_level, d_level)
7     density = w_level * density_upper + (1 - w_level)
   * density_lower
8     color = w_level * color_upper + (1 - w_level) *
   color_lower
9 # CORRECT: Proper interpolation between adjacent
   levels
```

Listing 112. C3 Implementation: Paper2Code (Score: 1.0)

1999
2000
2001
2002
2003
2004
2005
2006
2007
2008

```
1 # AutoP2C: Nonsensical interpolation
2 if self.multiscale_sampling:
3     # Linearly interpolate between the outputs of two
   closest hierarchy levels
4     if self.multiscale_sampling:
5         density = w * density + (1 - w) * density
6         color = w * color + (1 - w) * color
7 # CRITICAL ERROR: Interpolates with itself - no
   effect
```

Listing 113. C3 Implementation: AutoP2C (Score: 0.1)

Analysis: NERIFY, GPT-5, and Paper2Code correctly implement interpolation. AutoP2C's self-interpolation is a fundamental misunderstanding.

Component C4: Shared Multiresolution Features -

All Baselines The paper suggests reusing multi-resolution features across levels when the backbone already has a scale hierarchy.

```
1 # \nerify\ : Shared hash encoding across all levels
2 self.hash_map = HashEncoding(
3     num_levels=num_hash_levels, min_res=min_res,
   max_res=max_res,
4     log2_hashmap_size=log2_hashmap_size,
   features_per_level=features_per_level)
5 # All pyramid levels query the same shared encoding:
6 feats = self.hash_map(pos_flat) # Used by all
   density/color heads
7 # CORRECT: Single shared encoding as specified
```

Listing 114. C4 Implementation: NERIFY (Score: 1.0)

```
1 # GPT-5: Same encoder reference for all levels
2 enc = HashEncoder(cfg.n_levels, cfg.
   features_per_level,
3     cfg.base_resolution, cfg.
   scale_factor, cfg.hash_size)
4 models = [PyNeRFLevel(enc, cfg) for _ in range(cfg.
   n_levels)]
5 # ERROR: All levels share same encoder instance -
   problematic
```

Listing 115. C4 Implementation: GPT-5 (Score: 0.2)

```
1 # DeepSeek R1: Configurable shared features
2 def __init__(self, use_shared_features: bool = True):
3     if not use_shared_features:
4         self.feature_grid = self._create_feature_grid
   (base_config, level_idx)
5     else:
6         self.feature_grid = None # Will use shared
   features from base NeRF
7 # ISSUE: Structure present but implementation
   incomplete
```

Listing 116. C4 Implementation: DeepSeek R1 (Score: 0.6)

```
1 # Paper2Code: Single shared grid
2 self.shared_grid = True
3 self.feature_grid = FeatureGrid(resolution=self.
   base_resolution,
4     num_features=
5     hash_table_size=
6     hash_table_size)
6 # Each head references the same grid
7 head = PyNeRFHead(feature_grid=self.feature_grid)
8 # ISSUE: Simple hash table, not true multi-resolution
```

Listing 117. C4 Implementation: Paper2Code (Score: 0.6)

```
1 # AutoP2C: No feature sharing concept
2 self.positional_encoding = positional_encoding
3 x_encoded = self.apply_positional_encoding(x)
4 # MISSING: No shared features, uses positional
   encoding only
```

Listing 118. C4 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERIFY implements proper shared multi-resolution features. Others either misunderstand or omit this optimization.

2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027

2029
2030
2031
2032
2033
2034
2035
2036
2037
2038

2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050

2052
2053
2054
2055
2056
2057
2058
2059
2060
2061
2062
2063

2065
2066
2067
2068
2069
2070

2072
2073
2074

2075 Component C5: Per-Level Tiny MLPs - All Baselines
2076 The paper specifies tiny MLPs per level: 64-channel density
2077 MLP (1 hidden layer) and 128-channel color MLP (2 hidden
2078 layers).

```
2079
2080 1 # \nerfify\ : Exact specifications
2081 2 self.density_heads = nn.ModuleList([
2082 3     MLP(in_dim=self.feature_dim, num_layers=1,
2083 4         layer_width=density_hidden_dim, # 64
2084 5         out_dim=1, activation=nn.ReLU())
2085 6     for _ in range(self.num_pyramid_levels)])
2086 7 self.color_heads = nn.ModuleList([
2087 8     MLP(in_dim=color_in_dim, num_layers=2,
2088 9         layer_width=color_hidden_dim, # 128
2089 10        out_dim=3, activation=nn.ReLU(),
2090 11        out_activation=nn.Sigmoid())
2091 12     for _ in range(self.num_pyramid_levels)])
2092 # CORRECT: Matches paper architecture exactly
```

Listing 119. C5 Implementation: NERFIFY (Score: 1.0)

```
2093
2094 1 # GPT-5: Simplified MLPs
2095 2 class DensityMLP(nn.Module):
2096 3     def __init__(self, in_dim, hidden_dim):
2097 4         self.fc = nn.Sequential(nn.Linear(in_dim,
2098 5             hidden_dim),
2099 6                 nn.ReLU(), nn.Linear(
2100 7                     hidden_dim, 1))
2101 8 class ColorMLP(nn.Module):
2102 9     def __init__(self, in_dim, hidden_dim):
2103 10        self.fc = nn.Sequential(nn.Linear(in_dim + 3,
2104 11            hidden_dim),
2105 12                    nn.ReLU(), nn.Linear(
2106 13                        hidden_dim, 3))
2107 14 # ISSUE: Missing second hidden layer for color
```

Listing 120. C5 Implementation: GPT-5 (Score: 0.6)

```
2110
2111 1 # DeepSeek R1: TinyCUDANN networks
2112 2 self.density_mlp = tcnn.Network(
2113 3     n_input_dims=self._get_input_dims(config),
2114 4     n_output_dims=1,
2115 5     network_config={"otype": "FullyFusedMLP", "
2116 6         "activation": "ReLU",
2117 7             "n_neurons": 64, "n_hidden_layers"
2118 8                 : 1})
2119 9 self.color_mlp = tcnn.Network(
2120 10    n_input_dims=self._get_input_dims(config) + 3,
2121 11    n_output_dims=3,
2122 12    network_config={"otype": "FullyFusedMLP", "
2123 13        "n_neurons": 128,
2124 14            "n_hidden_layers": 2})
2125 15 # ISSUE: Using external library, not pure PyTorch
```

Listing 121. C5 Implementation: DeepSeek R1 (Score: 0.6)

```
2127
2128 1 # Paper2Code: Correct MLP specifications
2129 2 density_layers = []
2130 3 for hidden_dim in density_hidden_dims: # [64]
2131 4     density_layers.append(nn.Linear(current_dim,
2132 5         hidden_dim))
2133 6 density_layers.append(nn.ReLU())
2134 7 self.density_mlp = nn.Sequential(*density_layers)
2135 8 # Color MLP with 2 hidden layers [128, 128]
2136 9 for hidden_dim in color_hidden_dims:
2137 10    color_layers.append(nn.Linear(current_dim,
2138 11        hidden_dim))
2139 12 color_layers.append(nn.ReLU())
2140 13 # CORRECT: Proper layer counts and dimensions
```

Listing 122. C5 Implementation: Paper2Code (Score: 1.0)

```
2142
2143 1 # AutoP2C: Single MLPs for everything
2144 2 self.density_mlp = nn.Sequential(
2145 3     nn.Linear(2 * 10 * 3, 64), nn.ReLU(), nn.Linear
2146 4         (64, 1))
2147 5 # MISSING: No per-level MLPs, just one shared
```

Listing 123. C5 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY and Paper2Code match specifications exactly. GPT-5 and DeepSeek R1 have the right dimensions but implementation issues.

Component C6: Projected Area Computation - All Baselines The paper specifies computing $P(x) = pixel_area \times t^2$ where t is the sample distance along the ray.

```
2148
2149 1 # \nerfify\ : Exact paper formula
2150 2 t_mid = (ray_samples.frustums.starts + ray_samples.
2151 3 frustums.ends) * 0.5
2152 4 pixel_area = ray_samples.frustums.pixel_area
2153 5 if pixel_area.shape != t_mid.shape:
2154 6    pixel_area = pixel_area.expand_as(t_mid)
2155 7 proj_area = pixel_area * torch.clamp(t_mid, min=0.0)
2156 8        ** 2 + 1e-12
2157 9 # CORRECT: P(x) = pixel_area * t^2 as specified
```

Listing 124. C6 Implementation: NERFIFY (Score: 1.0)

```
2158
2159 1 # GPT-5: Simplified approach
2160 2 # In dummy loader:
2161 3 torch.ones(1_024, 1).cuda()*0.5 # projected area
2162 4 # Missing actual computation, uses placeholder
2163 5 # ISSUE: No real projected area calculation
```

Listing 125. C6 Implementation: GPT-5 (Score: 0.6)

```
2164
2165 1 # DeepSeek R1: Heuristic approximation
2166 2 sample_distances = torch.norm(sample_positions - rays
2167 3     ['origins'].unsqueeze(1), dim=-1)
2168 4 # Approximate pixel area as distance^2 * solid_angle
2169 5 pixel_areas = sample_distances ** 2 * 0.01 #
2170 6         heuristic scaling
2171 7 # ERROR: Heuristic formula, not paper specification
```

Listing 126. C6 Implementation: DeepSeek R1 (Score: 0.6)

```
2172
2173 1 # Paper2Code: WRONG FORMULA
2174 2 def compute_integration_area(self, depths):
2175 3     depths = depths.to(torch.float32)
2176 4     integration_area = (depths / self.focal) ** 2
2177 5     return integration_area
2178 6 # CRITICAL ERROR: Should be pixel_area * t^2, not (t/
2179 7     f)^2
```

Listing 127. C6 Implementation: Paper2Code (Score: 0.2)

```
2180
2181 1 # AutoP2C: No implementation
2182 2 def compute_projected_area(self, x):
2183 3     # Placeholder for actual projected area
2184 4         computation
2185 5     return torch.ones_like(x[..., 0])
2186 6 # CRITICAL ERROR: Always returns 1.0
```

Listing 128. C6 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERFIFY implements the correct formula. Paper2Code's error here is particularly damaging as it affects all level selections.

2203 Component C7: Adaptive Level Supervision - All Baselines The paper mentions maintaining an auxiliary structure to track supervised levels during training.

```
1 # \nerify\ : Occupancy grid from parent class
2 # Inherits from NerfactoModel which includes:
3 # self.occupancy_grid (from parent)
4 # self.proposal_networks (for sampling)
5 # Partial implementation through inheritance
6 # ISSUE: Not explicitly tracking supervised levels
```

Listing 129. C7 Implementation: NERIFY (Score: 0.6)

2214
2215
2216
2217
2218

```
1 # GPT-5: No adaptive supervision
2 # No occupancy grid or level tracking
3 # MISSING: Feature not implemented
```

Listing 130. C7 Implementation: GPT-5 (Score: 0.0)

2219
2220
2221
2222
2223
2224

```
1 # DeepSeek R1: Structure present
2 self.occupancy_grid = None
3 self.supervised_levels = None
4 # For tracking supervised levels during training
5 # ISSUE: Declared but never used
```

Listing 131. C7 Implementation: DeepSeek R1 (Score: 0.4)

2225
2226
2227
2228

```
1 # Paper2Code: No adaptive supervision
2 # MISSING: No occupancy grid or adaptive training
```

Listing 132. C7 Implementation: Paper2Code (Score: 0.0)

2229
2230
2231
2232
2233

```
1 # AutoP2C: No supervision concepts
2 # MISSING: Feature not implemented
```

Listing 133. C7 Implementation: AutoP2C (Score: 0.0)

2234 Analysis: None of the baselines fully implement adaptive supervision. NERIFY partially inherits it from Nerfacto.

2235 Component C8: Modern Backbone Integration - All Baselines The paper mentions integration with modern grid-based NeRFs like iNGP, K-Planes, TensoRF.

```
1 # \nerify\ : Full Nerfstudio integration
2 from nerfstudio.models.nerfacto import NerfactoModel,
    NerfactoModelConfig
3 class PynerfModel(NerfactoModel):
4     """Nerfacto-based model with PyNeRF multiscale
        field."""
5     def populate_modules(self):
6         super().populate_modules()
7         # Reuse scene contraction from Nerfacto
8         spatial_distortion = getattr(self.field, "spatial_distortion", None)
9 # CORRECT: Complete framework integration
```

Listing 134. C8 Implementation: NERIFY (Score: 1.0)

2253
2254
2255
2256
2257

```
1 # GPT-5: No framework integration
2 # Standalone implementation
3 # MISSING: Cannot leverage existing infrastructure
```

Listing 135. C8 Implementation: GPT-5 (Score: 0.0)

```
1 # DeepSeek R1: Multiple backbone options
2 def _create_base_nerf(self, config):
3     nerf_type = config.get('type', 'ingp')
4     if nerf_type == 'ingp':
5         return InstantNGPBackbone(config)
```

2264
2265
2266
2267
2268

```
6     elif nerf_type == 'tensorrf':
7         return TensoRFBBackbone(config)
8     elif nerf_type == 'kplanes':
9         return KPlanesBackbone(config)
10 # ISSUE: Stubs only, not implemented
```

Listing 136. C8 Implementation: DeepSeek R1 (Score: 0.8)

2270
2271
2272
2273

```
1 # Paper2Code: No backbone integration
2 # Custom standalone implementation
3 # MISSING: No framework compatibility
```

Listing 137. C8 Implementation: Paper2Code (Score: 0.0)

2275
2276
2277

```
1 # AutoP2C: No backbone concepts
2 # MISSING: Feature not implemented
```

Listing 138. C8 Implementation: AutoP2C (Score: 0.0)

2279 Analysis: Only NERIFY achieves real framework integration. DeepSeek R1 has the structure but lacks implementation.

2280 Component C9: Multiple Interpolation Modes - All Baselines The paper describes both linear interpolation (Equation 5) and Laplacian pyramid summation (Equation 4).

```
1 # \nerify\ : Both interpolation modes
2 if self.interp_mode == "sum":
3     # Laplacian pyramid style (Equation 4)
4     sigma_cum = torch.cumsum(sigma_all, dim=0)
5     sigma = torch.gather(sigma_cum, 0, idx_main).squeeze(0)
6 else:
7     # Linear interpolation (Equation 5)
8     sigma = w * sigma_main + (1.0 - w) * sigma_prev
9 # CORRECT: Both modes implemented
```

Listing 139. C9 Implementation: NERIFY (Score: 1.0)

2298
2299
2300
2301

```
1 # GPT-5: Only linear interpolation
2 c = w.unsqueeze(-1) * c2 + (1 - w).unsqueeze(-1) * c1
3 # MISSING: No Laplacian mode
```

Listing 140. C9 Implementation: GPT-5 (Score: 0.0)

2303
2304
2305
2306
2307
2308

```
1 # DeepSeek R1: Mode selection present
2 def __init__(self, interpolation_method: str = "linear"):
3     self.interpolation_method = interpolation_method
4 # ISSUE: Only linear mode implemented
```

Listing 141. C9 Implementation: DeepSeek R1 (Score: 0.4)

2310
2311
2312
2313

```
1 # Paper2Code: No mode selection
2 # Only implements linear interpolation
3 # MISSING: No Laplacian option
```

Listing 142. C9 Implementation: Paper2Code (Score: 0.0)

2315
2316
2317

```
1 # AutoP2C: No interpolation modes
2 # MISSING: Feature not implemented
```

Listing 143. C9 Implementation: AutoP2C (Score: 0.0)

2319 **Analysis:** Only NERIFY implements both interpolation
 2320 modes as described in the paper.

2321 **Component C10: Specific Hyperparameters - All Base-**

2322 lines

2323 The paper specifies: 8 pyramid levels, scale factor 2,
 base resolution 16, 4 features per level, hash table size 2^{20} .

```
1 # \nerify\ : All hyperparameters correct
2 num_pyramid_levels: int = 8
3 pyramid_scale_factor: float = 2.0
4 features_per_level: int = 4
5 log2_hashmap_size: int = 20
6 density_hidden_dim: int = 64
7 color_hidden_dim: int = 128
8 max_num_iterations=20000
9 train_num_rays_per_batch=8192
10 # CORRECT: Matches paper exactly
```

Listing 144. C10 Implementation: NERIFY (Score: 1.0)

```
1 # GPT-5: Configuration matches
2 @dataclass
3 class PyNeRFConfig:
4     n_levels: int = 8
5     base_resolution: int = 16
6     scale_factor: float = 2.0
7     features_per_level: int = 4
8     hash_size: int = 2**20
9     mlp_hidden_dim: int = 64
10    color_hidden_dim: int = 128
11 # CORRECT: All parameters specified
```

Listing 145. C10 Implementation: GPT-5 (Score: 1.0)

```
1 # DeepSeek R1: Complete configuration
2 'pyramid': {'num_levels': 8, 'scale_factor': 2.0,
3             'base_resolution': 16,
4             'use_shared_features': True},
4 'base_nerf': {'log2_hash_size': 19,
5               'features_per_level': 4}
5 'training': {'batch_size': 8192}
6 # CORRECT: Parameters match paper
```

Listing 146. C10 Implementation: DeepSeek R1 (Score: 1.0)

```
1 # Paper2Code: Most parameters correct
2 self.pyramid_levels: int = int(model_config.get("pyramid_levels", 8))
3 self.base_resolution: int = int(model_config.get("base_resolution", 256)) # Wrong!
4 self.scale_factor: int = int(model_config.get("scale_factor", 2))
5 # ERROR: Base resolution should be 16, not 256
```

Listing 147. C10 Implementation: Paper2Code (Score: 0.6)

```
1 # AutoP2C: Wrong parameters
2 hierarchy_levels=8 # Only this is correct
3 base_resolution=1 # Wrong! Should be 16
4 scaling_factor=2
5 # CRITICAL ERROR: Multiple wrong values
```

Listing 148. C10 Implementation: AutoP2C (Score: 0.0)

2376 **Analysis:** NERIFY , GPT-5, and DeepSeek R1 correctly
 2377 specify all hyperparameters. Paper2Code and AutoP2C have
 2378 critical errors in base resolution.

Table 38. PyNeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	Weighted Avg _{L1M}
NERIFY	1.0	1.0	1.0	1.0	1.0	0.6	1.0	1.0	1.0	1.0	0.97
GPT-5	1.0	1.0	1.0	0.2	0.6	0.6	0.0	0.0	1.0	1.0	0.52
DeepSeek R1	1.0	0.6	0.8	0.6	0.6	0.6	0.4	0.8	0.4	1.0	0.68
Paper2Code	1.0	0.8	1.0	0.6	1.0	0.2	0.0	0.0	0.0	0.6	0.58
AutoP2C	0.0	0.0	0.1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.03

3.3.6. Scoring Analysis

3.3.7. Why Baselines Fail Despite Component Scores

GPT-5 (Score: 0.52 components, 0% trainable)

- **Strengths:** Clean modular implementation with correct pyramid structure and interpolation logic
- **Fatal Issues:**

- No Nerfstudio integration - standalone code cannot leverage existing infrastructure
- Missing training loop and data loading pipeline
- Dummy data loader prevents actual training

- **Result:** Code compiles but cannot train on real data
- **DeepSeek R1 (Score: 0.68 components, 0% trainable)**

- **Strengths:** Comprehensive structure with multi-agent architecture consideration
- **Fatal Issues:**

- Many methods left unimplemented with pass statements
- Heuristic projected area computation (sample_distances**2 * 0.01)
- Incomplete PyNeRFLevel forward pass

- **Result:** Structural skeleton without functional implementation
- **Paper2Code (Score: 0.58 components, 0% trainable)**

- **Strengths:** Most complete baseline with proper pyramid architecture and detailed implementation
- **Fatal Issues:**

- Critical error in projected area formula: uses $(depth/focal)^2$ instead of $pixel_area * t^2$
- No framework integration - standalone implementation
- Missing adaptive supervision and occupancy grid

- **Result:** Despite 1250 lines of code, mathematical errors prevent convergence

AutoP2C (Score: 0.03 components, 0% trainable)

- **Strengths:** Attempts to create training infrastructure
- **Fatal Issues:**

- Complete architectural misunderstanding - no pyramid structure
- Placeholder functions that always return constants
- Import errors for non-existent modules (colmap Python module)
- Self-interpolation bug (interpolates values with themselves)

- **Result:** Complete failure to understand the paper's core concepts

Table 39. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5	R1	P2C	AutoP2C
Pyramid Levels (L)	8	✓	✓	✓	✓	✗
Scale Factor (s)	2	✓	✓	✓	✓	✗
Base Resolution	16	✓	✓	✓	~	✗
Features per Level	4	✓	✓	~	✓	✗
Hash Table Size	2^{20}	✓	✓	✗	✓	✗
Density MLP Hidden	64	✓	✓	✓	✓	✓
Color MLP Hidden	128	✓	✓	✓	✓	✓
Training Iterations	20,000	✓	✓	✗	✓	✓
Batch Size	8,192	✓	~	✓	✓	✓
W Score	—	1.00	0.80	0.70	0.60	0.00

2422

3.3.8. Hyperparameter Fidelity

2423

3.3.9. Conclusion

2424

All baselines attempt to implement PyNeRF with varying levels of sophistication. GPT-5 produces the cleanest modular code but lacks framework integration. DeepSeek R1 creates an ambitious architecture that remains largely unimplemented. Paper2Code achieves the most complete standalone implementation but fails due to a critical mathematical error in projected area computation. AutoP2C completely misunderstands the paper’s core pyramid concept. Only NERFIFY produces immediately trainable code as a complete Nerfstudio plugin with correct mathematical formulations, achieving a 97% semantic completeness score. This stark contrast - 100% trainable rate for NERFIFY versus 0% for all baselines - demonstrates that effective paper-to-code translation for complex vision research requires deep domain specialization beyond what general-purpose systems can provide.

2439

3.4. TensoRF: Tensorial Radiance Fields

2440

3.4.1. Paper Overview

2441

TensoRF [3] revolutionizes neural radiance fields by replacing computationally expensive MLPs with factorized 4D tensors using vector-matrix (VM) decomposition. The method decomposes the radiance field tensor $\mathcal{T} \in \mathbb{R}^{I \times J \times K \times P}$ into compact low-rank components, achieving $100\times$ compression with superior quality. This explicit grid representation with tensor factorization demonstrates that properly structured explicit methods can outperform implicit neural representations in both reconstruction quality and training efficiency, achieving 30-minute training on a single GPU.

2451

3.4.2. Implementation Overview

2452

3.4.3. Novel Components

2453

3.4.4. Quantitative Metrics

2454

3.4.5. Component-by-Component Analysis

2455

Component C1: Vector-Matrix Decomposition - All Baselines

2457

The paper defines VM decomposition as: $\mathcal{T} = \sum_{r=1}^R \mathbf{v}_r^X \circ \mathbf{M}_r^{YZ} + \mathbf{v}_r^Y \circ \mathbf{M}_r^{XZ} + \mathbf{v}_r^Z \circ \mathbf{M}_r^{XY}$ where \circ denotes outer product.

Table 40. TensoRF Implementation Summary Across All Baselines

Aspect	NERFIFY	GPT-5	DeepSeek R1	Paper2Code	AutoP2C
Lines of Code	467	298	412	856	234
File Organization	Plugin	Single	Single	Multi-file	Multi-file
VM Decomposition	✓	✓	Partial	✗	✗
Density Factorization	✓	✓	Partial	✗	✗
Appearance Matrix B	✓	✓	✗	✗	✗
L1 Regularization	✓	Partial	Partial	✗	✗
TV Regularization	✓	✗	✗	✗	✗
Trainable	✓	✗	✗	✗	✗

Note: NERFIFY always produces trainable code as a complete Nerfstudio plugin

Table 41. Novel Components in TensoRF with Importance Weights

ID	Component	Weight w_i
C1	Vector-Matrix (VM) Decomposition	0.20
C2	Density Grid Factorization	0.15
C3	Appearance Grid Factorization	0.15
C4	Global Appearance Matrix B	0.10
C5	Factor-level Trilinear Interpolation	0.10
C6	L1 Sparsity Regularization	0.08
C7	Total Variation Regularization	0.07
C8	Coarse-to-Fine Upsampling	0.05
C9	SH/MLP Decoder	0.05
C10	Separate Density/Appearance Architecture	0.05

Table 42. TensoRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	1.00	0.00	0.00	0.95	0.98
GPT-5	0.70	0.10	0.20	0.75	0.72
DeepSeek R1	0.60	0.20	0.20	0.70	0.65
Paper2Code	0.20	0.30	0.50	0.30	0.12
AutoP2C	0.10	0.20	0.70	0.15	0.28

```

1 # \nerfify\ : Complete VM decomposition with all
2     three modes
3 self.vx_sigma = nn.Parameter(0.1 * torch.randn(self.
4     R_sigma, self.Nx))
5 self.vy_sigma = nn.Parameter(0.1 * torch.randn(self.
6     R_sigma, self.Ny))
7 self.vz_sigma = nn.Parameter(0.1 * torch.randn(self.
8     R_sigma, self.Nz))
9 self.MyZ_sigma = nn.Parameter(0.1 * torch.randn(self.
10    R_sigma, self.Ny, self.Nz))
11 self.MxZ_sigma = nn.Parameter(0.1 * torch.randn(self.
12    R_sigma, self.Nx, self.Nz))
13 self.MxY_sigma = nn.Parameter(0.1 * torch.randn(self.
14    R_sigma, self.Nx, self.Ny))
15 Ax = (vx * MyZ).sum(dim=0) # v_r^X o M_r^YZ
16 Ay = (vy * MxZ).sum(dim=0) # v_r^Y o M_r^XZ
17 Az = (vz * MxY).sum(dim=0) # v_r^Z o M_r^XY
18 sigma_raw = Ax + Ay + Az + self.density_bias
19 # CORRECT: Exact VM decomposition from paper Eq. 7

```

Listing 149. C1 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Correct math but non-standard PyTorch usage
2 class VMDecomp(nn.Module):
3     def forward(self, xyz):
4         x, y, z = xyz
5         vx = F.linear(x, self.vectors[0].T)

```

2460
2461
2462
2463
2464
2465
2466
2467
2468
2469
2470
2471
2472
2473
2474
2475
2476
2477
2478
2479
2480

2487
2488
2489

```

6     myz = F.bilinear(y, z, self.matrices[0])
7     result += (vx * myz).sum(dim=-1)
8 # ISSUE: F.bilinear doesn't work this way in PyTorch

```

Listing 150. C1 Implementation: GPT-5 (Score: 0.9)

2491
2492
2493
2494
2495
2496
2497
2498

```

1 # DeepSeek R1: Structure present but flawed execution
2 class VMDecomposition(nn.Module):
3     def forward(self, xyz):
4         vx = F.linear(x, self.vectors[0].T)
5         myz = F.bilinear(y, z, self.matrices[0])
6         result += (vx * myz).sum(dim=-1)
7 # CRITICAL ERROR: Bilinear function misused

```

Listing 151. C1 Implementation: DeepSeek R1 (Score: 0.6)

2500
2501
2502
2503
2504
2505
2506
2507
2508

```

1 # Paper2Code: Wrong architecture - implements NeRF
2 class TensoRFModel(nn.Module):
3     def __init__(self, decomposition_type='VM'):
4         self.density_mlp = self._build_mlp(64, 1)
5         self.color_mlp = self._build_mlp(128, 2)
6 # CRITICAL ERROR: Uses MLPs, not tensor factorization !

```

Listing 152. C1 Implementation: Paper2Code (Score: 0.0)

2509
2510
2511
2512
2513
2514

```

1 # AutoP2C: No model at all
2 class BlenderDataset:
3     def load_images(self):
4         self.images.append(image)
5 # MISSING: Entire TensorF model absent

```

Listing 153. C1 Implementation: AutoP2C (Score: 0.0)

2516
2517
2518
2519
2520
2521

Analysis: VM decomposition is the core innovation. Only NERFIFY implements it perfectly. GPT-5 has correct math but PyTorch issues. DeepSeek R1 attempts but fails. Paper2Code/AutoP2C completely miss it.

Component C2: Density Grid Factorization - All Baselines

The density grid $\mathcal{G}_\sigma \in \mathbb{R}^{I \times J \times K}$ uses VM decomposition with $R_\sigma = 16$ components.

2522
2523
2524
2525
2526
2527
2528
2529
2530
2531
2532
2533
2534
2535
2536
2537
2538
2539
2540
2541
2542
2543
2544
2545
2546

```

1 # \nerfify\ : Perfect density factorization
2 def get_density(self, ray_samples: RaySamples):
3     x, y, z = self._flatten_positions(ray_samples)
4     vx = _interp_vector(self.vx_sigma, x)
5     vy = _interp_vector(self.vy_sigma, y)
6     vz = _interp_vector(self.vz_sigma, z)
7     MyZ = _interp_plane(self.MyZ_sigma, y, z)
8     MxZ = _interp_plane(self.MxZ_sigma, x, z)
9     MxY = _interp_plane(self.MxY_sigma, x, y)
10    Ax = (vx * MyZ).sum(dim=0)
11    Ay = (vy * MxZ).sum(dim=0)
12    Az = (vz * MxY).sum(dim=0)
13    sigma = trunc_exp(Ax + Ay + Az + self.
14        density_bias)
14 # CORRECT: Density grid with VM decomposition and
15     trunc_exp

```

Listing 154. C2 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Correct density factorization
2 def get_density(self, xyz):
3     return F.softplus(self.density(xyz))
4 # CORRECT: Proper density with softplus activation

```

Listing 155. C2 Implementation: GPT-5 (Score: 0.9)

```

1 # DeepSeek R1: Mostly correct but missing details
2 def forward_density(self, pos):
3     for r in range(self.R):
4         density_vals = density_vals + term_x + term_y
4         + term_z
5     return density_vals
6 # ISSUE: Missing activation function

```

Listing 156. C2 Implementation: DeepSeek R1 (Score: 0.8)

```

1 # Paper2Code: Wrong - uses MLPs
2 def forward(self, x, d):
3     encoded_x = self.encode_position(x)
4     density = self.density_mlp(encoded_x)
5 # CRITICAL ERROR: MLP with positional encoding

```

Listing 157. C2 Implementation: Paper2Code (Score: 0.0)

```

1 # AutoP2C: No density implementation
2 # MISSING: No density code found

```

Listing 158. C2 Implementation: AutoP2C (Score: 0.0)

Analysis: Density factorization correctly splits from appearance. NERFIFY and GPT-5 implement correctly. DeepSeek R1 partially, Paper2Code/AutoP2C fail completely.

Component C3: Appearance Grid Factorization - All Baselines

Appearance grid $\mathcal{G}_c \in \mathbb{R}^{I \times J \times K \times P}$ with $P = 27$ features and $R_c = 48$ components.

```

1 # \nerfify\ : Complete appearance factorization
2 self.vx_c = nn.Parameter(0.1 * torch.randn(self.R_c,
2      self.Nx))
3 self.vy_c = nn.Parameter(0.1 * torch.randn(self.R_c,
3      self.Ny))
4 self.vz_c = nn.Parameter(0.1 * torch.randn(self.R_c,
4      self.Nz))
5 self.MyZ_c = nn.Parameter(0.1 * torch.randn(self.R_c,
5      self.Ny, self.Nz))
6 self.MxZ_c = nn.Parameter(0.1 * torch.randn(self.R_c,
6      self.Nx, self.Nz))
7 self.MxY_c = nn.Parameter(0.1 * torch.randn(self.R_c,
7      self.Nx, self.Ny))
8 concat_contrib = torch.cat([Ax, Ay, Az], dim=0).
8     transpose(0, 1)
9 app_feat = self.B(concat_contrib) # [M, P=27]
10 # CORRECT: Separate appearance grid with feature
10     dimension 27

```

Listing 159. C3 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Present but simplified
2 self.appearance = VMDecomposition(
3     config.resolution,
4     config.num_appearance_components)
5 # ISSUE: Missing feature dimension details

```

Listing 160. C3 Implementation: GPT-5 (Score: 0.7)

```

1 # DeepSeek R1: Partial implementation
2 def forward_appearance(self, pos):
3     for r in range(self.R):
4         feat_list.append(term_x)
5         feat_list.append(term_y)
6         feat_list.append(term_z)
7 # ERROR: Missing matrix B transformation

```

Listing 161. C3 Implementation: DeepSeek R1 (Score: 0.6)

```

2548
2549
2550
2551
2552
2553
2554
2555
2556
2557
2558
2559
2560
2561
2562
2563
2564
2565
2566
2567
2568
2569
2570
2571
2572
2573
2574
2575
2576
2577
2578
2579
2580
2581
2582
2583
2584
2585
2586
2587
2588
2589
2590
2591
2592
2593
2594
2595
2596
2597
2598
2599
2600
2601
2602
2603
2604
2605
2606
2607
2608
2609
2610

```

2613
2614
2615

```
1 # Paper2Code: No appearance grid
2 color = self.color_mlp(encoded_x)
3 # CRITICAL ERROR: Uses MLP, not factorized grid
```

Listing 162. C3 Implementation: Paper2Code (Score: 0.0)

2617
2618
2619

```
1 # AutoP2C: Missing appearance
2 # MISSING: No appearance implementation
```

Listing 163. C3 Implementation: AutoP2C (Score: 0.0)

Analysis: Appearance requires separate VM decomposition with 48 components and 27 features. Only NERFIFY gets all details right.

Component C4: Global Appearance Matrix B - All Baselines

Matrix $\mathbf{B} \in \mathbb{R}^{P \times 3R_c}$ maps concatenated features to final appearance.

2628
2629
2630
2631
2632
2633
2634
2635

```
1 # \nerfify\ : Perfect matrix B implementation
2 self.B = nn.Linear(3 * self.R_C, self.P, bias=False)
3 concat_contrib = torch.cat([Ax, Ay, Az], dim=0).
    transpose(0, 1)
4 app_feat = self.B(concat_contrib) # [M, P]
5 # CORRECT: Matrix B as linear layer without bias
```

Listing 164. C4 Implementation: NERFIFY (Score: 1.0)

2636
2637
2638
2639
2640
2641
2642
2643
2644

```
1 # GPT-5: Correct matrix B
2 self.B = nn.Parameter(
3     torch.randn(config.appearance_dim,
4                 3 * config.num_appearance_components))
5 features = torch.matmul(vm_features, self.B.T)
6 # CORRECT: Proper matrix multiplication
```

Listing 165. C4 Implementation: GPT-5 (Score: 1.0)

2645
2646
2647
2648

```
1 # DeepSeek R1: Missing matrix B
2 # No global appearance matrix in code
3 # CRITICAL ERROR: Matrix B completely absent
```

Listing 166. C4 Implementation: DeepSeek R1 (Score: 0.0)

2650
2651
2652
2653
2654

```
1 # Paper2Code: No matrix B
2 # CRITICAL ERROR: Architecture doesn't include matrix
    B
```

Listing 167. C4 Implementation: Paper2Code (Score: 0.0)

2655
2656
2657

```
1 # AutoP2C: No matrix B
2 # MISSING: No model implementation
```

Listing 168. C4 Implementation: AutoP2C (Score: 0.0)

Analysis: Matrix B is crucial for appearance. Only NERFIFY and GPT-5 implement it. Its absence severely impacts quality in other baselines.

Component C5: Factor-level Trilinear Interpolation - All Baselines

Efficient interpolation at factor level rather than dense grid: $\mathcal{A}_r^X(\mathbf{x}) = \mathbf{v}_r^X(x) \cdot \mathbf{M}_r^{YZ}(y, z)$

2666
2667
2668
2669
2670
2671
2672
2673
2674
2675
2676
2677

```
1 # \nerfify\ : Efficient factor-level interpolation
2 def _interp_vector(vec: Tensor, coord: Tensor) ->
    Tensor:
3     scaled = coord.clamp(0.0, 1.0) * (N - 1)
4     i0 = torch.floor(scaled).long().clamp(0, N - 1)
5     il = (i0 + 1).clamp(0, N - 1)
6     w = (scaled - i0.float()).clamp(0.0, 1.0)
7     v0 = torch.gather(vec, 1, idx0)
8     vl = torch.gather(vec, 1, idx1)
9     return (1.0 - w)[None, :] * v0 + w[None, :] * vl
10 # CORRECT: Efficient linear interpolation on factors
```

Listing 169. C5 Implementation: NERFIFY (Score: 1.0)

2679
2680
2681
2682
2683
2684
2685

```
1 # GPT-5: Correct interpolation approach
2 def interp(self, v, x):
3     idx0 = torch.clamp(x.long(), 0, v.shape[1]-2)
4     w = x - idx0.float()
5     return (1-w)*v[:, idx0] + w*v[:, idx0+1]
6 # CORRECT: Factor interpolation implemented
```

Listing 170. C5 Implementation: GPT-5 (Score: 0.9)

2687
2688
2689
2690
2691
2692
2693
2694

```
1 # DeepSeek R1: Interpolation present but inefficient
2 def _interpld(self, vec: torch.Tensor, coord: torch.
    Tensor):
3     coord = coord.clamp(0, length - 1)
4     idx0 = torch.floor(coord).long()
5     t = coord - idx0.float()
6 # ISSUE: Suboptimal implementation
```

Listing 171. C5 Implementation: DeepSeek R1 (Score: 0.7)

2696
2697
2698

```
1 # Paper2Code: No factor interpolation
2 # CRITICAL ERROR: Uses positional encoding instead
```

Listing 172. C5 Implementation: Paper2Code (Score: 0.0)

2700
2701
2702

```
1 # AutoP2C: No interpolation
2 # MISSING: No model code
```

Listing 173. C5 Implementation: AutoP2C (Score: 0.0)

2704
2705
2706

Analysis: Factor-level interpolation is key to efficiency. NERFIFY and GPT-5 implement correctly, DeepSeek R1 partially.

Component C6: L1 Sparsity Regularization - All Baselines

L1 regularization on density factors with weight $\omega = 0.0004$.

2711
2712
2713
2714
2715
2716
2717
2718
2719
2720
2721
2722
2723
2724

```
1 # \nerfify\ : Complete L1 on all density factors
2 l1_terms = [
3     self.vx_sigma.abs().mean(),
4     self.vy_sigma.abs().mean(),
5     self.vz_sigma.abs().mean(),
6     self.MyZ_sigma.abs().mean(),
7     self.MxZ_sigma.abs().mean(),
8     self.MxY_sigma.abs().mean(),
9 ]
10 losses["l1_density"] = sum(l1_terms) / len(l1_terms)
11 loss_dict["l1_density"] = self.config.l1_density_mult
    * regs["l1_density"]
12 # CORRECT: L1 with exact weight 4e-4
```

Listing 174. C6 Implementation: NERFIFY (Score: 1.0)

```

2726 1 # GPT-5: L1 present but simplified
2727 2 ll_loss = 0.0
2728 3 for param in model.density.parameters():
2729 4     ll_loss += torch.abs(param).mean()
2730 5 total_loss = render_loss + config.ll_weight * ll_loss
2731 6 # ISSUE: Hardcoded weight
2732

```

Listing 175. C6 Implementation: GPT-5 (Score: 0.7)

```

2734 1 # DeepSeek R1: Wrong weight
2735 2 def compute_ll_reg(self):
2736 3     ll_reg += torch.mean(torch.abs(param))
2737 4 total_loss = mse_loss + 1e-4 * reg_loss
2738 5 # ERROR: Uses 1e-4 instead of 4e-4
2739

```

Listing 176. C6 Implementation: DeepSeek R1 (Score: 0.5)

```

2741 1 # Paper2Code: No L1 regularization
2742 2 # MISSING: No regularization terms
2743

```

Listing 177. C6 Implementation: Paper2Code (Score: 0.0)

```

2745 1 # AutoP2C: No regularization
2746 2 # MISSING: No training code
2747

```

Listing 178. C6 Implementation: AutoP2C (Score: 0.0)

Analysis: L1 sparsity crucial for compact models. Only NERFIFY gets weight exact, others have errors or missing.

Component C7: Total Variation Regularization - All Baselines

TV regularization for smooth factors, especially important for few-shot scenarios.

```

2755 1 # \nerify\ : Complete TV regularization
2756 2 def _tv_1d(self, vec: Tensor) -> Tensor:
2757 3     return (vec[:, 1:] - vec[:, :-1]).abs().mean()
2758 4 def _tv_2d(self, mat: Tensor) -> Tensor:
2759 5     tv_u = (mat[:, 1:, :] - mat[:, :-1, :]).abs().mean()
2760 6     tv_v = (mat[:, :, 1:] - mat[:, :, :-1]).abs().mean()
2761 7     return 0.5 * (tv_u + tv_v)
2762 8 tv_density = (self._tv_1d(self.vx_sigma) + self.
2763 9     _tv_1d(self.vy_sigma) +
2764 10    self._tv_1d(self.vz_sigma) + self.
2765 11    self._tv_2d(self.MyZ_sigma) +
2766        self._tv_2d(self.MxZ_sigma) + self.
2767        self._tv_2d(self.MxY_sigma)) / 6.0
2768
2769 # CORRECT: TV on both 1D vectors and 2D matrices
2770
2771

```

Listing 179. C7 Implementation: NERFIFY (Score: 1.0)

```

2773 1 # GPT-5: No TV regularization
2774 2 # MISSING: TV not implemented
2775

```

Listing 180. C7 Implementation: GPT-5 (Score: 0.0)

```

2777 1 # DeepSeek R1: No TV regularization
2778 2 # MISSING: TV completely absent
2779

```

Listing 181. C7 Implementation: DeepSeek R1 (Score: 0.0)

```

2781 1 # Paper2Code: No TV
2782 2 # MISSING: No regularization
2783

```

Listing 182. C7 Implementation: Paper2Code (Score: 0.0)

```

2785 1 # AutoP2C: No TV
2786 2 # MISSING: No model
2787

```

Listing 183. C7 Implementation: AutoP2C (Score: 0.0)

Analysis: TV regularization critical for few-shot. Only NERFIFY implements it correctly with both 1D and 2D variants.

Component C8: Coarse-to-Fine Upsampling - All Baselines

Progressive resolution increase from 128^3 to 300^3 during training.

```

2796 1 # \nerify\ : Upsampling via training schedule
2797 2 upsampling_schedule: [2000, 3000, 4000, 5500, 7000]
2798 3 grid_resolution: (128, 128, 128) # Initial
2799 4 final_resolution: 300 # Target
2800 5 # Training config handles progressive upsampling
2801 6 # ISSUE: Not explicit upsampling, handled by
2802 scheduler
2803

```

Listing 184. C8 Implementation: NERFIFY (Score: 0.8)

```

2805 1 # GPT-5: No coarse-to-fine
2806 2 # MISSING: No upsampling implementation
2807

```

Listing 185. C8 Implementation: GPT-5 (Score: 0.0)

```

2809 1 # DeepSeek R1: No upsampling
2810 2 # MISSING: No coarse-to-fine schedule
2811

```

Listing 186. C8 Implementation: DeepSeek R1 (Score: 0.0)

```

2813 1 # Paper2Code: No upsampling
2814 2 # MISSING: No progressive resolution
2815

```

Listing 187. C8 Implementation: Paper2Code (Score: 0.0)

```

2817 1 # AutoP2C: No upsampling
2818 2 # MISSING: No training pipeline
2819

```

Listing 188. C8 Implementation: AutoP2C (Score: 0.0)

Analysis: Coarse-to-fine speeds convergence. NERFIFY implements via schedule, others completely miss it.

Component C9: SH/MLP Decoder - All Baselines

View-dependent color decoder using spherical harmonics or small MLP.

```

2821 1 # \nerify\ : Complete SH encoding + MLP decoder
2822 2 self.direction_encoding = SHEncoding(levels=3,
2823     implementation="torch")
2824 3 self.color_head = MLP(
2825     in_dim=self.P + self.direction_encoding.
2826     get_out_dim(),
2827     num_layers=color_mlp_layers,
2828     layer_width=color_mlp_width,
2829     out_dim=3,
2830     activation=nn.ReLU(),
2831     out_activation=nn.Sigmoid())
2832 10 # CORRECT: SH encoding with 2-layer MLP decoder
2833

```

Listing 189. C9 Implementation: NERFIFY (Score: 1.0)

```

2840
2841 1 # GPT-5: MLP decoder present
2842 2 self.mlp = nn.Sequential(
2843 3     nn.Linear(3 * config.num_appearance_components,
2844 4         128),
2845 5     nn.ReLU(),
2846 6     nn.Linear(128, 3),
2847 7     nn.Sigmoid())
2848 7 # ISSUE: Missing SH option

```

Listing 190. C9 Implementation: GPT-5 (Score: 0.7)

```

2850
2851 1 # DeepSeek R1: Partial decoder
2852 2 self.mlp = nn.Sequential(
2853 3     nn.Linear(mlp_input_dim, mlp_hidden_size),
2854 4     nn.ReLU(),
2855 5     nn.Linear(mlp_hidden_size, 3))
2856 6 # ERROR: Missing SH, incomplete MLP

```

Listing 191. C9 Implementation: DeepSeek R1 (Score: 0.6)

```

2858
2859 1 # Paper2Code: Wrong decoder type
2860 2 self.color_mlp = self._build_mlp(128, 2)
2861 3 # CRITICAL ERROR: Wrong architecture

```

Listing 192. C9 Implementation: Paper2Code (Score: 0.2)

```

2863
2864 1 # AutoP2C: Evaluation mentions MLP
2865 2 # MISSING: No actual implementation

```

Listing 193. C9 Implementation: AutoP2C (Score: 0.1)

Analysis: Decoder maps appearance features to RGB. NERIFY implements both SH and MLP options correctly.

Component C10: Separate Density/Appearance Architecture - All Baselines

Independent factorizations for density and appearance with different ranks.

```

2874
2875 1 # \nerify\ : Completely separate density/appearance
2876 2 self.density_factor = TensorFactorization(
2877 3     grid_shape=grid_shape,
2878 4     num_components={"density": 16}, # R_sigma = 16
2879 5     decomposition=decomposition,
2880 6     grid_type="density")
2881 7 self.appearance_factor = TensorFactorization(
2882 8     grid_shape=grid_shape,
2883 9     num_components={"appearance": 48}, # R_c = 48
2884 10    decomposition=decomposition,
2885 11    grid_type="appearance")
2886 12 # CORRECT: Separate factors with different ranks

```

Listing 194. C10 Implementation: NERIFY (Score: 1.0)

```

2887
2888 1 # GPT-5: Separate but less clear
2889 2 self.density = VMDecomposition(
2890 3     config.resolution,
2891 4     config.num_density_components)
2892 5 self.appearance = VMDecomposition(
2893 6     config.resolution,
2894 7     config.num_appearance_components)
2895 8 # CORRECT: Separate architectures

```

Listing 195. C10 Implementation: GPT-5 (Score: 0.8)

```

2896
2897 1 # DeepSeek R1: Attempted separation
2898 2 self.density_factor = TensorFactorization(
2899 3     grid_type="density")
2900 4 self.appearance_factor = TensorFactorization(
2901 5     grid_type="appearance")
2902 6 # ISSUE: Implementation details incomplete

```

Listing 196. C10 Implementation: DeepSeek R1 (Score: 0.7)

```

2905
2906 1 # Paper2Code: No separation
2907 2 # CRITICAL ERROR: Single MLP for both

```

Listing 197. C10 Implementation: Paper2Code (Score: 0.0)

```

2909
2910 1 # AutoP2C: No architecture
2911 2 # MISSING: No model

```

Listing 198. C10 Implementation: AutoP2C (Score: 0.0)

Analysis: Separation allows different ranks for density (16) and appearance (48). Critical for efficiency/quality balance.

3.4.6. Scoring Analysis

Table 43. TensoRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	Weighted Avg _{LLM}
NERIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.8	1.0	1.0	0.98
GPT-5	0.9	0.9	0.7	1.0	0.9	0.7	0.0	0.0	0.7	0.8	0.72
DeepSeek R1	0.6	0.8	0.6	0.0	0.7	0.5	0.0	0.0	0.6	0.7	0.65
Paper2Code	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.2	0.0	0.12
AutoP2C	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.1	0.0	0.28

3.4.7. Why Baselines Fail Despite Component Scores

GPT-5 (Score: 0.72 components, 0% trainable)

- Strengths:** Correct VM decomposition, matrix B implementation, density/appearance separation

- Fatal Issues:**

- No Nerfstudio integration - requires complete pipeline assembly
- Missing TV regularization entirely (C7 = 0.0)
- No coarse-to-fine upsampling (C8 = 0.0)

- Result:** Mathematically sound but requires hours of engineering to make trainable

DeepSeek R1 (Score: 0.65 components, 0% trainable)

- Strengths:** Understands tensor decomposition concept, attempts proper separation

- Fatal Issues:**

- PyTorch tensor operations cause runtime crashes (F.bilinear misuse)
- Missing matrix B completely (C4 = 0.0)
- No TV regularization (C7 = 0.0) causes training instability

- Result:** Training crashes immediately on tensor shape mismatches

Paper2Code (Score: 0.12 components, 0% trainable)

- Strengths:** Basic volume rendering loop exists

- Fatal Issues:**

- Complete architectural misunderstanding - implements vanilla NeRF
- Zero correct components except partial decoder (C9 = 0.2)
- Uses MLPs with positional encoding throughout

- Result:** Trains as NeRF, not TensoRF - fundamentally wrong method

- 2949 **AutoP2C (Score: 0.28 components, 0% trainable)**
- 2950 • **Strengths:** Dataset loading code present
- 2951 • **Fatal Issues:**
- No model implementation whatsoever
 - Only evaluation stub mentions components ($C_9 = 0.1$)
 - Missing entire TensoRF architecture
- 2955 • **Result:** Code doesn't compile - no model to train

2956 3.4.8. Hyperparameter Fidelity

Table 44. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5	R1	P2C	AutoP2C
Grid Resolution	128 ³	✓	✓	✗	–	–
Density Rank R_σ	16	✓	✗	✓	–	–
Appearance Rank R_c	48	✓	✓	✗	–	–
Feature Dim P	27	✓	✓	✓	–	–
L1 Weight	4e-4	✓	✓	✗	–	–
TV Weight	Variable	✓	✗	✗	–	–
Learning Rate	0.02	✗	✓	✓	✗	–
Training Steps	30K	✓	✗	✗	–	–
MLP Width	128	✓	✓	✓	✗	–
W Score	–	0.95	0.75	0.70	0.30	0.15

2957 3.4.9. Conclusion

All baselines attempt to implement TensoRF with dramatically varying levels of success. GPT-5 achieves mathematically correct VM decomposition and matrix B but lacks critical TV regularization and framework integration. DeepSeek R1 understands the tensor factorization concept but fails catastrophically on PyTorch execution and missing matrix B. Paper2Code fundamentally misunderstands the method, implementing vanilla NeRF with MLPs and positional encoding instead of tensor factorization. AutoP2C provides only dataset loading code with no model implementation whatsoever. Despite GPT-5 achieving 72% component coverage and DeepSeek R1 reaching 65%, none produce trainable code due to missing Nerfstudio integration, incorrect tensor operations, or complete architectural misunderstandings. Only NERFIFY produces immediately trainable code as a complete Nerfstudio plugin, achieving 100% trainable rate versus 0% for all baselines, demonstrating the critical importance of domain-specific synthesis for complex computer vision research.

2977 3.5. Tetra-NeRF: Representing Neural Radiance Fields Using Tetrahedra

2979 3.5.1. Paper Overview

Tetra-NeRF [12] introduces an adaptive scene representation using Delaunay triangulation of point clouds, replacing uniform voxel grids with tetrahedra that naturally concentrate resolution near surfaces. The method achieves faster training and rendering by using barycentric interpolation within tetrahedra, eliminating the need for dense regular grids while maintaining quality comparable to state-of-the-art methods.

2987 3.5.2. Implementation Overview

Table 45. Tetra-NeRF Implementation Summary Across All Baselines

Aspect	NERFIFY	GPT-5	DeepSeek R1	Paper2Code	AutoP2C
Lines of Code	412	189	267	845	621
File Organization	Plugin	Single	Single	Multi-file	Multi-file
Delaunay Triangulation	✓	✓	✓	Attempted	✗
Barycentric Coords	✓	Simplified	✓	✗	✗
Feature Interpolation	✓	✓	✓	Partial	✗
Volume Rendering	✓	~	~	✗	✗
Trainable	✓	✗	✗	✗	✗

Note: NERFIFY always produces trainable code as a complete Nerfstudio plugin

2988 3.5.3. Novel Components

Table 46. Novel Components in Tetra-NeRF with Importance Weights

ID	Component	Weight w_i
C1	Delaunay triangulation for adaptive subdivision	0.20
C2	Barycentric coordinates using volume ratios (Eq. 2)	0.15
C3	Feature interpolation via barycentric weights	0.15
C4	Tetrahedral grid data structure	0.10
C5	Adaptive non-uniform resolution	0.10
C6	Hierarchical coarse-to-fine sampling	0.10
C7	Volume rendering with tetrahedral queries	0.08
C8	Tetrahedron volume computation	0.05
C9	Nerfstudio plugin integration	0.05
C10	RAdam optimizer with exponential decay	0.02

2989 3.5.4. Quantitative Metrics

Table 47. Tetra-NeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	1.00	0.00	0.00	1.00	1.00
GPT-5	0.50	0.25	0.25	0.60	0.58
DeepSeek R1	0.63	0.25	0.13	0.70	0.72
Paper2Code	0.13	0.25	0.63	0.20	0.22
AutoP2C	0.00	0.13	0.88	0.00	0.08

2990 3.5.5. Component-by-Component Analysis

Component C1: Delaunay Triangulation - All Baselines

The paper specifies: “We represent the scene as a dense triangulation of the input point cloud, where the scene is a set of non-overlapping tetrahedra whose union is the convex hull of the original point cloud.”

```
1 # \nerfify\ : Proper tetrahedral grid with simplex
      decomposition
2 class TetraNerfField(Field):
3     def _tetrahedral_interpolate(self, positions_norm
        ):
4         # 6-simplex decomposition of unit cube
5         scales = torch.tensor([self.nx-1, self.ny-1,
      self.nz-1])
```

```

3005      gpos = positions_norm * scales
3006      base = torch.floor(gpos).long()
3007      u = (gpos - base.to(gpos.dtype)).contiguous()
3008      # Sort u components for simplex determination
3009      u_sorted, perm = torch.sort(u, dim=1,
3010          descending=False)
3011      k1, k2, k3 = perm[:, 0], perm[:, 1], perm[:, 2]
3012
3013  # CORRECT: Full tetrahedral decomposition with proper
3014  # simplex

```

Listing 199. C1 Implementation: NERFIFY (Score: 1.0)

```

3016  # GPT-5: Basic Delaunay implementation
3017  def build_tetrahedra(points):
3018      tri = Delaunay(points)
3019      verts = torch.tensor(points, dtype=torch.float32,
3020          device=device)
3021      tets = torch.tensor(tri.simplices, dtype=torch.
3022          long, device=device)
3023      return verts, tets
3024
3025  # CORRECT: Uses scipy.spatial.Delaunay correctly

```

Listing 200. C1 Implementation: GPT-5 (Score: 1.0)

```

3027  # DeepSeek R1: Complete tetrahedral grid class
3028  class TetrahedralGrid:
3029      def __init__(self, points: torch.Tensor):
3030          self.points = points.cpu().numpy()
3031          self.tetrahedra = Delaunay(self.points)
3032          self.features = nn.Parameter(
3033              torch.randn(len(points), 64) * 1e-4)
3034
3035  # CORRECT: Proper Delaunay with feature
3036  # initialization

```

Listing 201. C1 Implementation: DeepSeek R1 (Score: 1.0)

```

3038  # Paper2Code: Wrong approach
3039  class DatasetLoader:
3040      def perform_delaunay_triangulation(self):
3041          delaunay = scipy.spatial.Delaunay(self.
3042              point_cloud)
3043          self.tetrahedra = delaunay.simplices
3044
3045  # ERROR: Not integrated with rendering pipeline

```

Listing 202. C1 Implementation: Paper2Code (Score: 0.2)

```

3047  # AutoP2C: No triangulation at all
3048  class ColmapDataset:
3049      def generate_point_cloud(self):
3050          self.point_cloud = np.array([point.xyz for
3051              point in
3052                  reconstruction.
3053                  points3D.values()])
3054
3055  # CRITICAL ERROR: Just loads points, no tetrahedra

```

Listing 203. C1 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY implements efficient simplex decomposition. GPT-5 and DeepSeek R1 correctly use Delaunay. Paper2Code attempts but fails to integrate. AutoP2C misses entirely.

Component C2: Barycentric Coordinates - All Baselines

The paper defines barycentric coordinates using volume ratios: $\lambda_i = \frac{V_i}{V_{\text{total}}}$ where V_i is the volume of sub-tetrahedron.

```

3063  # \nerfify\ : Efficient simplex-based computation
3064  def _tetrahedral_interpolate(self, positions_norm):
3065      u_sorted, perm = torch.sort(u, dim=1, descending=
3066          False)

```

```

3070  # Barycentric weights for simplex vertices
3071  w0 = 1.0 - u_sorted[:, 2]
3072  w1 = u_sorted[:, 2] - u_sorted[:, 1]
3073  w2 = u_sorted[:, 1] - u_sorted[:, 0]
3074  w3 = u_sorted[:, 0]
3075  # Weighted interpolation
3076  f = (f0 * w0.unsqueeze(-1) + f1 * w1.unsqueeze
3077      (-1) +
3078      f2 * w2.unsqueeze(-1) + f3 * w3.unsqueeze
3079      (-1))
3080
3081  # CORRECT: Efficient vectorized barycentric
3082  # computation

```

Listing 204. C2 Implementation: NERFIFY (Score: 1.0)

```

3083  # GPT-5: Linear system approach
3084  def barycentric_coords(p, tet_verts):
3085      V0, V1, V2, V3 = tet_verts
3086      M = torch.stack([V1-V0, V2-V0, V3-V0], dim=1)
3087      v = p - V0
3088      sol = torch.linalg.solve(M, v)
3089      l1, l2, l3 = sol
3090      l0 = 1 - l1 - l2 - l3
3091
3092  # ISSUE: No volume validation, less robust

```

Listing 205. C2 Implementation: GPT-5 (Score: 0.4)

```

3095  # DeepSeek R1: Volume-based as in paper
3096  def tetrahedron_volume(self, a, b, c, d):
3097      return torch.abs(torch.dot(a - d,
3098          torch.cross(b - d, c - d))) /
3099          6.0
3100
3101  def get_barycentric_coords(self, query_points,
3102      tetra_indices):
3103      V_total = self.tetrahedron_volume(v0, v1, v2, v3)
3104      V_x234 = self.tetrahedron_volume(query_points[i],
3105          v1, v2, v3)
3106      V_1x34 = self.tetrahedron_volume(v0, query_points
3107          [i], v2, v3)
3108      barycentric_coords[i] = torch.tensor([V_x234/
3109          V_total, ...])
3110
3111  # CORRECT: Exact paper formula with volume ratios

```

Listing 206. C2 Implementation: DeepSeek R1 (Score: 1.0)

```

3112  # Paper2Code: Missing implementation
3113  def compute_barycentrics(self, query_point,
3114      simplex_index):
3115      transform_matrix = self.delaunay.transform[
3116          simplex_index]
3117      T = transform_matrix[:ndim, :]
3118      offset = transform_matrix[ndim, :]
3119
3120  # Function incomplete, never returns values
3121
3122  # CRITICAL ERROR: Never computes actual coordinates

```

Listing 207. C2 Implementation: Paper2Code (Score: 0.0)

```

3123  # AutoP2C: No barycentric coordinates
3124  # Uses completely different approach
3125  # MISSING: Core mathematical concept entirely

```

Listing 208. C2 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY uses efficient simplex computation. DeepSeek R1 implements exact volume method. GPT-5 uses valid but less robust approach. Paper2Code and AutoP2C fail completely.

Component C3: Feature Interpolation - All Baselines

The paper specifies interpolating features: $f(\mathbf{q}) = \sum_{i=0}^3 \lambda_i \mathbf{f}_i$

3126
3127
3128
3129
3130
3131

3132
3133
3134

3135
3136
3137
3138
3139
3140
3141
3142
3143
3144
3145
3146
3147
3148
3149

```
1 # \nerify : Efficient vectorized interpolation
2 def _gather_grid(self, idx):
3     ix, iy, iz = idx.unbind(-1)
4     grid_flat = self.features.view(-1, self.
5         vertex_feature_dim)
5     lin = ix * (self.ny * self.nx) + iy * self.nx +
6         iz
6     return grid_flat[lin]
7 # Gather and interpolate
8 f0 = self._gather_grid(v0)
9 f1 = self._gather_grid(v1)
10 f = (f0 * w0.unsqueeze(-1) + f1 * w1.unsqueeze(-1) +
11      f2 * w2.unsqueeze(-1) + f3 * w3.unsqueeze(-1))
12 # CORRECT: Vectorized for efficiency
```

Listing 209. C3 Implementation: NERIFY (Score: 1.0)

3151
3152
3153
3154
3155
3156
3157
3158
3159
3160
3161
3162

```
1 # GPT-5: Correct but inefficient
2 def interpolate_feature(p, verts, feats, tets):
3     centroids = verts[tets].mean(dim=1)
4     idx = torch.argmin(torch.norm(centroids - p, dim
5         =1))
5     tet = tets[idx]
6     lambdas = barycentric_coords(p, verts[tet])
7     f = (lambdas[:, None] * feats[tet]).sum(0)
8     return f
9 # ISSUE: O(n) search instead of O(log n)
```

Listing 210. C3 Implementation: GPT-5 (Score: 0.8)

3163
3164
3165
3166
3167
3168
3169
3170
3171
3172
3173
3174
3175
3176

```
1 # DeepSeek R1: Correct interpolation
2 def interpolate_features(self, query_points):
3     for i in range(batch_size):
4         if tetra_indices[i] >= 0:
5             tetra_vertices = self.tetrahedra.
6             simplices[tetra_indices[i]]
6             vertex_features = self.features[
7                 tetra_vertices]
7             interpolated_features[i] = torch.sum(
8                 barycentric[i].unsqueeze(1) *
9                 vertex_features, dim=0)
9 # CORRECT: Proper weighted sum
```

Listing 211. C3 Implementation: DeepSeek R1 (Score: 1.0)

3177
3178
3179
3180
3181
3182
3183
3184
3185
3186
3187
3188
3189

```
1 # Paper2Code: Wrong method
2 def batch_barycentric_interpolation(self,
3     query_points):
4     for i in range(query_points.shape[0]):
5         try:
6             feat = self.barycentric_interpolation(pt,
7                 device=device)
7             except ValueError:
8                 feat = torch.zeros(self.vertex_features.
8                     shape[1])
8 # ERROR: Returns zeros instead of interpolating
```

Listing 212. C3 Implementation: Paper2Code (Score: 0.3)

3190
3191
3192
3193

```
1 # AutoP2C: No feature interpolation
2 # MISSING: Core rendering component
```

Listing 213. C3 Implementation: AutoP2C (Score: 0.0)

3194
3195
3196
3197
3198
3199
3200

Analysis: NERIFY provides vectorized implementation. GPT-5 correct but inefficient. DeepSeek R1 correct with loops. Paper2Code fails. AutoP2C missing.

Component C4: Tetrahedral Grid Structure - All Baselines

The paper requires a data structure to manage vertices, features, and connectivity.

3201
3202
3203
3204
3205
3206
3207
3208
3209
3210
3211
3212
3213
3214
3215
3216
3217
3218
3219

```
1 # \nerify : Complete TetraNerfField class
2 class TetraNerfField(Field):
3     def __init__(self, aabb, grid_resolution=32,
4         vertex_feature_dim=64,
5         bottleneck_dim=32):
6         super().__init__()
7         self.register_buffer("aabb", aabb)
8         self.nx = self.ny = self.nz = grid_resolution
9         self.features = nn.Parameter(torch.empty(
9             self.nx, self.ny, self.nz,
9             vertex_feature_dim))
10        nn.init.uniform_(self.features, a=-1e-4, b=1e
10        -4)
11        self.base_mlp = MLP(in_dim=vertex_feature_dim
12        ,
12        num_layers=3, layer_width
12        =128)
13 # CORRECT: Complete field with Nerfstudio integration
```

Listing 214. C4 Implementation: NERIFY (Score: 1.0)

3221
3222
3223
3224
3225
3226
3227
3228
3229

```
1 # GPT-5: Functions only, no proper class
2 def build_tetrahedra(points):
3     tri = Delaunay(points)
4     verts = torch.tensor(points)
5     tets = torch.tensor(tri.simplices)
6     def interpolate_feature(p, verts, feats, tets):
7         # Separate functions
8     # ISSUE: No unified data structure
```

Listing 215. C4 Implementation: GPT-5 (Score: 0.5)

3231
3232
3233
3234
3235
3236
3237
3238
3239
3240
3241
3242
3243
3244

```
1 # DeepSeek R1: Has structure but incomplete
2 class TetrahedralGrid:
3     def __init__(self, points):
4         self.points = points.cpu().numpy()
5         self.tetrahedra = Delaunay(self.points)
6         self.features = nn.Parameter(
6             torch.randn(len(points), 64) * 1e-4)
8     def find_tetrahedron(self, query_points):
9         query_np = query_points.cpu().numpy()
10        tetra_indices = self.tetrahedra.find_simplex(
10            query_np)
11 # CORRECT: Good structure but missing integration
```

Listing 216. C4 Implementation: DeepSeek R1 (Score: 0.8)

3245
3246
3247
3248
3249
3250
3251
3252
3253

```
1 # Paper2Code: Wrong structure
2 class TetrahedraField(nn.Module):
3     def __init__(self, point_cloud,
4         vertex_feature_dim=64):
4         self.points = point_cloud
5         self.delaunay = None # Never initialized
6     # CRITICAL ERROR: Empty module
```

Listing 217. C4 Implementation: Paper2Code (Score: 0.0)

3254
3255
3256

```
1 # AutoP2C: No grid structure
2 # MISSING: Core data structure
```

Listing 218. C4 Implementation: AutoP2C (Score: 0.0)

3258
3259
3260
3261
3262
3263

Analysis: NERIFY provides complete structure. GPT-5 lacks organization. DeepSeek R1 has structure but incomplete. Paper2Code and AutoP2C fail.

Component C5: Adaptive Resolution - All Baselines

The paper achieves adaptive resolution through Delaunay triangulation concentrating tetrahedra near surfaces.

```

3264
3265 1 # \nerify\ : Implicit adaptive via simplex
3266      decomposition
3267 2 def _tetrahedral_interpolate(self, positions_norm):
3268      # Scale to grid coordinates
3269      scales = torch.tensor([self.nx-1, self.ny-1, self
3270      .nz-1])
3271      gpos = positions_norm * scales
3272      base = torch.floor(gpos).long()
3273      # Smaller tetrahedra near surfaces due to
3274      # decomposition
3275      base[:, 0].clamp_(0, self.nx - 2)
3276      base[:, 1].clamp_(0, self.ny - 2)
3277      base[:, 2].clamp_(0, self.nz - 2)
3278 11 # CORRECT: Achieves adaptivity through decomposition

```

Listing 219. C5 Implementation: NERIFY (Score: 1.0)

```

3280
3281 1 # GPT-5: No adaptive handling
3282 2 t_vals = torch.linspace(0.0, 1.0, n_samples)
3283 3 # Uniform sampling without adaptation
3284 4 # MISSING: Adaptive resolution concept

```

Listing 220. C5 Implementation: GPT-5 (Score: 0.2)

```

3286
3287 1 # DeepSeek R1: Structure supports but not implemented
3288 2 self.tetrahedra = Delaunay(self.points)
3289 3 # Natural adaptivity from Delaunay but not utilized
3290 4 # ISSUE: Has potential but not exploited

```

Listing 221. C5 Implementation: DeepSeek R1 (Score: 0.5)

```

3292
3293 1 # Paper2Code: Opposite of adaptive
3294 2 # Uses uniform grid
3295 3 # CRITICAL ERROR: Misses key innovation

```

Listing 222. C5 Implementation: Paper2Code (Score: 0.0)

```

3297
3298 1 # AutoP2C: Random points
3299 2 random_points = np.random.normal(
3300 3     loc=self.point_cloud.mean(axis=0),
3301 4     scale=self.point_cloud.std(axis=0))
3302 5 # MISSING: Any spatial structure

```

Listing 223. C5 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERIFY achieves adaptive resolution. DeepSeek R1 has structure but doesn't exploit. Others miss innovation.

Component C6: Hierarchical Sampling - All Baselines

The paper uses coarse-to-fine sampling with importance weighting.

```

3304
3305 1 # \nerify\ : Inherits from Nerfacto's samplers
3306 2 class TetraNerfModel(NerfactoModel):
3307 3     def populate_modules(self):
3308         Model.populate_modules(self)
3309         super().populate_modules() # Gets proposal
3310         networks
3311         # Uses Nerfacto's production hierarchical
3312         # sampling
3313         self.field = TetraNerfField(aabb=self.
3314             scene_box.aabb)
3315 8 # CORRECT: Production-quality hierarchical sampling

```

Listing 224. C6 Implementation: NERIFY (Score: 1.0)

```

3323
3324 1 # GPT-5: Basic implementation
3325 2 def hierarchical_sampling(ray_o, ray_d, model,
3326      n_coarse=32, n_fine=32):
3327 3     t_coarse = torch.linspace(0, 1, n_coarse)

```

```

3328 4     # Fine sampling proportional to weights
3329 5     cdf = torch.cumsum(w_norm.squeeze(), 0)
3330 6     u = torch.rand(n_fine, device=device)
3331 7     t_fine = torch.interp(u, cdf, t_coarse)
3332 8 # ISSUE: Simplified, missing details

```

Listing 225. C6 Implementation: GPT-5 (Score: 0.6)

```

3334
3335 1 # DeepSeek R1: Complete but inefficient
3336 2 def volume_render(self, rays_o, rays_d, near, far):
3337 3     # Coarse sampling
3338 4     t_vals_coarse = torch.linspace(near, far, self.
3339         config.num_samples_coarse)
3340 5     weights = self.compute_weights(density_coarse,
3341         t_vals_coarse)
3342 6     # Fine sampling
3343 7     t_vals_fine = self.sample_fine(t_vals_coarse,
3344         weights)
3345 8 # CORRECT: Right idea but loop-based

```

Listing 226. C6 Implementation: DeepSeek R1 (Score: 0.8)

```

3347
3348 1 # Paper2Code: Only mentions
3349 2 def stratified_sampling(self, nears, fars,
3350      num_samples):
3351 3     t_bins = torch.linspace(0.0, 1.0, steps=
3352      num_samples + 1)
3353 4     # Only coarse, no fine
3354 5 # ERROR: Missing importance sampling

```

Listing 227. C6 Implementation: Paper2Code (Score: 0.2)

```

3356
3357 1 # AutoP2C: Random sampling
3358 2 def generate_rays(self):
3359 3     for pose in self.camera_poses:
3360 4         rays = self._generate_rays_for_pose(pose)
3361 5 # MISSING: Hierarchical concept

```

Listing 228. C6 Implementation: AutoP2C (Score: 0.0)

Analysis: NERIFY leverages production samplers. GPT-5 and DeepSeek R1 attempt with varying success. Paper2Code and AutoP2C fail.

Component C7: Volume Rendering - All Baselines

The paper requires volume rendering with tetrahedral field queries.

```

3366
3367 1 # \nerify\ : Complete volume rendering
3368 2 def get_outputs(self, ray_samples, density_embedding=
3369      None):
3370 3     dirs = get_normalized_directions(ray_samples.
3371         frustums.directions)
3372 4     dirs_flat = dirs.view(-1, 3)
3373 5     denc = self.direction_encoding(dirs_flat)
3374 6     h = torch.cat([denc, density_embedding.view(-1,
3375         self.bottleneck_dim)], dim=-1)
3376 7     rgb = self.color_head(h).view(*dirs.shape[:-1],
3377         -1).to(dirs)
3378 8     outputs[FieldHeadNames.RGB] = rgb
3379 9     return outputs
3380 10 # CORRECT: Full integration with Nerfstudio rendering

```

Listing 229. C7 Implementation: NERIFY (Score: 1.0)

```

3385
3386 1 # GPT-5: Basic volume rendering
3387 2 def volume_render(ray_o, ray_d, model, verts, feats,
3388      tets, n_samples=64):
3389 3     t_vals = torch.linspace(0.0, 1.0, n_samples)
3390 4     pts = ray_o[None, :] + ray_d[None, :] * t_vals[:, :
3391      None]
3392 5     weights = sigmas * T * deltas[:, None]
3393 6     rgb = (weights * colors).sum(0)

```

3394
3395

```
7     return rgb
8 # CORRECT: Basic but functional
```

Listing 230. C7 Implementation: GPT-5 (Score: 0.7)

3397
3398
3399
3400
3401
3402
3403
3404
3405

```
1 # DeepSeek R1: Volume rendering with issues
2 def compute_weights(self, density, t_vals):
3     delta = t_vals[..., 1:] - t_vals[..., :-1]
4     alpha = 1 - torch.exp(-density * delta)
5     transmittance = torch.cumprod(1 - alpha + 1e-10,
6         dim=-1)
6     weights = alpha * transmittance
7 # ISSUE: Incorrect transmittance computation
```

Listing 231. C7 Implementation: DeepSeek R1 (Score: 0.6)

3407
3408
3409
3410
3411

```
1 # Paper2Code: Wrong rendering
2 def volume_rendering(self, t_vals, sigma, rgb):
3     # Missing proper implementation
4 # CRITICAL ERROR: No actual rendering
```

Listing 232. C7 Implementation: Paper2Code (Score: 0.0)

3413
3414
3415

```
1 # AutoP2C: No volume rendering
2 # MISSING: Core rendering pipeline
```

Listing 233. C7 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY has full integration. GPT-5 basic but functional. DeepSeek R1 has errors. Paper2Code and AutoP2C missing.

Component C8: Tetrahedron Volume - All Baselines
The paper computes tetrahedron volume: $V = \frac{1}{6} |(\mathbf{a} - \mathbf{d}) \cdot ((\mathbf{b} - \mathbf{d}) \times (\mathbf{c} - \mathbf{d}))|$

3420
3421
3422
3423
3424
3425
3426
3427
3428
3429
3430
3431

```
1 # \nerify\ : Implicit in simplex decomposition
2 # Volume computation implicit in barycentric weights
3 w0 = 1.0 - u_sorted[:, 2] # Volume ratio
4 w1 = u_sorted[:, 2] - u_sorted[:, 1]
5 w2 = u_sorted[:, 1] - u_sorted[:, 0]
6 w3 = u_sorted[:, 0]
7 # CORRECT: Efficient volume-based weights
```

Listing 234. C8 Implementation: NERFIFY (Score: 1.0)

3432
3433
3434
3435
3436

```
1 # GPT-5: No explicit volume computation
2 # Uses linear algebra instead
3 # ISSUE: Missing volume formula
```

Listing 235. C8 Implementation: GPT-5 (Score: 0.3)

3437
3438
3439
3440
3441
3442
3443
3444

```
1 # DeepSeek R1: Exact formula
2 def tetrahedron_volume(self, a, b, c, d):
3     return torch.abs(torch.dot(a - d,
4                                torch.cross(b - d, c - d))) /
6.0
5 # CORRECT: Exact paper formula
```

Listing 236. C8 Implementation: DeepSeek R1 (Score: 1.0)

3445
3446
3447

```
1 # Paper2Code: No volume computation
2 # MISSING: Mathematical component
```

Listing 237. C8 Implementation: Paper2Code (Score: 0.0)

3449
3450
3451

```
1 # AutoP2C: Missing entirely
2 # MISSING: Core formula
```

Listing 238. C8 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY implicit in weights. DeepSeek R1 exact formula. GPT-5 missing. Paper2Code and AutoP2C absent.

Component C9: Nerfstudio Integration - All Baselines

The paper implementation requires Nerfstudio plugin structure for training.

```
1 # \nerify\ : Complete Nerfstudio plugin
2 tetra_nerf = MethodSpecification(
3     config=TrainerConfig(
4         method_name="tetra-nerf",
5         pipeline=TetraNerfPipelineConfig(
6             datamanager=TetraNerfDataManagerConfig(),
7             model=TetraNerfModelConfig(),
8             optimizers={
9                 "fields": {"optimizer": RAdamOptimizerConfig(lr=1e-3)}},
10            description="Tetra-NeRF: Nerfacto-based with
11            tetrahedral field")
11 # CORRECT: Complete plugin specification
```

Listing 239. C9 Implementation: NERFIFY (Score: 1.0)

```
1 # GPT-5: No Nerfstudio integration
2 if __name__ == "__main__":
3     # Standalone demo only
4     model, verts, feats, tets = train_tetra_nerf(pts,
5                                                 rays_o, rays_d)
5 # MISSING: Plugin structure
```

Listing 240. C9 Implementation: GPT-5 (Score: 0.0)

```
1 # DeepSeek R1: No framework integration
2 class TetraNerfTrainer:
3     def __init__(self, model, config):
4         self.model = model
5         self.optimizer = torch.optim.Adam(model.
6                                         parameters())
6 # MISSING: Nerfstudio plugin
```

Listing 241. C9 Implementation: DeepSeek R1 (Score: 0.0)

```
1 # Paper2Code: Custom training only
2 class Trainer:
3     def train(self):
4         for iter_num in range(1, self.iterations + 1):
5             # MISSING: Framework integration
```

Listing 242. C9 Implementation: Paper2Code (Score: 0.0)

```
1 # AutoP2C: No framework awareness
2 # MISSING: Plugin structure
```

Listing 243. C9 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERFIFY has Nerfstudio integration. All baselines lack plugin structure, making them untrainable in framework.

Component C10: RAdam Optimizer - All Baselines

The paper specifies RAdam optimizer with exponential decay from 1e-3 to 1e-4.

```
1 # \nerify\ : Exact paper specification
2 "fields": {
3     "optimizer": RAdamOptimizerConfig(lr=1e-3, eps=1e-
8),
4     "scheduler": ExponentialDecaySchedulerConfig(
5         lr_final=1e-4, max_steps=300000),
```

3516
3518

```
6 }
7 # CORRECT: RAdam with proper decay
```

Listing 244. C10 Implementation: NERFIFY (Score: 1.0)

3519
3520
3521
3522
3523
3524
3525
3526

```
1 # GPT-5: Correct optimizer
2 self.optimizer = torch.optim.RAdam(
3     list(model.parameters()) + [feats], lr=1e-3)
4 self.scheduler = torch.optim.lr_scheduler.
5     ExponentialLR(
6         self.optimizer, gamma=0.999)
7 # CORRECT: RAdam with decay
```

Listing 245. C10 Implementation: GPT-5 (Score: 1.0)

3528
3529
3530
3531
3532
3533
3534
3535
3536

```
1 # DeepSeek R1: Wrong optimizer
2 self.optimizer = torch.optim.Adam(model.parameters(),
3     lr=config.learning_rate)
4 self.scheduler = torch.optim.lr_scheduler.
5     ExponentialLR(
6         self.optimizer, gamma=0.1)
7 # ERROR: Uses Adam instead of RAdam
```

Listing 246. C10 Implementation: DeepSeek R1 (Score: 0.8)

3537
3538
3539
3540
3541
3542
3543
3544
3545
3546

```
1 # Paper2Code: Correct optimizer setup
2 self.optimizer = torch.optim.RAdam(self.model.
3     parameters(), lr=self.lr_initial)
4 def adjust_learning_rate(self, current_iter):
5     decay_factor = (self.lr_final / self.lr_initial)
      **
      (current_iter / self.
6     lr_decay_steps)
7 # CORRECT: RAdam with exponential decay
```

Listing 247. C10 Implementation: Paper2Code (Score: 1.0)

3548
3549
3550
3551
3552
3553
3554

```
1 # AutoP2C: Mentions RAdam
2 # Paper uses RAdam with exponential LR decay
3 # But implementation missing
4 # ISSUE: No actual optimizer code
```

Listing 248. C10 Implementation: AutoP2C (Score: 0.8)

3554
3555
3556

Analysis: NERFIFY, GPT-5, and Paper2Code correctly implement RAdam. DeepSeek R1 uses Adam. AutoP2C mentions but doesn't implement.

3557

3.5.6. Scoring Analysis

Table 48. Tetra-NeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	Weighted Avg _{LLM}
NERFIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.00
GPT-5	1.0	0.4	0.8	0.5	0.2	0.6	0.7	0.3	0.0	1.0	0.58
DeepSeek R1	1.0	1.0	1.0	0.8	0.5	0.8	0.6	1.0	0.0	0.8	0.72
Paper2Code	0.2	0.0	0.3	0.0	0.0	0.2	0.0	0.0	0.0	1.0	0.22
AutoP2C	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.8	0.08

3558

3.5.7. Why Baselines Fail Despite Component Scores

GPT-5 (Score: 0.58 components, 0% trainable)

- Strengths:** Implements core tetrahedral concepts correctly, proper Delaunay triangulation
- Fatal Issues:**
 - No Nerfstudio integration whatsoever
 - Inefficient O(n) tetrahedron search instead of O(log n)

- Missing training loop and data loading 3565
- Result:** Code runs in isolation but cannot be trained within Nerfstudio framework 3566
- DeepSeek R1 (Score: 0.72 components, 0% trainable)** 3567
- Strengths:** Mathematically correct volume-based barycentric coordinates, proper feature interpolation 3569
- Fatal Issues:**
 - No Nerfstudio plugin structure 3570
 - Loop-based processing causes memory overflow on real datasets 3571
 - Incomplete volume rendering implementation 3572
- Result:** Theoretically correct but practically unusable for training 3576
- Paper2Code (Score: 0.22 components, 0% trainable)** 3577
- Strengths:** Attempts multi-file organization, includes optimizer setup 3579
- Fatal Issues:**
 - Fundamental misunderstanding: implements voxel grid instead of tetrahedra 3582
 - Missing barycentric coordinate computation entirely 3583
 - PyTorch errors including device mismatches 3584
- Result:** Implements a completely different method than the paper specifies 3586
- AutoP2C (Score: 0.08 components, 0% trainable)** 3588
- Strengths:** Includes COLMAP data loading 3589
- Fatal Issues:**
 - Missing all core tetrahedral concepts 3591
 - No Delaunay triangulation or barycentric coordinates 3592
 - Non-functional Python with import errors 3593
- Result:** Non-functional code that fails at import stage 3594

3.5.8. Hyperparameter Fidelity

Table 49. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5	R1	P2C	AutoP2C
Feature dimension	64	✓	✓	✓	✗	✗
Hidden dimension	128	✓	✓	✓	✗	✗
MLP layers	3	✓	✓	✓	✗	✗
Coarse samples	64	✓	✗	✓	✗	—
Fine samples	128	✓	✗	✓	✗	—
Learning rate	1e-3	✓	✓	✓	✗	—
Optimizer	RAdam	✓	✗	✗	✗	—
W Score	—	1.00	0.60	0.70	0.20	0.00

3.5.9. Conclusion

All baselines attempt to implement Tetra-NeRF with varying levels of sophistication. GPT-5 and DeepSeek R1 demonstrate understanding of core tetrahedral concepts, correctly implementing Delaunay triangulation and barycentric interpolation. Paper2Code fundamentally misunderstands the paper, implementing a voxel-based approach instead of tetrahedra. AutoP2C fails to implement any meaningful components beyond basic data loading. Despite GPT-5 achieving 58% and DeepSeek R1 achieving 72% component coverage, 3596

359735983599360036013602360336043605

3606 neither produces trainable code due to missing Nerfstudio
 3607 integration, inefficient algorithms, and incomplete training
 3608 pipelines. Only NERFIFY produces immediately trainable
 3609 code as a complete Nerfstudio plugin with 100% compo-
 3610 nent fidelity. The 0% trainable rate for all baselines versus
 3611 100% for NERFIFY demonstrates the critical importance of
 3612 domain-specific understanding and framework integration in
 3613 neural rendering research. ““

3614 3.6. E-NeRF: Neural Radiance Fields from a Mov- 3615 ing Event Camera

3616 3.6.1. Paper Overview

3617 E-NeRF introduces a groundbreaking approach to neural
 3618 radiance field reconstruction using event cameras, which
 3619 capture pixel-level brightness changes rather than absolute
 3620 intensity. The method replaces traditional RGB photometric
 3621 losses with event-generation constraints, enabling NeRF
 3622 reconstruction from asynchronous event streams. This is par-
 3623 ticularly valuable for scenarios with fast motion, low light,
 3624 or high dynamic range where conventional cameras fail. The
 3625 paper demonstrates that event supervision alone can produce
 3626 high-quality 3D reconstructions, with optional RGB frame
 3627 integration for enhanced color fidelity.

3628 3.6.2. Implementation Overview

Table 50. E-NeRF Implementation Summary Across All Baselines

Aspect	NERFIFY	GPT-5	DeepSeek R1	Paper2Code	AutoP2C
Lines of Code	487	245	389	1650	156
File Organization	Plugin	Single	Single	Multi-file	Multi-file
Event Loss (Eq. 3)	✓	✓	✓	Partial	✗
Linlog Mapping	✓	✓	✓	✓	✗
No-Event Loss	✓	Partial	✓	✗	✗
Framework Integration	✓	✗	✗	✗	✗
Trainable	✓	✗	✗	✗	✗

Note: NERFIFY always produces trainable code as a complete Nerfstudio
 plugin

3629 3.6.3. Novel Components

Table 51. Novel Components in E-NeRF with Importance Weights

ID	Component	Weight w_i
C1	Event-based loss function (Eq. 3)	0.25
C2	Linlog brightness mapping (Eq. 2)	0.20
C3	Normalized event loss (Eq. 4)	0.15
C4	No-event loss (Eq. 5)	0.10
C5	Event + RGB combined loss (Eq. 6)	0.10
C6	Event pair sampling strategy	0.08
C7	Camera pose interpolation (slerp/cubic)	0.07
C8	Hash-based encoding (Instant-NGP)	0.05

Table 52. E-NeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	1.00	0.00	0.00	0.95	1.00
GPT-5	0.50	0.25	0.25	0.75	0.60
DeepSeek R1	0.63	0.25	0.13	0.80	0.72
Paper2Code	0.38	0.25	0.38	0.60	0.48
AutoP2C	0.00	0.13	0.88	0.00	0.05

3614 3.6.4. Quantitative Metrics

3615 3.6.5. Component-by-Component Analysis

3616 Component C1: Event-based Loss Function - All Baselines

3617 The paper’s core innovation is the event loss enforc-
 3618 ing brightness change constraints: $\mathcal{L}_{\text{evs}}(\Theta) = \|\Delta\hat{\mathbf{L}}(\Theta) -$
 3619 $\Delta\mathbf{L}(\Theta)\|_2^2$ where $\Delta L_k = p_k C$ with contrast threshold
 3620 $C = 0.2$.

```
1 # \nerfify\ : Full integration with event batch
  handling
2 def _event_losses(self, batch):
3     rb_t0 = self._bundle_from_batch_prefix(batch, "event_t0")
4     rb_t1 = self._bundle_from_batch_prefix(batch, "event_t1")
5
6     I0 = self._render_intensity(rb_t0)
7     I1 = self._render_intensity(rb_t1)
8     L0 = _linlog(I0, B)
9     L1 = _linlog(I1, B)
10    delta_hat = (L1 - L0).reshape(-1)
11
12    delta_gt = (C * batch["event_polarity"]).reshape
13        (-1)
14    ev_loss = F.mse_loss(delta_hat, delta_gt)
15    # CORRECT: Proper batch structure and loss
16        computation
```

Listing 249. C1 Implementation: NERFIFY (Score: 1.0)

```
1 # GPT-5: Mathematically correct but no framework
  integration
2 def event_loss(deltaL_hat, deltaL_true, normalized=
  False):
3     if normalized:
4         deltaL_hat = deltaL_hat / (torch.norm(
5             deltaL_hat) + 1e-8)
6         deltaL_true = deltaL_true / (torch.norm(
7             deltaL_true) + 1e-8)
8     return F.mse_loss(deltaL_hat, deltaL_true)
9 # ISSUE: No batch handling, standalone function
```

Listing 250. C1 Implementation: GPT-5 (Score: 0.8)

```
1 # DeepSeek R1: Comprehensive but missing integration
2 def generate_event_loss(self, pred_log_brightness1,
3                         pred_log_brightness2,
4                         measured_polarity, C=0.2):
5     pred_delta_L = pred_log_brightness2 -
6         pred_log_brightness1
7     target_delta_L = measured_polarity * C
8     event_loss = F.mse_loss(pred_delta_L,
9                           target_delta_L)
10    return event_loss
11 # CORRECT: Formula implementation
12 # MISSING: Training loop integration
```

3685

Listing 251. C1 Implementation: DeepSeek R1 (Score: 0.9)

3686
3687
3688
3689
3690
3691
3692
3693
3694
3695

```

1 # Paper2Code: Basic structure without proper batching
2 if self.event_loss_type == "standard":
3     event_loss = torch.mean((predicted_diff -
4         event_target_diff_tensor) ** 2)
5 elif self.event_loss_type == "normalized":
6     event_loss = Utils.normalized_event_loss(
7         predicted_diff, event_target_diff_tensor)
8 # ERROR: Missing batch structure, incomplete
9     integration

```

Listing 252. C1 Implementation: Paper2Code (Score: 0.5)

3697
3698
3699
3700

```

1 # AutoP2C: No implementation
2 # File only contains dataset placeholders
3 # CRITICAL ERROR: Event loss not implemented at all

```

Listing 253. C1 Implementation: AutoP2C (Score: 0.0)

3702
3703
3704
3705
3706
3707

Analysis: NERFIFY correctly implements the event loss with full Nerfstudio integration. GPT-5 and DeepSeek R1 have mathematically correct formulas but lack framework integration. Paper2Code attempts the loss but misses critical batch handling. AutoP2C completely lacks any loss implementation.

3708
3709

Component C2: Linlog Brightness Mapping - All Baselines

3710
3711

The paper specifies: $L(\mathbf{u}) = I(\mathbf{u}) \cdot \ln(B)/B$ if $I < B$, else $\ln(I)$ with $B = 20$.

3712
3713
3714
3715
3716
3717
3718
3719
3720
3721
3722
3723
3724

```

1 # \nerify\ : Exact paper formula with proper epsilon
2     handling
3 def _linlog(intensity: Tensor, B: float) -> Tensor:
4     eps = 1e-6
5     lnB_over_B = float(torch.log(torch.tensor(B)).item() / B)
6     I = torch.clamp(intensity, min=eps)
7     linear = I * lnB_over_B
8     logv = torch.log(I)
9     return torch.where(I < B, linear, logv)
9 # CORRECT: Exact implementation from paper

```

Listing 254. C2 Implementation: NERFIFY (Score: 1.0)

3725
3726
3727
3728
3729
3730
3731
3732
3733
3734
3735
3736
3737
3738
3739
3740
3741
3742
3743
3744

```

1 # GPT-5: Correct implementation with masking
2 def linlog(I, B=20.0):
3     mask = (I < B).float()
4     return mask * (I * torch.log(torch.tensor(B)) / B
5     ) + \
6         (1 - mask) * torch.log(I + 1e-6)
6 # CORRECT: Proper branching logic

```

Listing 255. C2 Implementation: GPT-5 (Score: 1.0)

3745
3746
3747
3748
3749
3750
3751
3752
3753
3754
3755
3756
3757
3758
3759
3760
3761

```

1 # DeepSeek R1: Correct with explicit masking
2 def linlog(self, I):
3     linear_mask = I < self.B
4     log_mask = ~linear_mask
5     L = torch.zeros_like(I)
6     L[linear_mask] = I[linear_mask] * torch.log(torch
7         .tensor(self.B)) / self.B
8     L[log_mask] = torch.log(I[log_mask])
9     return L
9 # CORRECT: Accurate implementation

```

Listing 256. C2 Implementation: DeepSeek R1 (Score: 1.0)

1 # Paper2Code: Verbose but correct	3746
2 def lin_log_mapping(self, intensity: torch.Tensor) ->	3747
3 torch.Tensor:	3748
4 B_tensor = torch.tensor(B, device=intensity.	3749
5 device)	3750
6 return torch.where(3751
7 intensity < B_tensor,	3752
8 intensity * torch.log(B_tensor) / B_tensor,	3753
9 torch.log(intensity + 1e-6)	3754
9)	3755
9 # CORRECT: Proper implementation	3756

Listing 257. C2 Implementation: Paper2Code (Score: 1.0)

1 # AutoP2C: Not implemented	3759
2 # CRITICAL ERROR: No linlog mapping found in codebase	3760

Listing 258. C2 Implementation: AutoP2C (Score: 0.0)

Analysis: All baselines except AutoP2C correctly implement linlog mapping. This relatively straightforward component shows that baselines can handle individual mathematical formulas but struggle with system integration.

Component C3: Normalized Event Loss - All Baselines

For unknown contrast threshold: $\mathcal{L}_{\text{norm}} = \left\| \frac{\Delta \hat{L}}{\|\Delta \hat{L}\|_2} - \frac{\Delta L}{\|\Delta L\|_2} \right\|_2^2$

1 # \nerify\ : Config-switchable normalized variant	3771
2 if self.config.use_normalized_event_loss:	3772
3 eps = 1e-8	3773
4 nhat = delta_hat / (torch.linalg.norm(delta_hat)	3774
5 + eps)	3775
6 ngt = delta_gt / (torch.linalg.norm(delta_gt) +	3776
7 eps)	3777
6 ev_loss = F.mse_loss(nhat, ngt)	3778
7 else:	3779
8 ev_loss = F.mse_loss(delta_hat, delta_gt)	3780
9 # CORRECT: Proper L2 normalization before MSE	3781

Listing 259. C3 Implementation: NERFIFY (Score: 1.0)

1 # GPT-5: Has normalized variant as parameter	3784
2 def event_loss(deltaL_hat, deltaL_true, normalized=	3785
3 False):	3786
3 if normalized:	3787
4 deltaL_hat = deltaL_hat / (torch.norm(3788
4 deltaL_hat) + 1e-8)	3789
5 deltaL_true = deltaL_true / (torch.norm(3790
5 deltaL_true) + 1e-8)	3791
6 return F.mse_loss(deltaL_hat, deltaL_true)	3792
7 # CORRECT: But not integrated with training	3793

Listing 260. C3 Implementation: GPT-5 (Score: 0.8)

1 # DeepSeek R1: Separate function for normalized loss	3796
2 def normalized_event_loss(self, pred_delta_L,	3797
2 target_delta_L):	3798
3 pred_norm = pred_delta_L / (torch.norm(3799
3 pred_delta_L, dim=-1, keepdim=True) + 1e-8)	3800
4 target_norm = target_delta_L / (torch.norm(3801
4 target_delta_L, dim=-1, keepdim=True) + 1e-8)	3802
5 return F.mse_loss(pred_norm, target_norm)	3803
6 # CORRECT: Proper implementation	3804

Listing 261. C3 Implementation: DeepSeek R1 (Score: 1.0)

3807
3808
3809
3810

```
1 # Paper2Code: No normalized variant
2 # Only has standard event loss
3 # CRITICAL ERROR: Missing normalized loss entirely
```

Listing 262. C3 Implementation: Paper2Code (Score: 0.0)

3812
3813
3814

```
1 # AutoP2C: Not implemented
2 # CRITICAL ERROR: No loss functions at all
```

Listing 263. C3 Implementation: AutoP2C (Score: 0.0)

3816
3817
3818
3819

Analysis: NERFIFY, GPT-5, and DeepSeek R1 correctly implement the normalized variant crucial for real-world data. Paper2Code and AutoP2C completely miss this component.

Component C4: No-Event Loss - All Baselines

3820

Penalizes brightness changes in areas without events:

$$\mathcal{L}_{\text{noevs}} = \sum_k \text{relu}(|\hat{L}_k - \hat{L}_{k-1}| - C)$$

3821
3822
3823
3824
3825
3826
3827
3828
3829
3830
3831
3832
3833
3834
3835
3836
3837
3838
3839
3840
3841
3842
3843

```
1 # \nerfify\ : Complete no-event handling
2 rb_n0 = self._bundle_from_batch_prefix(batch, "
    noevent_t0")
3 rb_n1 = self._bundle_from_batch_prefix(batch, "
    noevent_t1")
4 if rb_n0 is not None and rb_n1 is not None:
5     I0 = self._render_intensity(rb_n0)
6     I1 = self._render_intensity(rb_n1)
7     L0 = _linlog(I0, B)
8     L1 = _linlog(I1, B)
9     diff = torch.abs(L1 - L0) - C
10    noev_loss = torch.clamp(diff, min=0.0).mean()
11 # CORRECT: Full pipeline from sampling to loss
```

Listing 264. C4 Implementation: NERFIFY (Score: 1.0)

3837
3838
3839
3840
3841
3842
3843

```
1 # GPT-5: Formula correct but no sampling
2 def no_event_loss(Lhat_k, Lhat_kml, C):
3     return F.relu(torch.abs(Lhat_k - Lhat_kml) - C).
        mean()
4 # ERROR: Missing no-event location identification
```

Listing 265. C4 Implementation: GPT-5 (Score: 0.4)

3844
3845
3846
3847
3848
3849
3850
3851
3852
3853
3854
3855
3856

```
1 # DeepSeek R1: Class-based implementation
2 class NoEventLoss:
3     def __init__(self, C=0.2):
4         self.C = C
5     def __call__(self, pred_log_brightness1,
6                 pred_log_brightness2):
7         brightness_diff = torch.abs(
8             pred_log_brightness2 - pred_log_brightness1)
9         return F.relu(brightness_diff - self.C).mean()
10    ()
11 # CORRECT: Formula but missing sampling logic
```

Listing 266. C4 Implementation: DeepSeek R1 (Score: 0.8)

3857
3858
3859
3860

```
1 # Paper2Code: Not implemented
2 # Has no-event sampling but no loss computation
3 # CRITICAL ERROR: Missing no-event loss
```

Listing 267. C4 Implementation: Paper2Code (Score: 0.0)

3862
3863
3864

```
1 # AutoP2C: Not implemented
2 # CRITICAL ERROR: No loss implementation
```

Listing 268. C4 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERFIFY fully implements no-event loss with proper sampling. GPT-5 and DeepSeek R1 have correct formulas but miss the critical sampling component.

Component C5: Event + RGB Combined Loss - All Baselines

Combines supervision: $\mathcal{L} = \mathcal{L}_{\text{evs}} + \lambda_{\text{rgb}} \mathcal{L}_{\text{rgb}}$

```
1 # \nerfify\ : Integrated with Nerfacto pipeline
2 def get_loss_dict(self, outputs, batch, metrics_dict=None):
3     loss_dict = super().get_loss_dict(outputs, batch,
4         metrics_dict)
5     # Scale RGB loss
6     if "rgb_loss" in loss_dict and self.config.
7         lambda_rgb != 1.0:
8         loss_dict["rgb_loss"] = self.config.
9             lambda_rgb * loss_dict["rgb_loss"]
10    # Add event losses
11    ev_loss, noev_loss = self._event_losses(batch)
12    if ev_loss is not None:
13        loss_dict["event_loss"] = ev_loss
14    loss_dict["correct"] = True
15 # CORRECT: Proper loss combination with weights
```

Listing 269. C5 Implementation: NERFIFY (Score: 1.0)

```
1 # GPT-5: Mathematical combination only
2 def combined_loss(Levs, Lrgb=None, lambda_rgb=1.0):
3     if Lrgb is None:
4         return Levs
5     return Levs + lambda_rgb * Lrgb
6 # ERROR: No actual RGB rendering pipeline
```

Listing 270. C5 Implementation: GPT-5 (Score: 0.4)

```
1 # DeepSeek R1: Combined training step
2 def combined_training_step(self, event_batch,
3                             rgb_batch):
4     event_loss = self.event_training_step(event_batch)
5     if event_loss is not None:
6         # Simplified RGB loss computation
7         total_loss = event_loss + self.config.get('
8             lambda_rgb', 1.0) * rgb_batch
9 # ISSUE: Incomplete RGB pipeline
```

Listing 271. C5 Implementation: DeepSeek R1 (Score: 0.6)

```
1 # Paper2Code: Not implemented
2 # CRITICAL ERROR: No combined loss logic
```

Listing 272. C5 Implementation: Paper2Code (Score: 0.0)

```
1 # AutoP2C: Not implemented
2 # CRITICAL ERROR: No loss functions
```

Listing 273. C5 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERFIFY properly integrates event and RGB losses within a working framework. Others either lack RGB pipelines or miss the combination entirely.

Component C6: Event Pair Sampling Strategy - All Baselines

Samples event pairs within temporal windows (1ms average, 60ms window).

3866

3867

3868

3869

3870

3871

3872

3873

3874

3875

3876

3877

3878

3879

3880

3881

3882

3883

3884

3885

3886

3887

3888

3889

3890

3891

3892

3893

3894

3895

3896

3897

3898

3899

3900

3901

3902

3903

3904

3905

3906

3907

3908

3909

3910

3911

3912

3913

3914

3915

3916

3917

3918

```

3926
3927     1 # \nerify\ : Integrated with DataManager
3928     2 def _copy_event_tensors(self, src: Dict, dst: Dict)
3929         -> None:
3930     3     for k, v in src.items():
3931         if any(k.startswith(pref) for pref in
3932             EVENT_PREFIXES):
3933             dst[k] = v.to(self.device) if hasattr(v,
3934             "to") else v
3935     6 # CORRECT: Full event batch handling

```

Listing 274. C6 Implementation: NERIFY (Score: 1.0)

```

3937
3938     1 # GPT-5: Synthetic dummy dataset only
3939     2 class SyntheticEventDataset(Dataset):
3940         3     def __getitem__(self, idx):
3941             4         uk = torch.rand(2)
3942             5         tk, tkl = torch.rand(1), torch.rand(1)
3943             6         pk = random.choice([-1, 1])
3944     7 # ERROR: No real event sampling implementation

```

Listing 275. C6 Implementation: GPT-5 (Score: 0.2)

```

3946
3947     1 # DeepSeek R1: Placeholder function
3948     2 def sample_event_pairs(self, event_batch,
3949         max_time_window=60):
3950         3     # Implementation depends on your event data
3951             structure
3952             4     # This should sample neighboring events or events
3953                 within time windows
3954             5     pass
3955     6 # CRITICAL ERROR: Not implemented

```

Listing 276. C6 Implementation: DeepSeek R1 (Score: 0.3)

```

3957
3958     1 # Paper2Code: Detailed sampling logic
3959     2 def sample_event_pairs(self, n_samples: int) -> List[
3960         Dict[str, Any]]:
3961         3     for _ in range(n_samples):
3962             4         pixel = random.choice(valid_pixels)
3963             5         events_list = self.events_by_pixel[pixel]
3964             6         i = random.randint(0, len(events_list) - 2)
3965             7         # Accumulate polarities within window
3966             8         polarity_sum = sum([ev[1] for ev in
3967             events_list[i+1:chosen_index+1]])
3968     9 # CORRECT: Good sampling but not integrated

```

Listing 277. C6 Implementation: Paper2Code (Score: 0.8)

```

3970
3971     1 # AutoP2C: Not implemented
3972     2 # CRITICAL ERROR: No event handling

```

Listing 278. C6 Implementation: AutoP2C (Score: 0.0)

Analysis: NERIFY integrates event sampling with Nerfstudio's data pipeline. Paper2Code has detailed sampling logic but fails to integrate it. Others have minimal or no implementation.

Component C7: Camera Pose Interpolation - All Baselines

Uses spherical linear interpolation for rotation and cubic for translation.

```

3982
3983     1 # \nerify\ : Leverages Nerfstudio's built-in
3984         interpolation
3985     2 # Handled by Nerfstudio's camera system
3986     3 # CORRECT: Uses framework's optimized implementation

```

Listing 279. C7 Implementation: NERIFY (Score: 0.9)

```

3988
3989     1 # GPT-5: Correct slerp implementation
3990     2 def slerp(q0, q1, t):
3991         3     dot = torch.sum(q0 * q1, dim=-1, keepdim=True)
3992         4     dot = torch.clamp(dot, -1.0, 1.0)
3993         5     theta_0 = torch.acos(dot)
3994         6     sin_theta_0 = torch.sin(theta_0)
3995         7     s0 = torch.sin(theta_0 - theta_t) / sin_theta_0
3996         8     s1 = torch.sin(theta_t) / sin_theta_0
3997         9     return s0 * q0 + s1 * q1
3998     10 # ISSUE: Missing cubic interpolation for translation

```

Listing 280. C7 Implementation: GPT-5 (Score: 0.6)

```

4000
4001     1 # DeepSeek R1: Not implemented
4002     2 # Uses placeholder compute_rays function
4003     3 # CRITICAL ERROR: No pose interpolation

```

Listing 281. C7 Implementation: DeepSeek R1 (Score: 0.0)

```

4005
4006     1 # Paper2Code: Full implementation
4007     2 def _interpolate_pose(self, timestamp: float) ->
4008         torch.Tensor:
4009         3     # Slerp for rotation
4010         4     q_interp = Utils.slerp(q0, q1, weight)
4011         5     R_interp = self.quaternion_to_matrix(q_interp)
4012         6     # Cubic for translation
4013         7     t_interp = Utils.cubic_interpolation(t0_vec,
4014             t1_vec, weight)
4015     8 # CORRECT: Both slerp and cubic interpolation

```

Listing 282. C7 Implementation: Paper2Code (Score: 1.0)

```

4017
4018     1 # AutoP2C: Not implemented
4019     2 # CRITICAL ERROR: No pose handling

```

Listing 283. C7 Implementation: AutoP2C (Score: 0.0)

Analysis: Paper2Code surprisingly has the most complete pose interpolation. NERIFY relies on Nerfstudio's implementation. GPT-5 has partial slerp. Others lack this component.

Component C8: Hash-based Encoding - All Baselines

Instant-NGP style multiresolution hash encoding with 16 levels.

```

4028
4029     1 # \nerify\ : Uses Nerfstudio's optimized hash
4030         encoding
4031     2 # Configured in NerfactoField with proper parameters
4032     3 # L=16, log2_hashmap=19, N_max=6144
4033     4 # CORRECT: Framework's efficient implementation

```

Listing 284. C8 Implementation: NERIFY (Score: 1.0)

```

4035
4036     1 # GPT-5: Custom hash encoder implementation
4037     2 class HashEncoder(nn.Module):
4038         3     def __init__(self, n_levels=16,
4039             n_features_per_level=2,
4040                 log2_hashmap_size=19):
4041                     self.embeddings = nn.ParameterList([
4042                         nn.Parameter(torch.randn(self.hash_size,
4043                             n_features_per_level) * 1e-4)
4044                         for _ in range(n_levels)
4045                     ])
4046     9 # CORRECT: Full implementation from scratch

```

Listing 285. C8 Implementation: GPT-5 (Score: 1.0)

```

4048
4049 1 # DeepSeek R1: Uses tinycudann
4050 2 self.encoding = tcnn.Encoding(
4051 3     n_input_dims=3,
4052 4     encoding_config={
4053 5         "otype": "HashGrid",
4054 6         "n_levels": 16,
4055 7         "n_features_per_level": 2,
4056 8         "log2_hashmap_size": 19
4057 9     }
4058 10 )
4059 11 # CORRECT: Efficient tinycudann implementation

```

Listing 286. C8 Implementation: DeepSeek R1 (Score: 1.0)

```

4061
4062 1 # Paper2Code: Detailed hash grid implementation
4063 2 class HashGridEncoder(nn.Module):
4064 3     def __init__(self, num_levels=16, dims_per_level
4065 4         =2,
4066 5             hashmap_size_log2=19):
4067 6     for i in range(num_levels):
4068 7         resolution = N_min * ((N_max/N_min)**(i/
4069 8             num_levels-1))
4070 9     self.resolutions.append(resolution)
4071 10 # CORRECT: Complete custom implementation

```

Listing 287. C8 Implementation: Paper2Code (Score: 1.0)

```

4073
4074 1 # AutoP2C: Not implemented
4075 2 # CRITICAL ERROR: No model at all

```

Listing 288. C8 Implementation: AutoP2C (Score: 0.0)

Analysis: All baselines except AutoP2C correctly implement hash encoding, showing this is a well-understood component. The difference lies in integration quality.

3.6.6. Scoring Analysis

Table 53. E-NeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	Weighted Avg _{LLM}
NERFIFY	1.0	1.0	1.0	1.0	1.0	1.0	0.9	1.0	1.00
GPT-5	0.8	1.0	0.8	0.4	0.4	0.2	0.6	1.0	0.60
DeepSeek R1	0.9	1.0	1.0	0.8	0.6	0.3	0.0	1.0	0.72
Paper2Code	0.5	1.0	0.0	0.0	0.0	0.8	1.0	1.0	0.48
AutoP2C	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.05

3.6.7. Why Baselines Fail Despite Component Scores

GPT-5 (Score: 0.60 components, 0% trainable)

- **Strengths:** Near-perfect mathematical implementations, clean modular code
- **Fatal Issues:**
 - No Nerfstudio integration - standalone PyTorch implementation
 - Missing event data pipeline - uses synthetic dummy data
 - Training loop runs only 3 demonstration epochs then terminates
- **Result:** Code compiles but cannot train on real event data

DeepSeek R1 (Score: 0.72 components, 0% trainable)

- **Strengths:** Comprehensive component coverage, uses efficient tinycudann
- **Fatal Issues:**

- Event sampling returns empty - placeholder function not implemented
 - Missing camera pose interpolation entirely
 - No actual training step - only function signatures
 - **Result:** Most complete baseline but critical functions are stubs
- Paper2Code (Score: 0.48 components, 0% trainable)**
- **Strengths:** Extensive 1650-line codebase, detailed event sampling logic
 - **Fatal Issues:**
 - Import errors - references non-existent colmap Python module
 - Missing normalized event loss variant
 - Event sampling not connected to training loop
 - **Result:** Extensive boilerplate that doesn't execute
- AutoP2C (Score: 0.05 components, 0% trainable)**
- **Strengths:** Basic dataset structure attempted
 - **Fatal Issues:**
 - No model implementation whatsoever
 - No loss functions defined
 - Only contains data loading placeholders
 - **Result:** Complete failure - non-functional code

3.6.8. Hyperparameter Fidelity

Table 54. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5	R1	P2C	AutoP2C
Contrast threshold C	0.2	✓	✓	✓	✓	✗
Linlog threshold B	20.0	✓	✓	✓	✓	✗
Batch size	30096	4096*	8	Variable	30096	✗
Hash levels L	16	✓	✓	✓	✓	✗
\log_2 hashmap	19	✓	✓	✓	✓	✗
N_{\max}	$B \cdot 2048$	✓	✓	Variable	✓	✗
Learning rate	$1e-3$	$1e-2^*$	✓	✓	✓	✗
W Score	–	0.95	0.75	0.80	0.60	0.00

*NERFIFY adapts parameters for Nerfstudio optimization

4119

3.6.9. Conclusion

All baselines attempt to implement E-NeRF with varying sophistication, from AutoP2C's complete failure (0.05 score) to DeepSeek R1's comprehensive but non-functional attempt (0.72 score). GPT-5 achieves mathematically correct components but lacks system integration, while Paper2Code generates extensive boilerplate missing critical training logic. Only NERFIFY produces immediately trainable code by properly integrating event-specific losses into Nerfstudio's proven NeRF framework, achieving perfect coverage (1.00 score) and full functionality. This stark contrast—0% trainable rate for all baselines versus 100% for NERFIFY—demonstrates that effective paper-to-code translation for event-based vision demands deep domain specialization over generic approaches.

4120

4121

4122

4123

4124

4125

4126

4127

4128

4129

4130

4131

4132

4133

4134

4135 **3.7. StyleNeRF: A Style-Based 3D-Aware Genera-**
 4136 **tor for High-Resolution Image Synthesis**

4137 **3.7.1. Paper Overview**

4138 StyleNeRF [7] introduces a groundbreaking 3D-aware gener-
 4139 ative model that integrates neural radiance fields with style-
 4140 based generation for high-resolution image synthesis. The
 4141 paper addresses the computational bottleneck in NeRF-based
 4142 image generation by proposing an efficient approximation
 4143 that aggregates features in 2D before final color computation,
 4144 enabling progressive upsampling. This approach achieves
 4145 both high visual quality and 3D consistency while reducing
 4146 rendering time from minutes to seconds.

4147 **3.7.2. Implementation Overview**

Table 55. StyleNeRF Implementation Summary Across All Baselines

Aspect	NERFIFY	GPT-5	DeepSeek R1	Paper2Code	AutoP2C
Lines of Code	742	385	412	1847	623
File Organization	Plugin	Single	Single	Multi-file	Multi-file
Style Modulation	✓	Simplified	✓	✗	✗
2D Upsampling Path	✓	Partial	✓	✗	✗
NeRF-path Regularization	✓	Wrong	✗	✗	✗
Custom Upsampler	✓	~	Partial	✗	✗
View Direction Removal	✓	✗	✗	✗	✗
Nerfstudio Integration	✓	✗	✗	✗	✗
Trainable	✓	✗	✗	✗	✗

Note: NERFIFY always produces trainable code as a complete Nerfstudio plugin

4148 **3.7.3. Novel Components**

Table 56. Novel Components in StyleNeRF with Importance Weights

ID	Component	Weight w_i
C1	Style-conditioned NeRF field (Eq. 2-3)	0.20
C2	Fourier positional encoding (Eq. 1)	0.08
C3	Low-res feature rendering + 2D upsampling (Eq. 5-6)	0.22
C4	Custom upsampler with blur kernel (Eq. 7)	0.15
C5	NeRF-path regularization loss (Eq. 9)	0.15
C6	View direction removal for consistency	0.05
C7	Mapping network (StyleGAN2-based)	0.05
C8	NeRF++ foreground/background split	0.05
C9	Progressive training strategy (3-stage)	0.03
C10	Hyperparameter fidelity	0.02

4149 **3.7.4. Quantitative Metrics**

4150 **3.7.5. Component-by-Component Analysis**

4151 **Component C1: Style-Conditioned NeRF Field - All Base-**
 4152 **lines**

StyleNeRF introduces style modulation for NeRF MLPs using FiLM-style conditioning:

$$\phi_w^n(\mathbf{x}) = g_w^n \circ g_w^{n-1} \circ \dots \circ g_w^1 \circ \zeta(\mathbf{x})$$

where g_w^i are style-modulated layers with weight matrices modulated by style vector \mathbf{w} .

Table 57. StyleNeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	1.00	0.00	0.00	1.00	0.98
GPT-5	0.40	0.30	0.30	0.55	0.52
DeepSeek R1	0.50	0.30	0.20	0.64	0.62
Paper2Code	0.30	0.40	0.30	0.46	0.28
AutoP2C	0.00	0.10	0.90	0.00	0.00

```

1 # \nerify\ : Complete FiLM-style modulation with per
2   -layer affine transforms
3 class StyleMLP(nn.Module):
4   def __init__(self, in_dim, width, depth,
5    style_dim):
6     super().__init__()
7     self.layers = nn.ModuleList([nn.Linear(...),
8      for _ in range(depth)])
9     self.style_affines = nn.ModuleList([
10      nn.Linear(style_dim, 2 * width) for _ in
11      range(depth)
12    ])
13   def forward(self, x, style):
14     for lin, aff in zip(self.layers, self.
15     style_affines):
16       gb = aff(style)
17       gamma, beta = torch.chunk(gb, 2, dim=-1)
18       x = (1.0 + gamma) * lin(x) + beta # FiLM
19       modulation
20       x = self.act(x)
21   return x

```

Listing 289. C1 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Simplified modulation missing +1 formulation
2 class ModulatedMLP(nn.Module):
3   def forward(self, x, w):
4     s = self.style(w)
5     x = x * (s.unsqueeze(1) + 1e-6) # ERROR:
6     Should be (1 + s)
7     x = self.linear(x)
8     if self демодулятор:
9       d = torch.rsqrt((x ** 2).mean(dim=-1,
10      keepdim=True) + 1e-8)
11     x = x * d
12   return self.act(x)

```

Listing 290. C1 Implementation: GPT-5 (Score: 0.6)

```

1 # DeepSeek R1: Correct modulation with demodulation
2 class StyleLinear(nn.Module):
3   def forward(self, x, style):
4     style = self.affine(style)
5     weight = self.weight.unsqueeze(0)
6     weight = weight * (style.unsqueeze(1) + 1) #
7     CORRECT: (s + 1)
8     if self демодулятор:
9       d = torch.rsqrt((weight ** 2).sum(dim=2,
10      keepdim=True) + 1e-8)
11     weight = weight * d
12   return (x.unsqueeze(1) * weight).sum(dim=2) +
13   self.bias

```

Listing 291. C1 Implementation: DeepSeek R1 (Score: 1.0)

```

1 # Paper2Code: No style modulation at all
2 class StyleNerfField(nn.Module):
3   def __init__(self, hidden_units=256):
4     super().__init__()
5     # CRITICAL ERROR: Standard MLP without style
6     conditioning

```

```

4214         self.density_network = nn.Sequential(
4215             nn.Linear(63, hidden_units),
4216             nn.ReLU(),
4217             nn.Linear(hidden_units, 1)
4218         )
4219     # MISSING: Style vector integration entirely
4220 
```

Listing 292. C1 Implementation: Paper2Code (Score: 0.0)

```

4221 # AutoP2C: No model implementation at all
4222 # CRITICAL ERROR: Only dataset loading code present
4223 class AFHQDataset(data.Dataset):
4224     def __init__(self, data_path, image_size=512):
4225         # MISSING: Entire model, no style modulation
4226         implemented
4227         self.dataset = []
4228 
```

Listing 293. C1 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY and DeepSeek R1 correctly implement FiLM-style modulation with proper $(1 + \gamma)$ formulation. GPT-5 has the structure but uses incorrect modulation. Paper2Code and AutoP2C completely miss this core innovation.

Component C2: Fourier Positional Encoding - All Baselines

The paper uses standard NeRF positional encoding with $L = 10$ for positions and $L = 4$ for directions:

$$\zeta^L(x) = [\sin(2^0x), \cos(2^0x), \dots, \sin(2^{L-1}x), \cos(2^{L-1}x)]$$

```

4237 # \nerfify\ : Correct Fourier encoding with proper
4238 # frequencies
4239 self.pos_encoding = NeREncoding(
4240     in_dim=3, num_frequencies=10, # L=10 for
4241     # positions
4242     min_freq_exp=0.0, max_freq_exp=9.0, include_input
4243     =True
4244 )
4245 self.dir_encoding = NeREncoding(
4246     in_dim=3, num_frequencies=4, # L=4 for
4247     # directions
4248     min_freq_exp=0.0, max_freq_exp=3.0, include_input
4249     =True
4250 )
4251 
```

Listing 294. C2 Implementation: NERFIFY (Score: 1.0)

```

4253 # GPT-5: Correct implementation
4254 class FourierEmbedding(nn.Module):
4255     def __init__(self, num_freqs=10):
4256         super().__init__()
4257         self.freq_bands = 2.0 ** torch.linspace(0,
4258             num_freqs - 1, num_freqs)
4259     def forward(self, x):
4260         out = [x]
4261         for f in self.freq_bands:
4262             out.append(torch.sin(f * x))
4263             out.append(torch.cos(f * x))
4264         return torch.cat(out, dim=-1) # CORRECT
4265 
```

Listing 295. C2 Implementation: GPT-5 (Score: 1.0)

```

4267 # DeepSeek R1: Correct encoding
4268 class FourierFeatureEncoding(nn.Module):
4269     def __init__(self, L=10, include_input=True):
4270         super().__init__()
4271         self.L = L
4272         self.include_input = include_input
4273 
```

```

7         self.output_dim = 2 * L * 3 * 2 + (3 if
8         include_input else 0)
4274 
```

Listing 296. C2 Implementation: DeepSeek R1 (Score: 1.0)

```

1 # Paper2Code: Correct formula, wrong usage
2 def positional_encoding(x: torch.Tensor, L: int) ->
3     torch.Tensor:
4     encoding = [x]
5     for i in range(L):
6         for fn in [torch.sin, torch.cos]:
7             encoding.append(fn((2.0 ** i) * x))
8     return torch.cat(encoding, dim=-1)
4275 # ISSUE: Not integrated into model properly
4276 
```

Listing 297. C2 Implementation: Paper2Code (Score: 0.8)

```

1 # AutoP2C: No positional encoding
2 # CRITICAL ERROR: Missing entirely
4277 
```

Listing 298. C2 Implementation: AutoP2C (Score: 0.0)

Analysis: Most baselines correctly implement Fourier encoding as it's a standard NeRF component. Only AutoP2C completely misses it.

Component C3: Low-Resolution Feature Rendering + 2D Upsampling - All Baselines

The core efficiency innovation renders features at 32×32 resolution then upsamples:

$$I_w^{\text{Approx}}(\mathbf{r}) \approx h_c \circ [\phi_w^{n_c, n_\sigma}(\mathcal{A}(\mathbf{r})), \zeta(\mathbf{d})]$$

```

1 # \nerfify\ : Complete 2D aggregation for efficient
2 # upsampling
3 def get_outputs(self, ray_bundle):
4     # Low-res volumetric rendering
5     field_outputs = self.field(ray_samples)
6     weights = ray_samples.get_weights(sigmas)
7     # Aggregate features in 2D (early aggregation)
8     style_feat = field_outputs["style_feat"] # [N, s
8     , C]
9     feat_agg = torch.sum(weights * style_feat, dim
10 =-2) # [N, C]
11     # Decode once from aggregated features
12     approx_rgb = self.field.
13     decode_color_from_features(
14         features=feat_agg, directions=ray_bundle.
15         directions
16     )
17     # Then upsample via synthesis network
4297 
```

Listing 299. C3 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Partial - has low-res rendering but wrong
2 # upsampling
3 def forward(self, z, camera_poses, render_full=True):
4     rays_o, rays_d = self.generate_rays(camera_poses,
4     res=32)
5     feature_map, _, _ = self.generator.volume_render(
5     rays_o, rays_d, styles)
6     # ISSUE: Missing proper 2D aggregation pipeline
7     for block in self.synthesis_blocks:
7         x = block(x, w) # Direct upsampling without
7         proper aggregation
4318 
```

Listing 300. C3 Implementation: GPT-5 (Score: 0.6)

```

4330 1 # DeepSeek R1: Correct approach with progressive
4331    upsampling
4332 2 def forward(self, z, camera_poses, render_full=True):
4333    if not render_full:
4334        rays_o, rays_d = self.generate_rays(
4335            camera_poses, res=32)
4336            feature_map, depth_map, weights = self.
4337            generator.volume_render(
4338                rays_o, rays_d, styles)
4339        return feature_map, depth_map
4340    else:
4341        # Progressive upsampling from low-res
4342        features
4343        feature_map = feature_map.reshape(B, 32, 32,
4344            C)
4345        for upsample_layer, block in zip(self.
4346            upsample_layers, self.blocks):
4347            feature_map = upsample_layer(feature_map,
4348            styles)
4349

```

Listing 301. C3 Implementation: DeepSeek R1 (Score: 0.8)

```

4351 1 # Paper2Code: Direct high-res rendering (wrong
4352    approach)
4353 2 def forward(self, rays_o, rays_d, style):
4354    # CRITICAL ERROR: Renders directly at target
4355    resolution
4356    points = sample_points_on_rays(rays_o, rays_d,
4357        n_samples=64)
4358    densities = self.density_network(points)  # Expensive!
4359    colors = self.color_network(points)
4360    rgb = volume_render(densities, colors, points)
4361    # MISSING: Low-res rendering + upsampling
4362    strategy entirely
4363

```

Listing 302. C3 Implementation: Paper2Code (Score: 0.0)

```

4366 1 # AutoP2C: No rendering implementation
4367 2 # CRITICAL ERROR: Missing entirely
4368

```

Listing 303. C3 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY correctly implements 2D feature aggregation for efficient upsampling. DeepSeek R1 has the right approach. GPT-5 partially implements it. Paper2Code and AutoP2C miss this critical efficiency innovation.

Component C4: Custom Upsampler with Blur Kernel - All Baselines

StyleNeRF's hybrid upsampler combines pixel-shuffle with learnable MLPs and fixed blur:

$$\text{Upsample}(X) = \text{Conv2d}(\text{PixelShuffle}(\text{Repeat}(X, 4) + \psi_\theta(X), 2), K)$$

```

4376 1 # \nerfify\ : Complete hybrid upsampler
4377 2 class HybridUpsampler(nn.Module):
4378 3     def forward(self, x):
4379        x_repeat = x.repeat_interleave(4, dim=1)  # Repeat 4x
4380        psi = self.mlp(x.permute(0, 2, 3, 1))  # Learnable variations
4381        x_combined = x_repeat + psi.permute(0, 3, 1, 2)
4382        x_shuffled = F.pixel_shuffle(x_combined, 2)  # PixelShuffle
4383        x_blurred = self.blur_conv(x_shuffled)  # Fixed blur kernel
4384        return x_blurred
4385

```

Listing 304. C4 Implementation: NERFIFY (Score: 1.0)

```

4391 1 # GPT-5: Partial - missing blur kernel
4392 2 class Upsampler(nn.Module):
4393 3     def forward(self, x):
4394        x_rep = x.repeat(1, 4, 1, 1)
4395        psi = self.mlp(x.permute(0, 2, 3, 1)).permute
4396        (0, 3, 1, 2)
4397        x = F.pixel_shuffle(x_rep + psi, 2)
4398        # MISSING: Fixed blur convolution
4399        return x  # ERROR: No blur kernel applied
4400

```

Listing 305. C4 Implementation: GPT-5 (Score: 0.4)

```

4402 1 # DeepSeek R1: Has blur but simplified
4403 2 class CustomUpsampler(nn.Module):
4404 3     def __init__(self, in_ch, out_ch, style_dim):
4405        super().__init__()
4406        self.mlp = nn.Sequential(StyleLinear(...))
4407        self.blur = nn.Conv2d(in_ch, out_ch, 3,
4408            padding=1, bias=False)
4409        nn.init.dirac_(self.blur.weight)  # Initialize as identity
4410
4411    def forward(self, x, style):
4412        x_repeated = F.pixel_shuffle(x.repeat(1, 4,
4413            1, 1, 2))
4414        x_mlp = self.mlp(x_repeated, style)
4415        x_combined = x_repeated + F.pixel_shuffle(
4416            x_mlp, 2)
4417        return self.blur(x_combined)  # Apply blur
4418

```

Listing 306. C4 Implementation: DeepSeek R1 (Score: 0.8)

```

4420 1 # Paper2Code: No custom upsampler
4421 2 # CRITICAL ERROR: Missing entirely, using standard
4422    upsampling
4423

```

Listing 307. C4 Implementation: Paper2Code (Score: 0.0)

```

4425 1 # AutoP2C: No upsampler implementation
4426 2 # CRITICAL ERROR: Missing entirely
4427

```

Listing 308. C4 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERFIFY fully implements the hybrid upsampler with all components. DeepSeek R1 has most elements. GPT-5 misses the blur kernel. Paper2Code and AutoP2C don't implement it.

Component C5: NeRF-Path Regularization Loss - All Baselines

Novel loss enforcing 3D consistency:

$$\mathcal{L}_{\text{NeRF-path}} = \frac{1}{|S|} \sum_{(i,j) \in S} (I_w^{\text{Approx}}[i,j] - I_w^{\text{NeRF}}[i,j])^2$$

```

4435 1 # \nerfify\ : Correct NeRF-path regularization
4436 2 def get_loss_dict(self, outputs, batch):
4437    if self.config.nerf_path_loss_mult > 0.0:
4438        rgb_full = outputs["rgb"]  # Full volumetric
4439        RGB
4440        rgb_approx = outputs["approx_rgb"]  # Aggregated-feature RGB
4441        # Subsample pixels (25% by default)
4442        num = int(self.config.nerf_path_sample_frac *
4443            rgb_full.shape[0])
4444        idx = torch.randperm(rgb_full.shape[0])[:num]
4445        nerf_path = F.mse_loss(rgb_approx[idx],
4446            rgb_full[idx])
4447        loss_dict["nerf_path_reg"] = 0.2 * nerf_path
4448        # = 0.2
4449

```

Listing 309. C5 Implementation: NERFIFY (Score: 1.0)

4452
4453
4454
4455
4456
4457
4458
4459
4460
4461
4462
4463
4464

```
1 # GPT-5: Wrong formulation entirely
2 def nerf_path_reg(fake_high, fake_low):
3     # CRITICAL ERROR: Comparing high-res vs
4     # downsampled
5     # Should be full volumetric vs aggregated
6     # features!
7     return F.mse_loss(fake_high, fake_low.detach())
8 # In training:
9 reg = 0.2 * nerf_path_reg(fake_img,
10    F.interpolate(fake_img, scale_factor=0.5)) #
```

WRONG!

Listing 310. C5 Implementation: GPT-5 (Score: 0.2)

4465
4466
4467
4468
4469
4470
4471
4472
4473

```
1 # DeepSeek R1: Missing NeRF-path regularization
2 # CRITICAL ERROR: Not implemented
3 class NeRFPathRegularization(nn.Module):
4     def forward(self, approx_output, nerf_output,
5                 mask=None):
6         # ERROR: Defined but never used in training
7         # loop
```

Listing 311. C5 Implementation: DeepSeek R1 (Score: 0.0)

4474
4475
4476
4477
4478
4479
4480
4481
4482
4483
4484

```
1 # Paper2Code: No NeRF-path regularization
2 def nerf_path_regularization(predictions, targets):
3     # CRITICAL ERROR: Wrong concept - not the paper's
4     # formulation
5     squared_differences = (sampled_predictions -
6     sampled_targets) ** 2
7     # MISSING: Comparison between full and approx
8     # paths
```

Listing 312. C5 Implementation: Paper2Code (Score: 0.0)

4485
4486
4487
4488
4489

```
1 # AutoP2C: No loss implementation
2 # CRITICAL ERROR: Missing entirely
```

Listing 313. C5 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERFIFY correctly implements the NeRF-path regularization comparing full volumetric vs aggregated-feature RGB. GPT-5 implements a completely wrong loss. Others don't implement it at all.

Component C6: View Direction Removal - All Baselines

Paper removes view direction conditioning from color prediction to improve 3D consistency.

4496
4497
4498
4499
4500
4501
4502
4503
4504
4505
4506
4507
4508
4509
4510
4511
4512
4513

```
1 # \nerify\ : Correctly removes view direction
2 def __init__(self, remove_viewdir: bool = True):
3     self.remove_viewdir = remove_viewdir
4     def _encode_directions(self, directions):
5         if self.remove_viewdir:
6             # Return empty tensor - no view direction
7             return torch.zeros((directions.shape[:-1],
8                               device=directions.device))
9         return self.dir_encoding(directions)
```

Listing 314. C6 Implementation: NERFIFY (Score: 1.0)

```
1 # GPT-5: Still uses view direction
2 self.color_out = nn.Linear(hidden_dim, 3)
3 # ERROR: Color prediction still depends on view
4 # direction
5 # MISSING: remove_viewdir flag and implementation
```

Listing 315. C6 Implementation: GPT-5 (Score: 0.0)

1 # DeepSeek R1: No view direction removal	4515
2 def forward_color(self, x, style, is_background=False	4517
3):	4518
4 # ERROR: Still conditions on view direction	4519
5 implicitly	4520
6 color = self.color_head(h) # No removal	4521
7 implemented	4523

Listing 316. C6 Implementation: DeepSeek R1 (Score: 0.0)

1 # Paper2Code: No view direction removal	4524
2 # MISSING: Not implemented	4525

Listing 317. C6 Implementation: Paper2Code (Score: 0.0)

1 # AutoP2C: No model implementation	4528
2 # MISSING: Entirely	4529

Listing 318. C6 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERFIFY implements view direction removal. All baselines miss this 3D consistency improvement.

Component C7: Mapping Network - All Baselines

StyleGAN2-style mapping from z to w with 8 FC layers.

1 # \nerify\ : Complete mapping network	4536
2 self.mapping_network = MappingNetwork(4537
3 latent_dim=512, style_dim=512,	4538
4 num_layers=8, lr_multiplier=0.01 # 100x lower LR	4539
5)	4540

Listing 319. C7 Implementation: NERFIFY (Score: 1.0)

1 # GPT-5: Correct mapping network	4543
2 class MappingNetwork(nn.Module):	4544
3 def __init__(self, z_dim=512, w_dim=512, n_layers	4545
4 =8):	4546
5 layers = []	4547
6 for i in range(n_layers):	4548
7 layers.append(nn.Linear(z_dim if i==0	4549
8 else w_dim, w_dim))	4550
9 layers.append(nn.LeakyReLU(0.2))	4551
10 self.mapping = nn.Sequential(*layers)	4552

Listing 320. C7 Implementation: GPT-5 (Score: 1.0)

1 # DeepSeek R1: Correct implementation	4555
2 class MappingNetwork(nn.Module):	4556
3 def forward(self, z):	4557
4 z = F.normalize(z, dim=-1) # Normalize input	4558
5 for layer in self.layers:	4559
6 z = F.leaky_relu(layer(z), 0.2)	4560
7 return z	4561

Listing 321. C7 Implementation: DeepSeek R1 (Score: 1.0)

1 # Paper2Code: Has mapping but wrong integration	4564
2 class MappingNetwork(nn.Module):	4565
3 def __init__(self, latent_dim, style_dim,	4566
4 num_layers):	4567
5 # ISSUE: Correct structure but not properly	4568
6 used	4569
7 self.mapping = nn.Sequential(*layers)	4570

Listing 322. C7 Implementation: Paper2Code (Score: 0.6)

1 # AutoP2C: No mapping network	4573
2 # MISSING: Entirely	4574

Listing 323. C7 Implementation: AutoP2C (Score: 0.0)

4577 Analysis: Most baselines implement the mapping network correctly as it's a standard StyleGAN2 component.

4579 Component C8: NeRF++ Foreground/Background Split - All Baselines

4581 NeRF++ architecture with inverted sphere parameteriza-
4582 tion for background.

```
1 # \nerfify\ : Complete NeRF++ implementation
2 # Foreground in unit sphere, background with inverted
   sphere
3 self.spatial_distortion = None # Use scene box
   normalization
4 # Handles fg/bg split properly in field
   implementation
```

Listing 324. C8 Implementation: NERFIFY (Score: 1.0)

```
1 # GPT-5: Correct fg/bg split
2 # Foreground network
3 self.field = StyleNeRFField(hidden_dim=256)
4 # Background network (smaller)
5 self.bg_field = StyleNeRFField(hidden_dim=128)
```

Listing 325. C8 Implementation: GPT-5 (Score: 1.0)

```
1 # DeepSeek R1: Correct implementation
2 def forward_geometry(self, x, style, is_background=
   False):
3     if is_background:
4         r = torch.sqrt((x ** 2).sum(dim=-1, keepdim=
   True))
5         x_inv = torch.cat([x / r, 1.0 / r], dim=-1)
   # Inverted sphere
```

Listing 326. C8 Implementation: DeepSeek R1 (Score: 1.0)

```
1 # Paper2Code: Mentions but doesn't implement properly
2 # ISSUE: No actual fg/bg separation in rendering
```

Listing 327. C8 Implementation: Paper2Code (Score: 0.4)

```
1 # AutoP2C: No NeRF++ implementation
2 # MISSING: Entirely
```

Listing 328. C8 Implementation: AutoP2C (Score: 0.0)

4617 Analysis: NERFIFY , GPT-5, and DeepSeek R1 correctly
4618 implement NeRF++. Paper2Code mentions it but doesn't
4619 implement properly.

4620 Component C9: Progressive Training Strategy - All
4621 Baselines

4622 Three-stage training: 500k low-res, 5M progressive, 25M
4623 high-res.

```
1 # \nerfify\ : Has progressive structure
2 self.stage1_images = 500000
3 self.stage2_images = 5000000
4 self.stage3_images = 25000000
5 # Implemented through config but simplified
```

Listing 329. C9 Implementation: NERFIFY (Score: 0.8)

```
1 # GPT-5: No progressive training
2 # ERROR: Trains at target resolution from start
3 for iteration in range(num_iterations):
4     # MISSING: Resolution scheduling
```

Listing 330. C9 Implementation: GPT-5 (Score: 0.0)

```
1 # DeepSeek R1: Has stages but incomplete
2 stages = [(32, 500000), (1024, 5000000), (1024,
   25000000)]
3 # ISSUE: Structure exists but not fully implemented
```

Listing 331. C9 Implementation: DeepSeek R1 (Score: 0.4)

```
1 # Paper2Code: Defines stages but wrong usage
2 class ProgressiveTrainingConfig:
3     stage1_images: int = 500_000
4     # ISSUE: Not integrated with resolution changes
```

Listing 332. C9 Implementation: Paper2Code (Score: 0.6)

```
1 # AutoP2C: No training implementation
2 # MISSING: Entirely
```

Listing 333. C9 Implementation: AutoP2C (Score: 0.0)

4653 Analysis: No baseline fully implements progressive train-
4654 ing. NERFIFY has the structure. Paper2Code defines it but
4655 doesn't use it properly.

4656 Component C10: Hyperparameter Fidelity - All Base-
4657 lines

Matching paper's exact hyperparameters for reproducibil-
ity.

```
1 # \nerfify\ : All hyperparameters match paper
2 style_dim=512, trunk_layers=4, trunk_hidden_dim=256,
3 color_layers=2, color_hidden_dim=128,
4 nerf_path_loss_mult=0.2, nerf_path_sample_frac=0.25
```

Listing 334. C10 Implementation: NERFIFY (Score: 1.0)

```
1 # GPT-5: Some parameters correct
2 style_dim=512, hidden_dim=256
3 # ISSUE: Missing several key parameters
```

Listing 335. C10 Implementation: GPT-5 (Score: 0.6)

```
1 # DeepSeek R1: Most parameters correct
2 style_dim=512, hidden_dim=256, bg_hidden_dim=128
3 # ISSUE: Missing nerf_path parameters
```

Listing 336. C10 Implementation: DeepSeek R1 (Score: 0.7)

```
1 # Paper2Code: Partial match
2 latent_dim: int = 512
3 # ISSUE: Many incorrect or missing
```

Listing 337. C10 Implementation: Paper2Code (Score: 0.5)

```
1 # AutoP2C: No hyperparameters
2 # MISSING: Entirely
```

Listing 338. C10 Implementation: AutoP2C (Score: 0.0)

4685 Analysis: Only NERFIFY matches all hyperparameters
4686 exactly. Others have partial matches.

Table 58. StyleNeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	Weighted Avg _{LLM}
NERIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.8	1.0	0.98
GPT-5	0.6	1.0	0.6	0.4	0.2	0.0	1.0	1.0	0.0	0.6	0.52
DeepSeek R1	1.0	1.0	0.8	0.8	0.0	0.0	1.0	1.0	0.4	0.7	0.62
Paper2Code	0.0	0.8	0.0	0.0	0.0	0.0	0.6	0.4	0.6	0.5	0.28
AutoP2C	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.00

4687 **3.7.6. Scoring Analysis**4688 **3.7.7. Why Baselines Fail Despite Component Scores**4689 **GPT-5 (Score: 0.52 components, 0% trainable)**

- **Strengths:** Correct Fourier encoding, mapping network, NeRF++ split
- **Fatal Issues:**
 - Wrong NeRF-path regularization (compares high-res vs downsampled instead of volumetric vs aggregated)
 - No Nerfstudio integration - standalone implementation incompatible with benchmark
 - Missing view direction removal and blur kernel in up-sampler
- **Result:** Cannot train on benchmark datasets, incorrect 3D consistency enforcement

4701 **DeepSeek R1 (Score: 0.62 components, 0% trainable)**

- **Strengths:** Perfect style modulation with demodulation, correct NeRF++ architecture
- **Fatal Issues:**
 - Missing NeRF-path regularization entirely - loses key 3D consistency benefit
 - No Nerfstudio plugin structure - cannot integrate with training pipeline
 - Missing view direction removal
- **Result:** Strong mathematical understanding but incomplete system, cannot execute training

4712 **Paper2Code (Score: 0.28 components, 0% trainable)**

- **Strengths:** Configuration management, dataset loading structure
- **Fatal Issues:**
 - No style modulation - uses standard MLPs without conditioning
 - Missing 2D upsampling path - renders directly at high resolution (computationally infeasible)
 - No NeRF-path regularization or custom upsampler
- **Result:** Fundamental misunderstanding of StyleNeRF architecture, cannot scale beyond 256x256

4723 **AutoP2C (Score: 0.00 components, 0% trainable)**

- **Strengths:** Attempted dataset loading (with incorrect URLs)
- **Fatal Issues:**
 - No model implementation whatsoever
 - Only generates dataset and evaluation code
 - Import errors prevent execution
- **Result:** Complete failure to implement any StyleNeRF components

3.7.8. Hyperparameter Fidelity

4732

Table 59. Hyperparameter Implementation Accuracy

Parameter	Paper	NERIFY	GPT-5	R1	P2C	AutoP2C
Style dim	512	✓	✓	✓	✓	✗
Latent dim	512	✓	✓	✓	✓	✗
Mapping layers	8	✓	✓	✓	✓	✗
FG hidden dim	256	✓	✓	✓	✓	✗
BG hidden dim	128	✓	✓	✓	256	✗
L_{pos}	10	✓	✓	✓	✓	✗
L_{dir}	4	✓	✓	✓	✓	✗
Initial resolution	32	✓	✓	✓	64	✗
β (NeRF-path)	0.2	✓	Wrong	✗	✗	✗
λ (R1 reg)	0.5	✓	✗	✗	0.5	✗
W Score	—	1.00	0.55	0.64	0.46	0.00

3.7.9. Conclusion

4733

All baselines attempt to implement StyleNeRF with varying degrees of sophistication. GPT-5 and DeepSeek R1 demonstrate understanding of core architectural concepts but fail on critical details - GPT-5 implements the wrong NeRF-path regularization while DeepSeek R1 omits it entirely. Paper2Code fundamentally misunderstands the method, implementing standard NeRF without style conditioning or the efficiency innovations. AutoP2C completely fails, producing only non-functional dataset code. Despite some baselines achieving reasonable component scores (DeepSeek R1: 0.62, GPT-5: 0.52), none produce trainable code due to missing Nerfstudio integration, incorrect loss formulations, or architectural failures. Only NERIFY produces immediately trainable code as a complete Nerfstudio plugin, demonstrating the critical importance of domain-specific synthesis with a stark 100% vs 0% trainable rate highlighting the necessity of specialized paper-to-code approaches for complex vision research.

4734

4735

4736

4737

4738

4739

4740

4741

4742

4743

4744

4745

4746

4747

4748

4749

4750

4751

3.8. iNeRF: Inverting Neural Radiance Fields for Pose Estimation

4752

4753

3.8.1. Paper Overview

4754

iNeRF [31] introduces a framework for 6DoF pose estimation by inverting a trained NeRF model through gradient-based optimization. The method performs mesh-free pose refinement by backpropagating photometric loss through differentiable volume rendering, optimizing camera poses on the SE(3) manifold. Key innovations include interest-region ray sampling for computational efficiency, exponential map parameterization for valid pose updates, and self-supervised NeRF training by estimating poses for unlabeled images.

4755

4756

4757

4758

4759

4760

4761

4762

4763

Table 60. iNeRF Implementation Summary Across All Baselines

Aspect	NERFIFY	GPT-5	DeepSeek R1	Paper2Code	AutoP2C
Lines of Code	580	180	650	2100	950
File Organization	Plugin	Single	Single	Multi-file	Multi-file
SE(3) Exp Map	✓	Partial	✓	✓	✗
Interest Sampling	✓	Simplified	✓	✓	✗
YUV Loss	✓	✗	✓	✓	✗
LR Schedule	✓	✗	Wrong	✓	✗
Nerfstudio Integration	✓	✗	✗	✗	✗
Trainable	✓	✗	✗	✗	✗

Note: NERFIFY always produces trainable code as a complete Nerfstudio plugin

Table 61. Novel Components in iNeRF with Importance Weights

ID	Component	Weight w_i
C1	SE(3) exponential map parameterization	0.20
C2	Interest region sampling with dilation	0.18
C3	Interest point sampling (corner detection)	0.12
C4	Photometric RGB MSE loss	0.10
C5	YUV-UV loss variant (LineMOD)	0.08
C6	Camera LR schedule ($\alpha_t = \alpha_0 \cdot 0.8^{t/100}$)	0.10
C7	Small ray batch (2048 rays)	0.08
C8	Gradient-based pose updates	0.07
C9	Volume rendering integration	0.04
C10	Pose-only optimization mode	0.03

Table 62. iNeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	1.00	0.00	0.00	1.00	0.97
GPT-5	0.50	0.30	0.20	0.60	0.58
DeepSeek R1	0.60	0.30	0.10	0.70	0.68
Paper2Code	0.70	0.20	0.10	0.80	0.75
AutoP2C	0.00	0.10	0.90	0.00	0.05

3.8.2. Implementation Overview

3.8.3. Novel Components

3.8.4. Quantitative Metrics

3.8.5. Component-by-Component Analysis

Component C1: SE(3) Exponential Map Parameterization - All Baselines

The paper specifies pose updates via the exponential map from Lie algebra $\mathfrak{se}(3)$ to $\text{SE}(3)$: $T = \exp([\xi]_{\times})T_0$ where $\xi = [\omega, v]^T \in \mathbb{R}^6$ and the rotation matrix uses Rodrigues' formula: $R = I + \frac{\sin \theta}{\theta}[\omega]_{\times} + \frac{1-\cos \theta}{\theta^2}[\omega]_{\times}^2$.

```
1 # \nerfify\ : Uses Nerfstudio's camera optimizer with
2     # correct SE(3)
3 camera_optimizer=CameraOptimizerConfig(mode="SO3xR3")
4 # In callbacks:
5 def set_camera_lr(step: int) -> None:
6     T_update = se3_exponential_map(self.pose_params)
7     T_current = T_update @ self.initial_pose
8     # CORRECT: Full SE(3) manifold parameterization
```

Listing 339. C1 Implementation: NERFIFY (Score: 1.0)

```
1 # GPT-5: SE(3) exp map without small angle handling
2 def se3_expmmap(screw):
3     omega = screw[:, :3]; v = screw[:, 3:]
4     theta = torch.norm(omega, dim=1, keepdim=True) +
5         1e-8
6     # MISSING: No small angle approximation
7     R = I + torch.sin(theta) * omega_hat / theta +
8         (1 - torch.cos(theta)) * (omega_hat @
9             omega_hat) / theta**2
10    K = I * theta + (1 - torch.cos(theta)) *
11        omega_hat
12    t = (K @ v [..., None]).squeeze(-1)
13    # ERROR: Numerical instability for small rotations
```

4784
4785
4786
4787
4788
4789
4790
4791
4792
4793
4794
4795
4796
4797
4798

Listing 340. C1 Implementation: GPT-5 (Score: 0.8)

```
1 # DeepSeek R1: Complete SE(3) with small angle
2     # handling
3 def se3_exp_map(self, screw_axis):
4     omega = screw_axis[:3]; v = screw_axis[3:]
5     theta = torch.norm(omega)
6     if theta < 1e-8: # CORRECT: Small angle
7         approximation
8             R = torch.eye(3) + self.skew_symmetric(omega)
9             t = v
10    else:
11        # CORRECT: Full Rodrigues formula
12        omega_hat = omega / theta
13        omega_hat_skew = self.skew_symmetric(
14            omega_hat)
15        R = torch.eye(3) + torch.sin(theta) *
16            omega_hat_skew +
17            (1 - torch.cos(theta)) * omega_hat_skew @
18            omega_hat_skew
19        V = (torch.sin(theta)/theta) * torch.eye(3) +
20        ...
21        t = V @ v.unsqueeze(-1).squeeze(-1)
```

4799
4800
4801
4802
4803
4804
4805
4806
4807
4808
4809
4810
4811
4812
4813
4814
4815
4816
4817
4818
4819
4820

Listing 341. C1 Implementation: DeepSeek R1 (Score: 1.0)

```
1 # Paper2Code: Textbook-correct SE(3) implementation
2 def se3_exponential_map(xi: torch.Tensor) -> torch.
3     Tensor:
4     omega = xi[0:3]; nu = xi[3:6]
5     theta = torch.norm(omega)
6     eps: float = 1e-6
7     if theta.item() < eps: # CORRECT: Small angle
8         handling
9         R = I + A; t = nu
10    else:
11        sin_theta = torch.sin(theta); cos_theta =
12            torch.cos(theta)
13        R = I + (sin_theta/theta)*A + ((1-cos_theta)/
14            theta**2)*A2
15        t = (I*theta + (1-cos_theta)*A + (theta-
16            sin_theta)*A2) @ (nu/theta)
17    T = torch.eye(4); T[0:3, 0:3] = R; T[0:3, 3] = t
18    # CORRECT: Complete SE(3) with proper
19    parameterization
```

4822
4823
4824
4825
4826
4827
4828
4829
4830
4831
4832
4833
4834
4835
4836
4837
4838
4839
4840
4841

Listing 342. C1 Implementation: Paper2Code (Score: 1.0)

```
1 # AutoP2C: No SE(3) parameterization at all
2 def optimize_pose(self, image, initial_pose):
3     pose = initial_pose.clone().detach()
4     requires_grad_(True)
5     optimizer = optim.Adam([pose], lr=0.01)
6     # CRITICAL ERROR: Direct matrix optimization
7     for _ in range(100):
8         loss.backward()
9         optimizer.step()
10        with torch.no_grad():
11            pose[:] = self.project_to_se3(pose) # Empty placeholder
12    def project_to_se3(self, pose):
```

4843
4844
4845
4846
4847
4848
4849
4850
4851
4852
4853
4854
4855
4856

4858

```
12     return pose # WRONG: No actual projection
```

Listing 343. C1 Implementation: AutoP2C (Score: 0.0)

Analysis: NERIFY, DeepSeek R1, and Paper2Code correctly implement the full SE(3) exponential map with proper Rodrigues formula and small angle handling. GPT-5 implements the correct formulas but lacks small angle approximation, causing numerical instability. AutoP2C completely fails by optimizing the pose matrix directly without manifold constraints.

Component C2: Interest Region Sampling with Dilatation - All Baselines

Interest region sampling computes gradient magnitudes using Sobel filters, thresholds to create a binary mask, applies morphological dilation (5x5 kernel, 2 iterations), and samples rays from the dilated regions.

```
1 # \nerify\ : Complete interest region with Sobel and
2     dilation
3 def _compute_interest_mask(image_hw3, threshold,
4     dilate_iters):
5     # Sobel filters for edge detection
6     kx = torch.tensor([[1, 0, -1], [2, 0, -2], [1, 0,
7         -1]])
8     ky = torch.tensor([[1, 2, 1], [0, 0, 0], [-1, -2,
9         -1]])
10    gx = F.conv2d(gray_, kx, padding=1)
11    gy = F.conv2d(gray_, ky, padding=1)
12    grad = torch.sqrt(gx.pow(2) + gy.pow(2))
13    mask = (grad_norm >= threshold).float()
14    # CORRECT: Morphological dilation via max pooling
15    for _ in range(dilate_iters):
16        mask = F.max_pool2d(mask, kernel_size=5,
17            stride=1, padding=2)
```

Listing 344. C2 Implementation: NERIFY (Score: 1.0)

```
1 # GPT-5: Simplified gradient without Sobel or
2     dilation
3 def interest_region_sampling(image, num_rays=2048):
4     # WRONG: Simple difference instead of Sobel
5     grad_x = torch.abs(image[:, :, 1:] - image[:, :, :-1])
6     grad_y = torch.abs(image[:, 1:, :] - image[:, :-1, :])
7     grad = F.pad(grad_x[..., :-1] + grad_y[..., :-1],
8         (0, 1, 0, 1))
9     # MISSING: No morphological dilation
10    prob = grad.flatten() / grad.sum()
11    idx = torch.multinomial(prob, num_rays,
12        replacement=False)
13    # ERROR: Missing core algorithm components
```

Listing 345. C2 Implementation: GPT-5 (Score: 0.4)

```
1 # DeepSeek R1: Complete with OpenCV Sobel and
2     dilation
3 def interest_region_sampling(self, pose, image,
4     num_rays):
5     gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)
6     grad_x = cv2.Sobel(gray, cv2.CV_64F, 1, 0, ksize
7         =3)
8     grad_y = cv2.Sobel(gray, cv2.CV_64F, 0, 1, ksize
9         =3)
10    grad_mag = np.sqrt(grad_x**2 + grad_y**2)
11    mask = (grad_mag > 0.2).astype(np.uint8)
12    # CORRECT: 5x5 kernel, 2 iterations
13    kernel = np.ones((5, 5), np.uint8)
14    mask = cv2.dilate(mask, kernel, iterations=2)
```

4923

Listing 346. C2 Implementation: DeepSeek R1 (Score: 1.0)

```
1 # Paper2Code: Comprehensive interest region
2     implementation
3 def interest_region_sampling(self, observed_img,
4     num_rays):
5     sobelx = cv2.Sobel(gray, cv2.CV_64F, 1, 0, ksize
6         =3)
7     sobely = cv2.Sobel(gray, cv2.CV_64F, 0, 1, ksize
8         =3)
9     gradient_magnitude = np.sqrt(sobelx**2 + sobely
10        **2)
11    binary_mask = (gradient_magnitude >= 0.2).astype(
12        np.uint8)
13    # CORRECT: Proper dilation parameters
14    kernel = cv2.getStructuringElement(cv2.MORPH_RECT
15        , (5, 5))
16    dilated_mask = cv2.dilate(binary_mask, kernel,
17        iterations=2)
```

Listing 347. C2 Implementation: Paper2Code (Score: 1.0)

```
1 # AutoP2C: No interest sampling at all
2 def generate_rays(self):
3     for pose in self.camera_poses:
4         directions = self.compute_ray_directions(self
5             .image_size,
6             self
7                 .intrinsics['focal_length'])
8         rays_world = self.transform_rays(directions,
9             pose)
10    # CRITICAL ERROR: Generates rays for ALL pixels
11    # MISSING: No gradient computation, no sampling
12        strategy
```

Listing 348. C2 Implementation: AutoP2C (Score: 0.0)

Analysis: NERIFY, DeepSeek R1, and Paper2Code correctly implement complete interest region sampling with Sobel filters and morphological dilation. GPT-5 uses simplified gradients without proper edge detection or dilation. AutoP2C completely lacks any interest sampling mechanism.

Component C3: Interest Point Sampling - All Baselines

Interest point sampling detects corners/features and samples rays from high-gradient locations without dilation.

```
1 # \nerify\ : Interest point via top-k gradient
2     magnitudes
3 elif strategy == "interest_point":
4     # Sample from top-k gradient magnitudes (no
5         dilation)
6     mag = torch.sqrt(gx.pow(2) + gy.pow(2))
7     flat = mag.view(-1)
8     k = min(N * 10, flat.numel())
9     topk = torch.topk(flat, k=k, largest=True)
10    idxs = topk.indices
11    ys, xs = idxs // W, idxs % W
12    # CORRECT: Distinct from region sampling
```

Listing 349. C3 Implementation: NERIFY (Score: 1.0)

```
1 # GPT-5: No separate interest point implementation
2 # Uses same simplified gradient approach for all
3     sampling
3 def interest_region_sampling(image, num_rays=2048):
4     grad = F.pad(grad_x[..., :-1] + grad_y[..., :-1],
5         (0, 1, 0, 1))
6     prob = grad.flatten() / grad.sum()
6     # ISSUE: No distinction between point and region
7         sampling
```

4924
4925
4926
4927
4928
4929
4930
4931
4932
4933
4934
4935
4936
4937
4938
4939
4940
49414943
4944
4945
4946
4947
4948
4949
4950
4951
4952
4953
4954
49554957
4958
4959
4960
4961
4962
4963
4964
4965
4966
4967
4968
4969
4970
4971
4972
4973
4974
4975
4976
4977
4978

4989

Listing 350. C3 Implementation: GPT-5 (Score: 0.5)

```

4990 1 # DeepSeek R1: Harris corner detection for interest
4991   points
4992 2 def interest_point_sampling(self, pose,
4993     observed_image, batch_size):
4994   3 gray = cv2.cvtColor(observed_image, cv2.
4995     COLOR_RGB2GRAY)
4996 4 corners = cv2.cornerHarris(gray.astype(np.float32)
4997   ), 2, 3, 0.04)
4998 5 corners = cv2.dilate(corners, None)  # Local
4999   maxima
5000 6 threshold = 0.01 * corners.max()
5001 7 interest_mask = corners > threshold
5002 8 interest_coords = np.argwhere(interest_mask)
5003 9 # CORRECT: Proper corner detection

```

Listing 351. C3 Implementation: DeepSeek R1 (Score: 1.0)

```

5006 1 # Paper2Code: goodFeaturesToTrack for interest points
5007 2 def _interest_point_sampling(self, image_np, H, W):
5008   3 corners = cv2.goodFeaturesToTrack(gray,
5009     maxCorners=self.batch_size, qualityLevel
5010       =0.01, minDistance=10)
5011 5 if corners is not None:
5012   6 keypoints = corners.reshape(-1, 2)
5013 7 # CORRECT: Standard feature detection

```

Listing 352. C3 Implementation: Paper2Code (Score: 1.0)

```

5016 1 # AutoP2C: No interest point sampling
5017 2 # MISSING: No feature detection or corner detection
5018 3 # MISSING: No sampling strategy differentiation

```

Listing 353. C3 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY , DeepSeek R1, and Paper2Code implement distinct interest point sampling methods. GPT-5 doesn't differentiate between point and region sampling. AutoP2C lacks any interest-based sampling.

Component C4: Photometric RGB MSE Loss - All Baselines

The paper uses standard RGB mean squared error loss:

$$\mathcal{L} = \sum_{r \in \mathcal{R}} \|\hat{C}(r) - C(r)\|_2^2.$$

```

5029 1 # \nerify\ : Standard RGB MSE via Nerfacto base
5030   class
5031 2 loss_dict = super().get_loss_dict(outputs, batch,
5032     metrics_dict)
5033 3 # Inherits correct RGB loss: torch.mean((pred - gt)
5034   ** 2)

```

Listing 354. C4 Implementation: NERFIFY (Score: 1.0)

```

5037 1 # GPT-5: Correct RGB MSE
5038 2 loss = F.mse_loss(rendered, observed)
5039 3 # CORRECT: Standard photometric loss

```

Listing 355. C4 Implementation: GPT-5 (Score: 1.0)

```

5042 1 # DeepSeek R1: RGB MSE implementation
5043 2 if loss_type == '12':
5044   3 loss = torch.mean((rendered_rgb - observed_rgb)
5045     ** 2)
5046 4 # CORRECT: Standard MSE loss

```

Listing 356. C4 Implementation: DeepSeek R1 (Score: 1.0)

```

5049 1 # Paper2Code: Photometric loss with proper
5050   normalization
5051 2 def photometric_loss(predicted, observed):
5052   3 difference = predicted - observed
5053   4 squared_difference = difference ** 2
5054   5 photometric_loss = total_error / (H * W * 3)
5055 6 # CORRECT: Normalized MSE

```

Listing 357. C4 Implementation: Paper2Code (Score: 1.0)

```

5058 1 # AutoP2C: Generic loss without context
5059 2 def photometric_loss(predicted, observed):
5060   3 difference = predicted - observed
5061   4 squared_difference = difference ** 2
5062 5 # ISSUE: Oversimplified, no ray context

```

Listing 358. C4 Implementation: AutoP2C (Score: 0.5)

Analysis: All methods implement RGB MSE loss, but AutoP2C's implementation is overly generic without ray-based context.

Component C5: YUV-UV Loss Variant - All Baselines

LineMOD variant uses only U and V channels in YUV space with BT.709 conversion matrix.

```

5065 1 # \nerify\ : YUV-UV loss with BT.709 standard
5066 2 if self.config.use_yuv_uv_loss:
5067   3 yuv_gt = rgb_to_yuv_bt709(image)  # BT.709
5068     coefficients
5069   4 yuv_pr = rgb_to_yuv_bt709(pred)
5070   5 gt_uv = yuv_gt[:, :, 1:]  # U and V only
5071   6 pr_uv = yuv_pr[:, :, 1:]
5072   7 uv_loss = torch.mean((gt_uv - pr_uv) ** 2)
5073 8 # CORRECT: Exact paper specification

```

Listing 359. C5 Implementation: NERFIFY (Score: 1.0)

```

5074 1 # GPT-5: No YUV loss implementation
5075 2 # MISSING: Only RGB MSE, no YUV variant

```

Listing 360. C5 Implementation: GPT-5 (Score: 0.0)

```

5076 1 # DeepSeek R1: YUV with different coefficients
5077 2 elif loss_type == 'yuv':
5078   3 rgb_to_yuv = torch.tensor([
5079     [0.299, 0.587, 0.114],  # ISSUE: BT.601, not
5080       BT.709
5081     [-0.14713, -0.28886, 0.436],
5082     [0.615, -0.51499, -0.10001]
5083   ])
5084 7 loss = torch.mean((rendered_yuv[:, :, 1:] -
5085     observed_yuv[:, :, 1:]) ** 2)
5086 9 # ERROR: Wrong color space standard

```

Listing 361. C5 Implementation: DeepSeek R1 (Score: 0.8)

```

5087 1 # Paper2Code: Correct BT.709 YUV conversion
5088 2 conversion_matrix = np.array(
5089   [[0.2126, 0.7152, 0.0722],  # CORRECT: BT.709
5090     [-0.09991, -0.33609, 0.436],
5091     [0.615, -0.55861, -0.05639]])
5092 6 yuv_image = torch.tensordot(image, conversion_matrix,
5093   T, dims=([2], [1]))

```

Listing 362. C5 Implementation: Paper2Code (Score: 1.0)

```

5094 1 # AutoP2C: No YUV implementation
5095 2 # MISSING: No color space conversion

```

Listing 363. C5 Implementation: AutoP2C (Score: 0.0)

5112 Analysis: NERFIFY and Paper2Code correctly implement BT.709 YUV conversion. DeepSeek R1 uses wrong coefficients (BT.601). GPT-5 and AutoP2C lack YUV loss entirely.

5116 Component C6: Camera LR Schedule - All Baselines

5117 Paper specifies exponential decay: $\alpha_t = \alpha_0 \cdot 0.8^{t/100}$.

```
5118 1 # \nerfify\ : Exact formula via training callback
5119 2 def set_camera_lr(step: int) -> None:
5120 3     lr = self._camera_lr0 * (0.8 ** (step / 100.0))
5121 4     for g in optim.param_groups:
5122 5         g["lr"] = lr
5123 6 # CORRECT: Precise paper formula
```

Listing 364. C6 Implementation: NERFIFY (Score: 1.0)

```
5126 1 # GPT-5: Fixed learning rate
5127 2 optimizer = torch.optim.Adam([pose], lr=lr)
5128 3 # MISSING: No learning rate schedule
```

Listing 365. C6 Implementation: GPT-5 (Score: 0.0)

```
5131 1 # DeepSeek R1: Wrong exponential schedule
5132 2 self.scheduler = torch.optim.lr_scheduler.
5133     ExponentialLR(
5134         self.optimizer, gamma=0.8)
5135 4 # ERROR: gamma^t instead of gamma^(t/100)
```

Listing 366. C6 Implementation: DeepSeek R1 (Score: 0.5)

```
5138 1 # Paper2Code: Correct custom schedule
5139 2 def get_learning_rate(initial_lr, decay_factor,
5140     decay_steps, iteration):
5141     lr = initial_lr * (decay_factor ** (iteration /
5142         decay_steps))
5143 4 return lr
5144 5 # CORRECT: Matches paper formula
```

Listing 367. C6 Implementation: Paper2Code (Score: 1.0)

```
5147 1 # AutoP2C: No learning rate schedule
5148 2 optimizer = optim.Adam([pose], lr=0.01)
5149 3 # MISSING: Fixed LR throughout
```

Listing 368. C6 Implementation: AutoP2C (Score: 0.0)

5152 Analysis: Only NERFIFY and Paper2Code implement the exact LR schedule. DeepSeek R1 uses wrong decay rate. GPT-5 and AutoP2C have no scheduling.

5155 Component C7: Small Ray Batch Size - All Baselines

5156 Paper uses 2048 rays per optimization step for efficiency.

```
5157 1 # \nerfify\ : Exact batch size
5158 2 datamanager=TemplateDataManagerConfig(
5159 3     train_num_rays_per_batch=2048, # CORRECT: Paper
5160     value
5161 4     eval_num_rays_per_batch=4096,
5162 5 )
```

Listing 369. C7 Implementation: NERFIFY (Score: 1.0)

```
5165 1 # GPT-5: Correct batch size
5166 2 def interest_region_sampling(image, num_rays=2048):
5167 3 # CORRECT: 2048 rays
```

Listing 370. C7 Implementation: GPT-5 (Score: 1.0)

```
5170 1 # DeepSeek R1: Wrong default
5171 2 def __init__(self, batch_size=512): # ERROR: Default
5172 3     512
5173 3     self.batch_size = batch_size
```

Listing 371. C7 Implementation: DeepSeek R1 (Score: 0.5)

```
5176 1 # Paper2Code: Correct from config
5177 2 self.batch_size = config["ray_sampling"].get("batch_size", 2048)
5178 3 # CORRECT: 2048 default
```

Listing 372. C7 Implementation: Paper2Code (Score: 1.0)

```
5182 1 # AutoP2C: No batch control
5183 2 rays_world = self.transform_rays(directions, pose)
5184 3 # CRITICAL ERROR: Processes full image
```

Listing 373. C7 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY , GPT-5, and Paper2Code use correct batch size. DeepSeek R1 defaults to 512. AutoP2C processes entire images.

Component C8: Gradient-based Pose Updates - All Baselines

All methods use gradient descent, but implementation quality varies.

```
5194 1 # \nerfify\ : Proper gradient flow through volume
5195     rendering
5196 2 loss = self.get_loss_dict(outputs, batch)
5197 3 loss.backward()
5198 4 optimizer.step()
5199 5 # CORRECT: Full differentiable pipeline
```

Listing 374. C8 Implementation: NERFIFY (Score: 1.0)

```
5202 1 # Standard gradient descent present in all baselines
5203 2 loss.backward()
5204 3 optimizer.step()
```

Listing 375. C8 Implementation: All Others (Score: 0.8-1.0)

Analysis: All methods implement basic gradient descent, with varying quality in the overall pipeline.

Component C9: Volume Rendering Integration - All Baselines

Standard NeRF volume rendering with alpha compositing.

```
5213 1 # \nerfify\ : Uses Nerfacto's volume renderer
5214 2 # Inherits complete volume rendering pipeline
```

Listing 376. C9 Implementation: NERFIFY (Score: 1.0)

```
5217 1 # All baselines have some volume rendering
5218 2 alpha = 1.0 - torch.exp(-sigma * deltas)
5219 3 weights = alpha * transmittance
5220 4 final_color = torch.sum(weights * rgb, dim=1)
```

Listing 377. C9 Implementation: Others (Score: 0.6-1.0)

Analysis: All methods include volume rendering with varying completeness.

Component C10: Pose-only Optimization Mode - All Baselines

5227 Freezing NeRF weights and only optimizing camera
5228 poses.

```
5229 1 # \nerfify\ : Explicit pose-only mode
5230 2 if self.config.invert_poses_only:
5231 3     for p in self.parameters():
5232 4         p.requires_grad = False
5233 5     for p in self.camera_optimizer.parameters():
5234 6         p.requires_grad = True
5235 7 # CORRECT: Clear separation
```

Listing 378. C10 Implementation: NERFIFY (Score: 1.0)

5236 1 # No explicit pose-only mode in other baselines
5237 2 # MISSING: All assume fixed NeRF implicitly

Listing 379. C10 Implementation: Others (Score: 0.0)

5242 **Analysis:** Only NERFIFY has explicit pose-only optimization
5243 control.

5244 3.8.6. Scoring Analysis

Table 63. iNeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	Weighted Avg _{LLM}
NERFIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.7	0.97
GPT-5	0.8	0.4	0.5	1.0	0.0	0.0	1.0	1.0	0.8	0.0	0.58
DeepSeek R1	1.0	1.0	1.0	1.0	0.8	0.5	0.5	1.0	1.0	0.0	0.68
Paper2Code	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.8	0.0	0.75
AutoP2C	0.0	0.0	0.0	0.5	0.0	0.0	0.3	0.6	0.0		0.05

5245 3.8.7. Why Baselines Fail Despite Component Scores

5246 GPT-5 (Score: 0.58 components, 0% trainable)

- 5247 • **Strengths:** Implements SE(3) structure, basic volume
5248 rendering, correct batch size

- 5249 • **Fatal Issues:**

- 5250 – No Nerfstudio integration - standalone script cannot be
5251 deployed
- 5252 – Missing small angle handling causes numerical instability
- 5253 – Simplified gradient computation without proper edge
5254 detection
- 5255 – No learning rate schedule leads to poor convergence

- 5256 • **Result:** Code runs but produces unstable optimization
5257 with divergent poses

5258 DeepSeek R1 (Score: 0.68 components, 0% trainable)

- 5259 • **Strengths:** Complete SE(3) implementation, proper interest
5260 sampling, YUV loss variant

- 5261 • **Fatal Issues:**

- 5262 – No framework integration - requires complete rewrite
5263 for Nerfstudio
- 5264 – Wrong LR schedule causes 100x faster decay than intended
- 5265 – Default batch size mismatch reduces gradient quality

- 5266 • **Result:** Mathematically sound but practically unusable
5267 without major refactoring

5268 Paper2Code (Score: 0.75 components, 0% trainable)

- 5269 • **Strengths:** Most complete component implementation,
5270 correct hyperparameters

- 5271 • **Fatal Issues:**

- 5272 – 2100+ lines of framework-agnostic code incompatible
5273 with Nerfstudio
- 5274 – Over-engineered module structure prevents simple integration
- 5275 – Requires specific config.yaml format not matching Nerf
5276 studio conventions

- 5277 • **Result:** High-quality standalone system that cannot be
5278 deployed as intended plugin

5279 AutoP2C (Score: 0.05 components, 0% trainable)

- 5280 • **Strengths:** None - fundamental algorithmic failures
5281 throughout

- 5282 • **Fatal Issues:**

- 5283 – No SE(3) manifold constraints - optimizes invalid rotation
5284 matrices
- 5285 – Missing all sampling strategies - processes entire images
- 5286 – Import errors prevent compilation (non-existent
5287 colmap_wrapper)
- 5288 – Empty placeholder functions throughout codebase

- 5289 • **Result:** Code neither compiles nor demonstrates any do
5290 main understanding

5291 3.8.8. Hyperparameter Fidelity

Table 64. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5	R1	P2C	AutoP2C
Ray batch size	2048	✓	✓	✗	✓	✗
Initial LR (α_0)	0.01	✓	✗	✓	✓	✓
LR decay factor	0.8	✓	✗	✗	✓	✗
LR decay steps	100	✓	✗	✗	✓	✗
Optimizer	Adam	✓	✓	✓	✓	✓
Adam β_1	0.9	✓	✗	✓	✗	✗
Adam β_2	0.999	✓	✗	✓	✗	✗
Edge threshold	0.2	✓	✗	✓	✓	✗
Dilation iterations	2	✓	✗	✓	✓	✗
Kernel size	5x5	✓	✗	✓	✓	✗
W Score	–	1.00	0.20	0.60	0.80	0.10

5292 3.8.9. Conclusion

5293 All baselines attempt to implement iNeRF with varying so
5294 phistication. Paper2Code achieves the highest component
5295 coverage (75%) with correct SE(3) parameterization and
5296 complete interest sampling but produces framework-agnostic
5297 code. DeepSeek R1 implements core algorithms correctly
5298 (68%) but uses wrong hyperparameters that prevent conver
5299 gence. GPT-5 provides a functional skeleton (58%) missing
5300 critical details like small angle handling and proper edge
5301 detection. AutoP2C completely fails (5%) with no SE(3)
5302 constraints, missing sampling strategies, and import errors.
5303 Only NERFIFY produces immediately trainable code as a
5304 complete Nerfstudio plugin with 97% component fidelity
5305 and exact hyperparameter matching, demonstrating that the
5306 0% trainable rate for baselines versus 100% for NERFIFY val
5307 idates the necessity of domain-specific synthesis for complex
5308 vision research.

5312 **3.9. SIGNeRF: Scene Integrated Generation for** 5327
 5313 **Neural Radiance Fields**

5314 **3.9.1. Paper Overview**

5315 SIGNeRF introduces a framework for editing neural radiance fields by generating multi-view consistent reference sheets through depth-conditioned ControlNet inpainting.
 5316 The method enables scene-integrated generation by first creating a reference sheet with multiple views arranged in a
 5317 grid, processing them through ControlNet with depth conditioning, and then propagating the edits to all training views.
 5318 This addresses the critical challenge of maintaining 3D consistency when editing NeRF scenes with diffusion models.
 5319

5320 **3.9.2. Implementation Overview**

Table 65. SIGNeRF Implementation Summary Across All Baselines

Aspect	NERFIFY	GPT-5	DeepSeek R1	Paper2Code	AutoP2C
Lines of Code	482	876	624	1842	658
File Organization	Plugin	Single	Single	Multi-file	Multi-file
Reference Sheet	✓	✓	Partial	✓	✗
ControlNet Integration	✓	✓	Wrong Type	✓	✗
Dataset Updating	✓	Partial	Simplified	✓	✗
Two-Stage Training	✓	✓	Simplified	Partial	✗
Selection Modes	✓	Partial	Partial	Partial	✗
Trainable	✓	✗	✗	✗	✗

Note: NERFIFY always produces trainable code as a complete Nerfstudio plugin

5325 **3.9.3. Novel Components**

Table 66. Novel Components in SIGNeRF with Importance Weights

ID	Component	Weight w_i
C1	Reference sheet generation with grid layout	0.20
C2	ControlNet depth-conditioned inpainting	0.20
C3	Multi-view dataset updating	0.15
C4	Two-stage NeRF training (initial + fine-tune)	0.10
C5	Selection modes (bounding box and proxy)	0.10
C6	Iterative refinement with parameter adjustment	0.10
C7	ControlNet hyperparameters (scale, guidance)	0.08
C8	LPIPS loss for consistency	0.07

5326 **3.9.4. Quantitative Metrics**

Table 67. SIGNeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	1.00	0.00	0.00	1.00	1.00
GPT-5	0.50	0.25	0.25	0.63	0.58
DeepSeek R1	0.63	0.25	0.12	0.75	0.72
Paper2Code	0.38	0.38	0.24	0.50	0.52
AutoP2C	0.00	0.13	0.87	0.00	0.08

5327 **3.9.5. Component-by-Component Analysis**

5328 **Component C1: Reference Sheet Generation - All Baselines**

5329 The paper specifies creating a reference sheet by arranging M views in a grid with one empty slot:

$$\mathcal{R} \leftarrow \text{ControlNet}(\bar{\mathcal{I}}_R, \bar{\mathcal{D}}_R, \bar{\mathcal{M}}_R, y) \quad (1)$$

```
1 # \nerify\ : Complete grid generation with
  Nerfstudio integration
2 class ControlNetEditor:
3     def create_reference_sheet(self, rgb_images,
4                               depth_images,
5                               masks, prompt,
6                               grid_size=(2, 3)):
7         """Create edited reference sheet with M views
8         in grid."""
9         # Arrange images into grid with empty slot
10        reference_sheet = self._arrange_grid(
11            rgb_images, depth_images, masks,
12            grid_size
13        )
14        # Process with ControlNet
15        edited_sheet = self._run_controlnet_inpainting(
16            reference_sheet['rgb_grid'],
17            reference_sheet['depth_grid'],
18            reference_sheet['mask_grid'], prompt
19        )
20        return edited_sheet
21 # CORRECT: Proper grid layout with empty slot
```

Listing 380. C1 Implementation: NERFIFY (Score: 1.0)

```
1 # GPT-5: Standalone implementation
2 def generate_reference_sheet(color_imgs, depth_imgs,
3                               mask_imgs, prompt, ncols
4                               =3):
5     """Combine M views into tiled reference sheet."""
6     grid_color = make_grid(colors, ncols)
7     grid_depth = make_grid(depths, ncols)
8     grid_mask = make_grid([cv2.merge([m,m,m]) for m in masks], ncols)
9     gen = pipe(prompt=prompt, image=grid_color,
10                control_image=grid_depth, mask_image=
11                grid_mask)
12    return gen.images[0]
13 # ERROR: Not integrated with Nerfstudio framework
```

Listing 381. C1 Implementation: GPT-5 (Score: 0.8)

```
1 # DeepSeek R1: Simplified grid generation
2 def generate_reference_sheet(self, prompt,
3                               selection_mode,
4                               num_reference_views=5):
5     reference_cameras = self._select_reference_cameras(
6         num_reference_views, selection_params)
7     color_grid, depth_grid, mask_grid = [], [], []
8     for cam in reference_cameras:
9         color, depth = self.nerf_model.render(cam)
10        mask = self._generate_mask(cam,
11                                  selection_mode,
12                                  selection_params,
13                                  depth)
14        color_grid.append(color)
15        grid_size = int(np.ceil(np.sqrt(
16            num_reference_views)))
17        # ISSUE: Missing actual ControlNet processing
18        reference_sheet = self._create_reference_sheet(
19            color_grid, depth_grid, mask_grid, grid_size,
20            prompt)
```

Listing 382. C1 Implementation: DeepSeek R1 (Score: 0.7)

```

5397
5398 1 # Paper2Code: Complete but standalone
5399 2 def generate_reference_sheet(self, image_grid,
5400     depth_grid,
5401         mask_grid, prompt=None):
5402     """Generate reference sheet using diffusion
5403     pipeline."""
5404     image_pil = self._np_to_pil(image_grid)
5405     depth_pil = self._np_to_pil(depth_grid)
5406     mask_pil = self._np_to_pil(mask_grid)
5407     generator = torch.Generator(device=self.device)._
5408         manual_seed(
5409             self.seed)
5410     result = self.pipeline(prompt=prompt, image=
5411         image_pil,
5412             mask_image=mask_pil,
5413
5414         controlnet_conditioning_image=depth_pil)
5415     return self._pil_to_np(result.images[0])
5416 # ERROR: Not a Nerfstudio plugin

```

Listing 383. C1 Implementation: Paper2Code (Score: 1.0)

```

5418
5419 1 # AutoP2C: No reference sheet generation
5420 2 class BlenderDataset:
5421     def load_data(self):
5422         """Loads images and camera poses."""
5423         with open('transforms.json', 'r') as f:
5424             metadata = json.load(f)
5425             self.camera_poses = [np.array(frame['
5426                 transform_matrix'])]
5427             for frame in metadata[
5428                 'frames']]
5429 # MISSING: No reference sheet or ControlNet
5430 implementation

```

Listing 384. C1 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY and Paper2Code correctly implement grid-based reference sheet generation, but only NERFIFY integrates it with Nerfstudio. GPT-5 has correct logic but lacks framework integration. DeepSeek R1 has the structure but misses actual ControlNet processing. AutoP2C completely lacks this component.

Component C2: ControlNet Depth-Conditioned In-painting - All Baselines

The paper specifies using SDXL with depth ControlNet for multi-view consistency.

```

5442
5443 1 # \nerfify\ : Correct depth ControlNet setup
5444 2 from diffusers import (
5445     StableDiffusionXLControlNetInpaintPipeline,
5446         ControlNetModel)
5447 4 self.controlnet_editor = ControlNetEditor(
5448     controlnet_scale=self.config.controlnet_scale,
5449         # 0.7
5450     guidance_scale=self.config.guidance_scale,
5451         # 7.5
5452     denoising_strength=self.config.denoising_strength
5453         # 0.8
5454 )
5455 # CORRECT: Uses depth ControlNet as specified

```

Listing 385. C2 Implementation: NERFIFY (Score: 1.0)

```

5457
5458 1 # GPT-5: Correct depth ControlNet
5459 2 controlnet = ControlNetModel.from_pretrained(
5460     "diffusers/controlnet-depth-sdxl-1.0",
5461     torch_dtype=torch.float16)
5462 5 self.pipe =
5463     StableDiffusionXLControlNetInpaintPipeline\
5464         .from_pretrained("stabilityai/stable-diffusion-xl
5465         -base-1.0",

```

```

7
8         controlnet=controlnet,
9             torch_dtype=torch.float16).to(
device)
# CORRECT: Proper depth conditioning

```

Listing 386. C2 Implementation: GPT-5 (Score: 1.0)

```

1 # DeepSeek R1: Wrong ControlNet type
2 if self.config['use_sdxl']:
3     controlnet = ControlNetModel.from_pretrained(
4         "diffusers/controlnet-canny-sdxl-1.0", #
WRONG!
5         torch_dtype=torch.float16)
6 # CRITICAL ERROR: Uses Canny instead of Depth

```

Listing 387. C2 Implementation: DeepSeek R1 (Score: 0.5)

```

1 # Paper2Code: Correct but hardcoded parameters
2 self.pipeline =
3     StableDiffusionControlNetInpaintPipeline\
4         .from_pretrained("stabilityai/stable-diffusion-xl
-base-1.0",
5             torch_dtype=torch.float16,
6                 revision="fp16")
7 self.controlnet_scale = float(
8     diffusion_config.get("controlnet_scale", [0.4,
1.0])[0]
) # Uses 0.4 instead of paper's 0.7-0.8
# ISSUE: Wrong hyperparameters

```

Listing 388. C2 Implementation: Paper2Code (Score: 0.9)

```

1 # AutoP2C: No ControlNet implementation
2 class ControlNet(nn.Module):
3     def __init__(self, latent_dim=512, control_scale
=0.7):
4         super(ControlNet, self).__init__()
5         self.encoder = self.create_encoder()
6         # Just a placeholder NN module
7 # MISSING: No actual diffusion model integration

```

Listing 389. C2 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY and GPT-5 correctly use depth ControlNet. DeepSeek R1 critically uses the wrong ControlNet type (Canny), breaking 3D consistency. Paper2Code has correct implementation but wrong hyperparameters. AutoP2C has no real ControlNet.

Component C3: Multi-View Dataset Updating - All Baselines

The paper describes propagating edits to all views by compositing each into the reference sheet.

```

1 # \nerfify\ : Complete view propagation
2 def propagate_to_views(self, reference_sheet,
3     view_rgbs,
4         view_depths, view_masks, prompt
5             ):
6     edited_views = []
7     for rgb, depth, mask in zip(view_rgbs,
8         view_depths, view_masks):
8         composite_sheet = self._replace_grid_slot(
9             reference_sheet, rgb, depth, mask,
10             slot_index=-1)
11         edited_view = self._run_controlnet_inpainting(
12             (
13                 composite_sheet['rgb_grid'],
14                 composite_sheet['depth_grid'],
15                 composite_sheet['mask_grid'], prompt
16             )
17             edited_slot = self._extract_grid_slot(
18                 edited_view,

```

```

5532
5533     slot_index=-1)
5534     edited_views.append(edited_slot)
5535
5536     return edited_views
5537
5538 # CORRECT: Proper view compositing and extraction

```

Listing 390. C3 Implementation: NERFIFY (Score: 1.0)

```

5539
5540     # GPT-5: Dataset update but saves to disk
5541     def update_dataset(ref_sheet_path, color_imgs,
5542                         depth_imgs,
5543                         mask_imgs, prompt, out_dir):
5544         ref_sheet = Image.open(ref_sheet_path).convert("RGB")
5545         for idx, (ci, di, mi) in enumerate(zip(color_imgs,
5546                                         depth_imgs,
5547                                         mask_imgs)):
5548             grid = make_grid([np.array(ref_sheet), np.
5549                             array(color)],
5550                             ncols=2)
5551             gen = pipe(prompt=prompt, image=grid,
5552                         control_image=depth,
5553                         mask_image=mask).images[0]
5554             gen.save(os.path.join(out_dir, f"view_{idx:04d}.png"))
5555
5556 # ERROR: Saves to disk instead of returning for
5557 # training

```

Listing 391. C3 Implementation: GPT-5 (Score: 0.7)

```

5560
5561     # DeepSeek R1: Incomplete update logic
5562     def update_image_set(self, reference_sheet):
5563         for i, cam in self.camera_poses.items():
5564             color, depth = self.nerf_model.render(cam)
5565             mask = self._get_camera_mask(cam)
5566             # MISSING: No actual ControlNet propagation
5567             edited_view = reference_sheet # WRONG!
5568             self.edited_images[i] = edited_view

```

Listing 392. C3 Implementation: DeepSeek R1 (Score: 0.4)

```

5570
5571     # Paper2Code: Complete implementation
5572     def update_view(self, reference_sheet, view_rgb,
5573                     view_depth,
5574                     view_mask, prompt):
5575         composite_rgb = self._inject_view_into_grid(
5576             self.ref_sheet, view_rgb, self.
5577             empty_slot_position)
5578         result = self.pipeline(prompt=prompt, image=
5579             composite_rgb_pil,
5580                         mask_image=
5581             composite_mask_pil,
5582
5583             controlnet_conditioning_image=
5584             composite_depth_pil)
5585         updated_view = updated_composite_np[y_start:y_end
5586
5587             ,
5588             x_start:x_end,
5589             :]
5590     return updated_view
5591
5592 # ERROR: Not integrated with Nerfstudio training

```

Listing 393. C3 Implementation: Paper2Code (Score: 1.0)

```

5593
5594     # AutoP2C: No dataset updating
5595     def update_dataset(ref_sheet_path, view_list, prompt)
5596         :
5597         # Function defined but no implementation
5598         pass
5599 # MISSING: No implementation at all

```

Listing 394. C3 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY and Paper2Code correctly implement view propagation with compositing. GPT-5 has the logic but saves to disk. DeepSeek R1 lacks actual ControlNet propagation. AutoP2C has no implementation.

Component C4: Two-Stage NeRF Training - All Baselines

The paper specifies training initial NeRF then fine-tuning on edited data.

```

1 # \nerify\ : Integrated two-stage with Nerfstudio
2 class SIGNeRFPipeline(VanillaPipeline):
3     def get_train_loss_dict(self, step: int):
4         ray_bundle, batch = self.datamanager.
5         next_train(step)
6         model_outputs = self.model(ray_bundle)
7         loss_dict = self.model.get_loss_dict(
8             model_outputs, batch, metrics_dict)
9         return model_outputs, loss_dict, metrics_dict
10 # Config: initial_training: 30k iterations
11 #           fine_tuning: load checkpoint, continue
12 # CORRECT: Proper two-stage pipeline

```

Listing 395. C4 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Explicit two stages but standalone
2 def train(self, data_dir, output_name="nerf", steps
3           =30000):
4     config = TrainerConfig(method_name="nerfacto",
5                           max_num_iterations=steps,
5                           data=data_dir)
6     trainer = Trainer(config)
7     trainer.train()
8
9 def finetune(self, pretrained_dir, edited_data_dir,
10            steps=10000):
11    config = TrainerConfig(load_dir=pretrained_dir,
12                           max_num_iterations=steps,
13                           data=edited_data_dir)
14
15 # ERROR: Not a Nerfstudio plugin

```

Listing 396. C4 Implementation: GPT-5 (Score: 0.9)

```

1 # DeepSeek R1: Simplified placeholder
2 def train_nerf(self, images, camera_poses):
3     self.nerf_model.train_model(images, camera_poses)
4
5 def refine_nerf(self, num_iterations=1000):
6     self.nerf_model.fine_tune(edited_images_list,
7                               camera_poses_list,
8                               num_iterations)
9 # ISSUE: SimplifiedNeRF is dummy implementation

```

Listing 397. C4 Implementation: DeepSeek R1 (Score: 0.6)

```

1 # Paper2Code: Generic training only
2 class NeRFTrainer:
3     def train(self):
4         for iter_idx in range(1, self.iterations + 1)
5             :
5         batch = self.training_data[iter_idx %
6         num_batches]
6         prediction = self.model(batch)
7         loss = F.mse_loss(prediction, batch)
8
9 # MISSING: No fine-tuning stage

```

Listing 398. C4 Implementation: Paper2Code (Score: 0.4)

```

1 # AutoP2C: Generic NeRF training
2 class Trainer:
3     def train(self):
4         for epoch in range(self.num_epochs):

```

```

5667         predictions = self.model(rays['coords'],
5668             rays['view_dirs'])
5669             loss = self.compute_loss(predictions,
5670                 batch['images'])
5671 # MISSING: No two-stage process at all

```

Listing 399. C4 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERFIFY properly integrates two-stage training with Nerfstudio. GPT-5 has explicit stages but standalone. DeepSeek R1 has placeholder implementation. Paper2Code and AutoP2C lack fine-tuning.

Component C5: Selection Modes - All Baselines

The paper specifies bounding box and proxy mesh selection modes.

```

1 # \nerfify\ : Complete selection modes
2 def generate_mask(self, depth_map, selection):
3     mode = selection.get("mode", self.default_mode)
4     if mode == "bounding_box":
5         bbox = selection["bbox"]
6         region = depth_map[bbox["y_min"] : bbox["y_max"]
7             ],
7             bbox["x_min"] : bbox["x_max"]
8         mask = (depth_map >= th_min) & (depth_map <
9             th_max)
10    elif mode == "proxy":
11        proxy_mask = selection.get("proxy_mask")
12        proxy_depth = selection.get("proxy_depth")
13        # Full implementation for both options
13 # CORRECT: Both modes fully implemented

```

Listing 400. C5 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Basic selection modes
2 def generate_masks_from_bbox(depths, bbox):
3     x1,y1,x2,y2 = bbox
4     masks=[]
5     for d in depths:
6         m=np.zeros_like(d,dtype=np.uint8)
7         m[y1:y2,x1:x2]=255
8         masks.append(m)
9     def proxy_mask_from_mesh(depths, proxy_depths):
10        masks.append(((p>0) & (p<d)).astype(np.uint8)*255
11 # ISSUE: Simplified, not integrated

```

Listing 401. C5 Implementation: GPT-5 (Score: 0.7)

```

1 # DeepSeek R1: Both modes but simplified
2 def _shape_selection_mask(self, camera,
3     selection_params,
4         depth_map, h, w):
5     bbox = selection_params['bounding_box']
6     # Project bounding box to image space
6     def _proxy_selection_mask(self, camera,
7         selection_params,
8             depth_map, h, w):
9     proxy_mesh = selection_params['proxy_mesh']
10    # Render proxy mesh from camera view
10 # ISSUE: Incomplete implementations

```

Listing 402. C5 Implementation: DeepSeek R1 (Score: 0.7)

```

1 # Paper2Code: Only bounding box mode
2 if mode == "bounding_box":
3     bbox = selection["bbox"]
4     region = depth_map[y_min:y_max, x_min:x_max]
5     mask_bool = (depth_map >= th_min) & (depth_map <=
6         th_max)
6 # MISSING: No proxy mode implementation

```

Listing 403. C5 Implementation: Paper2Code (Score: 0.3)

```

1 # AutoP2C: No selection modes
2 # No implementation found

```

Listing 404. C5 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY fully implements both selection modes. GPT-5 and DeepSeek R1 have both but simplified. Paper2Code only has bounding box. AutoP2C has none.

Component C6: Iterative Refinement - All Baselines

The paper describes iterative refinement with parameter adjustment.

```

1 # \nerfify\ : Full iterative refinement support
2 # Config allows multiple iterations:
3 trainer_config = TrainerConfig(
4     pipeline=SIGNeRFPipelineConfig(
5         model=SIGNeRFModelConfig(
6             use_lpips_loss=True,
7             lpips_loss_mult=0.01,
8             depth_supervision_mult=0.1
9         )
10    )
11 # CORRECT: Supports iterative refinement

```

Listing 405. C6 Implementation: NERFIFY (Score: 1.0)

```

1 # GPT-5: Optional second iteration
2 if second_iter:
3     self.ref_gen.controlnet_scale=1.0
4     self.ref_gen.denoising_strength=0.5
5     ref_sheet2=os.path.join(base_dir, "reference_sheet_iter2.png")
6     self.ref_gen.generate_reference_sheet(...)
7     self.trainer.finetune...
8 # ISSUE: Only supports one refinement

```

Listing 406. C6 Implementation: GPT-5 (Score: 0.6)

```

1 # DeepSeek R1: Iterative refinement method
2 def iterative_refinement(self, num_iterations=2):
3     for iteration in range(num_iterations):
4         if iteration > 0:
5             self.config['controlnet_scale'] = min
6             (1.0,
7                 self.config.get('controlnet_scale',
8                     0.8) + 0.2)
7             reference_sheet = self.
8             generate_reference_sheet...
8             self.update_image_set(reference_sheet)
9             self.refine_nerf(500)
10 # CORRECT: Good iterative structure

```

Listing 407. C6 Implementation: DeepSeek R1 (Score: 0.8)

```

1 # Paper2Code: No iterative refinement
2 # Single pass only, no iteration support

```

Listing 408. C6 Implementation: Paper2Code (Score: 0.0)

```

1 # AutoP2C: No iterative refinement
2 # No implementation found

```

Listing 409. C6 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY supports full iterative refinement through config. DeepSeek R1 has good iterative structure. GPT-5 only supports one refinement. Paper2Code and AutoP2C lack this feature.

5796 Component C7: ControlNet Hyperparameters - All
5797 Baselines

The paper specifies $\text{controlnet_scale} = 0.4 - 1.0$, $\text{guidance_scale} = 6 - 10$, $\text{denoising_strength} = 0.5 - 0.95$.

```
5798
5799 1 # \nerify\ : Correct hyperparameters
5800 2 self.controlnet_editor = ControlNetEditor(
5801 3     controlnet_scale=self.config.controlnet_scale,
5802 4         # 0.7
5803 5     guidance_scale=self.config.guidance_scale,
5804 6         # 7.5
5805 7     denoising_strength=self.config.denoising_strength
5806 8         # 0.8
5807
5808 9 # CORRECT: All within paper ranges
```

Listing 410. C7 Implementation: NERIFY (Score: 1.0)

```
5810
5811 1 # GPT-5: Correct parameters
5812 2 self.controlnet_scale = controlnet_scale # 0.8
5813 3 self.guidance_scale = guidance_scale # 7.0
5814 4 self.denoising_strength = denoising_strength # 0.8
5815 5 # CORRECT: All within specified ranges
```

Listing 411. C7 Implementation: GPT-5 (Score: 1.0)

```
5816
5817 1 # DeepSeek R1: Mostly correct
5818 2 controlnet_conditioning_scale=self.config.get('
5819     controlnet_scale', 0.8)
5820 3 guidance_scale=self.config.get('guidance_scale', 7.5)
5821 4 num_inference_steps=15 # Too few, paper uses 40-50
5822 5 # ISSUE: Inference steps too low
```

Listing 412. C7 Implementation: DeepSeek R1 (Score: 0.8)

```
5823
5824 1 # Paper2Code: Wrong values from config
5825 2 self.controlnet_scale = float(controlnet_scale_range
5826     [0]) # 0.4
5827 3 self.guidance_scale = float(guidance_scale_range[0])
5828 4 self.denoising_strength = float(denoising_range[0])
5829 5 # ERROR: Takes first value from ranges, suboptimal
```

Listing 413. C7 Implementation: Paper2Code (Score: 0.5)

```
5830
5831 1 # AutoP2C: No ControlNet parameters
5832 2 control_scale = 0.7 # Defined but never used
5833 3 # MISSING: No actual ControlNet usage
```

Listing 414. C7 Implementation: AutoP2C (Score: 0.0)

Analysis: NERIFY and GPT-5 use correct hyperparameters. DeepSeek R1 mostly correct but inference steps too low. Paper2Code uses suboptimal values. AutoP2C has no real parameters.

5844 Component C8: LPIPS Loss for Consistency - All
5845 Baselines

The paper specifies using LPIPS loss for enhanced texture consistency during fine-tuning.

```
5846
5847 1 # \nerify\ : LPIPS loss integrated
5848 2 from torchmetrics.image.lpip import
5849     LearnedPerceptualImagePatchSimilarity
5850 3 if self.config.use_lips_loss:
5851 4     self.lips =
5852         LearnedPerceptualImagePatchSimilarity(normalize=
5853             True)
```

```
5854
5855 5     lips_loss = self.lips(pred_rgb_permuted,
5856     gt_rgb_permuted)
5857 6     loss_dict["lips_loss"] = self.config.
5858     lips_loss_mult * lips_loss
5859 7 # CORRECT: Proper LPIPS integration
```

Listing 415. C8 Implementation: NERIFY (Score: 1.0)

```
5860
5861 1 # GPT-5: No LPIPS loss
5862 2 # Uses standard MSE loss only
5863 3 # MISSING: No LPIPS implementation
```

Listing 416. C8 Implementation: GPT-5 (Score: 0.0)

```
5864
5865 1 # DeepSeek R1: No LPIPS loss
5866 2 # Standard NeRF losses only
5867 3 # MISSING: No LPIPS implementation
```

Listing 417. C8 Implementation: DeepSeek R1 (Score: 0.0)

```
5868
5869 1 # Paper2Code: No LPIPS in training
5870 2 # Only mentions LPIPS in evaluation metrics
5871 3 # MISSING: Not used in training loss
```

Listing 418. C8 Implementation: Paper2Code (Score: 0.0)

```
5872
5873 1 # AutoP2C: LPIPS defined but misused
5874 2 def lips_loss(predictions, targets, lips_model):
5875 3     lips_value = lips_model(predictions, targets)
5876 4     return lips_value.mean()
5877 5 # ERROR: Defined but never integrated in training
```

Listing 419. C8 Implementation: AutoP2C (Score: 0.5)

Analysis: Only NERIFY properly integrates LPIPS loss in training. AutoP2C defines it but doesn't use it. Others completely lack LPIPS implementation.

3.9.6. Scoring Analysis

Table 68. SIGNeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	Weighted AvgLLM
NERIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.00
GPT-5	0.8	1.0	0.7	0.9	0.7	0.6	1.0	0.0	0.71
DeepSeek R1	0.7	0.5	0.4	0.6	0.7	0.8	0.8	0.0	0.56
Paper2Code	1.0	0.9	1.0	0.4	0.3	0.0	0.5	0.0	0.51
AutoP2C	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.5	0.06

3.9.7. Why Baselines Fail Despite Component Scores

GPT-5 (Score: 0.58 overall, 0% trainable)

- Strengths:** Correct depth ControlNet, proper hyperparameters, explicit two-stage training
- Fatal Issues:**
 - No Nerfstudio integration - standalone script only
 - Saves edited images to disk instead of training pipeline
 - Missing LPIPS loss for consistency
- Result:** Cannot be loaded as ‘ns-train signerf’, requires 3-5 hours manual integration

DeepSeek R1 (Score: 0.72 overall, 0% trainable)

- Strengths:** Good iterative refinement structure, both selection modes

- 5901 • **Fatal Issues:**
- Uses Canny ControlNet instead of Depth - fundamentally breaks 3D consistency
 - SimplifiedNeRF is placeholder with no real implementation
 - Missing actual ControlNet propagation in dataset update
- 5902 • **Result:** Wrong ControlNet type makes method unusable for 3D-consistent editing
- 5903 **Paper2Code (Score: 0.52 overall, 0% trainable)**
- 5904 • **Strengths:** Complete reference sheet generation, correct depth ControlNet, proper view updating
- 5905 • **Fatal Issues:**
- Not a Nerfstudio plugin - multi-file standalone code
 - Missing fine-tuning stage in training
 - Wrong hyperparameter values from config
- 5906 • **Result:** Core algorithm correct but cannot integrate with Nerfstudio framework
- 5907 **AutoP2C (Score: 0.08 overall, 0% trainable)**
- 5908 • **Strengths:** Basic dataset loading structure
- 5909 • **Fatal Issues:**
- No reference sheet generation or ControlNet implementation
 - No dataset updating or editing pipeline
 - Generic NeRF code without any SIGNeRF-specific features
- 5910 • **Result:** Complete failure to understand paper, produces generic NeRF only

5917 3.9.8. Hyperparameter Fidelity

5918 Table 69. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5	R1	P2C	AutoP2C
ControlNet scale	0.4-1.0	✓	✓	✓	✗	✗
Guidance scale	6-10	✓	✓	✓	✓	✗
Denoising strength	0.5-0.95	✓	✓	✓	✓	✗
Inference steps	40-50	✓	✓	✗	✓	✗
Initial iterations	30000	✓	✓	✗	✗	✗
Fine-tune iterations	10000	✓	✓	✗	✗	✗
LPIPS weight	0.01	✓	✗	✗	✗	✗
W Score	–	1.00	0.71	0.43	0.43	0.00

5929 3.9.9. Conclusion

5930 All baselines attempt to implement SIGNeRF with varying
5931 levels of sophistication. GPT-5 correctly implements the
5932 core algorithm including depth ControlNet and two-stage
5933 training but lacks Nerfstudio integration. DeepSeek R1 has
5934 good structure and iterative refinement but critically uses
5935 the wrong ControlNet type (Canny instead of Depth). Pa-
5936 per2Code implements reference sheet generation and view
5937 updating correctly but misses fine-tuning and uses wrong
5938 hyperparameters. AutoP2C completely fails to understand
5939 the paper, producing only generic NeRF code without any
5940 editing capabilities. Only NERFIFY produces immediately
5941 trainable code.

3.10. MCNeRF: Monte Carlo Rendering and De-noising for Real-Time NeRFs

5942 **3.10.1. Paper Overview**

5943 MCNeRF [8] introduces a general-purpose acceleration
5944 technique for volumetric NeRF rendering through Monte
5945 Carlo importance sampling. The method replaces dense ray
5946 marching with stochastic sampling based on transmittance-
5947 weighted probability distributions, achieving $7\times$ reduction
5948 in color MLP evaluations while maintaining visual qual-
5949 ity through a lightweight image-space denoiser. This work
5950 addresses the critical bottleneck of real-time NeRF render-
5951 ing by demonstrating that careful importance sampling with
5952 only 5 samples per pixel can match the quality of traditional
5953 128-sample quadrature.

5954 3.10.2. Implementation Overview

5955 Table 70. MCNeRF Implementation Summary Across All Baselines

Aspect	NERFIFY	GPT-5	DeepSeek R1	Paper2Code	AutoP2C
Lines of Code	476	287	892	1243	892
File Organization	Plugin	Single	Single	Multi-file	Multi-file
Monte Carlo Sampling	✓	✓	Partial	✗	✗
Two-Pass Rendering	✓	✓	Partial	✗	✗
Importance Sampling	✓	✓	✓	✗	✗
Denoiser Network	✓	✓	✓	Simplified	Attempted
Trainable	✓	✗	✗	✗	✗

5956 Note: NERFIFY always produces trainable code as a complete Nerfstudio
5957 plugin

5958 3.10.3. Novel Components

5959 Table 71. Novel Components in MCNeRF with Importance Weights

ID	Component	Weight w_i
C1	Monte Carlo integration replacing dense quadrature	0.25
C2	Importance sampling via transmittance CDF	0.20
C3	Stratified sampling for variance reduction	0.15
C4	Two-pass rendering (density then color)	0.15
C5	Lightweight 3-layer CNN denoiser	0.10
C6	Combined MSE + SSIM loss	0.05
C7	Per-scene denoiser training	0.05
C8	Real-time performance (5 spp target)	0.05

5958 3.10.4. Quantitative Metrics

5959 Table 72. MCNeRF Implementation Coverage Metrics

Method	C	I	M	W	Score _{LLM}
NERFIFY (Ours)	1.00	0.00	0.00	1.00	0.95
GPT-5	0.75	0.25	0.00	0.85	0.95
DeepSeek R1	0.50	0.38	0.13	0.80	0.74
Paper2Code	0.00	0.13	0.88	0.20	0.15
AutoP2C	0.00	0.25	0.75	0.10	0.08

5959 **3.10.5. Component-by-Component Analysis**
 5960 **Component C1: Monte Carlo Integration - All Baselines**
 5961 The paper replaces standard quadrature $\hat{C}(\mathbf{r}) =$
 5962 $\sum_{i=1}^N T_i \alpha_i \mathbf{c}_i$ with Monte Carlo estimation: $\hat{C}(\mathbf{r}) =$
 5963 $\mathbb{E}_{\mathbf{i} \sim p_i} [\mathbf{c}_i W]$ where $p_i = w_i/W$ and $W = \sum_i T_i \alpha_i$.

```
1 # \nerify\ : Unbiased MC estimator with proper
  normalization
2 rgb_sel = field_outputs_sel[FieldHeadNames.RGB]  # [R,
  , M, 3]
3 # Unbiased estimator: mean_j( c_{i,j} ) * W
4 rgb_est = rgb_sel.mean(dim=1) * W  # [R, 3]
5 # Composite with background using (1 - W)
6 bg = self._background_color(R, device)
7 rgb = rgb_est + (1.0 - W_clamped) * bg  # [R, 3]
8 # CORRECT: Proper MC formula W * mean(c_samples)
```

Listing 420. C1 Implementation: NERIFY (Score: 1.0)

```
1 # GPT-5: Correct MC integration with weight
  normalization
2 for ep in tqdm(range(epochs)):
3     rays_o, rays_d = rays
4     preds = monte_carlo_render(rays_o, rays_d, model)
5     # Proper MC estimator in monte_carlo_render
6     rgb_out = rgb_out / n_samples
7     return rgb_out
8 # CORRECT: Proper averaging over MC samples
```

Listing 421. C1 Implementation: GPT-5 (Score: 1.0)

```
1 # DeepSeek R1: MC sampling but simplified averaging
2 noisy_colors = torch.mean(colors, dim=1)  # (B, 3)
3 alpha = torch.mean(torch.sigmoid(densities), dim=1,
  keepdim=True)
4 # ISSUE: Missing proper weight normalization
```

Listing 422. C1 Implementation: DeepSeek R1 (Score: 0.8)

```
1 # Paper2Code: No MC implementation, only metrics
2 def calculate_psnr(predictions, ground_truth):
3     mse = F.mse_loss(predictions, ground_truth)
4     psnr = 20 * torch.log10(max_pixel_value) - 10 *
  torch.log10(mse)
5 # MISSING: No Monte Carlo sampling implementation
```

Listing 423. C1 Implementation: Paper2Code (Score: 0.0)

```
1 # AutoP2C: Dataset loading only, no rendering
2 def generate_rays(self):
3     for pose in self.camera_poses:
4         directions = np.stack([(i - cx) / fx, (j - cy)
  ) / fy])
5 # MISSING: No MC integration, only ray generation
```

Listing 424. C1 Implementation: AutoP2C (Score: 0.0)

6010 **Analysis:** NERIFY and GPT-5 correctly implement the
 6011 unbiased MC estimator with $W \cdot \text{mean}(c)$. DeepSeek R1
 6012 attempts MC but lacks proper weight normalization. Pa-
 6013 per2Code and AutoP2C completely miss the MC algorithm.

6014 **Component C2: Importance Sampling via Trans-
 6015 mittance CDF - All Baselines**

6016 The paper samples according to $p_i = w_i/W$ using in-
 6017 verse CDF transform on cumulative weights.

```
1 # \nerify\ : Complete CDF-based importance sampling
2 W = weights.sum(dim=1)  # [R, 1]
3 # CDF over samples
4 cdf = torch.cumsum(weights[..., 0], dim=1)  # [R, S]
5 # searchsorted returns indices in [0, S]
6 idx = torch.searchsorted(cdf, thresholds, right=True)
  .clamp(max=S - 1)
7 # CORRECT: Proper inverse CDF sampling
```

Listing 425. C2 Implementation: NERIFY (Score: 1.0)

```
1 # GPT-5: Proper cumulative weight building
2 W = weights.sum(dim=-1, keepdim=True)
3 W[W == 0] = 1e-6
4 pdf = weights / W
5 cdfs = torch.cumsum(pdf, dim=-1)
6 idx = torch.searchsorted(cdfs, us, right=True).clamp(
  max=n_coarse - 1)
7 # CORRECT: Textbook inverse CDF implementation
```

Listing 426. C2 Implementation: GPT-5 (Score: 1.0)

```
1 # DeepSeek R1: CDF sampling with minor issues
2 cdf = torch.cumsum(probs, dim=-1)
3 indices = torch.searchsorted(cdf, u, right=True)
4 sampled_depths = depth_below + t * (depth_above -
  depth_below)
5 # ISSUE: Overcomplicated interpolation logic
```

Listing 427. C2 Implementation: DeepSeek R1 (Score: 0.8)

```
1 # Paper2Code: No importance sampling
2 def render_view(self, camera_pose, image_size=[800,
  800]):
3     rays = self.generate_rays(camera_pose, image_size
  )
4     rendered_image = self.model.render(rays)
5 # MISSING: No CDF-based sampling
```

Listing 428. C2 Implementation: Paper2Code (Score: 0.0)

```
1 # AutoP2C: Basic ray generation only
2 ray_directions = directions @ pose_matrix[:, :, 3].T
3 ray_origins = np.broadcast_to(pose_matrix[:, :, 3],
  ray_directions.shape)
4 # MISSING: No importance sampling logic
```

Listing 429. C2 Implementation: AutoP2C (Score: 0.0)

Analysis: NERIFY and GPT-5 correctly implement in-
 verse CDF sampling. DeepSeek R1 has the right idea but
 overcomplicates interpolation. Paper2Code and AutoP2C
 lack any importance sampling.

**Component C3: Stratified Sampling for Variance Re-
 duction - All Baselines**

The paper uses stratified random sampling: $t_k = \frac{k+u_k}{M}$.
 W where $u_k \sim U(0, 1)$.

```
1 # \nerify\ : Perfect stratified sampling
2 if self.config.mc_stratified and M > 1:
3     u = torch.rand(R, M, device=device)
4     k = torch.arange(M, device=device)[None, :].
      expand(R, -1).float()
5     thresholds = ((k + u) / float(M)) * W  # [R, M]
6 # CORRECT: Proper jittered uniform sampling
```

Listing 430. C3 Implementation: NERIFY (Score: 1.0)

6079
6080
6081
6082
6083
6084
6085
6086

```
1 # GPT-5: Correct stratification
2 us = torch.rand(N_rays, n_samples), device=device)
3 # Stratified MC sampling
4 cdfs = torch.cumsum(pdf, dim=-1)
5 idx = torch.searchsorted(cdfs, us, right=True).clamp(
    max=n_coarse - 1)
6 # CORRECT: Proper stratified sampling implementation
```

Listing 431. C3 Implementation: GPT-5 (Score: 1.0)

6088
6089
6090
6091
6092
6093

```
1 # DeepSeek R1: Partial stratification
2 u = torch.rand(B, num_samples, device=depths.device)
3 u = (u + torch.arange(num_samples).float().view(1,
    -1)) / num_samples
4 # ISSUE: Missing device specification for arange
```

Listing 432. C3 Implementation: DeepSeek R1 (Score: 0.6)

6095
6096
6097
6098
6099

```
1 # Paper2Code: No stratified sampling
2 # Placeholder for ray generation logic
3 return []
4 # MISSING: No stratification implementation
```

Listing 433. C3 Implementation: Paper2Code (Score: 0.0)

6101
6102
6103
6104
6105

```
1 # AutoP2C: Uniform sampling only
2 i, j = np.meshgrid(np.arange(w), np.arange(h),
    indexing='xy')
3 # MISSING: No stratified sampling
```

Listing 434. C3 Implementation: AutoP2C (Score: 0.0)

6107
6108
6109
6110
6111

Analysis: NERFIFY and GPT-5 correctly implement stratified sampling with jittered uniforms. DeepSeek R1 partially implements but has device handling issues. Paper2Code and AutoP2C lack stratification.

Component C4: Two-Pass Rendering - All Baselines

The paper specifies Pass 1 for density-only queries to build weights, Pass 2 for selective color evaluation.

6115
6116
6117
6118
6119
6120
6121
6122
6123
6124
6125

```
1 # \nerify : Clean two-pass separation
2 # First pass: densities, weights, accumulations
3 density, density_embedding = self.field.get_density(
    ray_samples)
4 weights = ray_samples.get_weights(density)
5 # Second pass: MC sample color indices
6 idx = torch.searchsorted(cdf, thresholds, right=True)
7 field_outputs_sel = self.field.get_outputs(
    selected_ray_samples, de_sel)
8 # CORRECT: Clear separation of passes
```

Listing 435. C4 Implementation: NERFIFY (Score: 1.0)

6126
6127
6128
6129
6130
6131
6132
6133
6134
6135

```
1 # GPT-5: Explicit two-pass with comments
2 # First pass: compute opacity weights
3 for i in range(n_coarse):
4     _, sigma = model(pts, rays_d)
5     weights[:, i] = T * alpha
6 # Second pass: evaluate radiance at sampled indices
7 for j in range(n_samples):
8     rgb, _ = model(pts, rays_d)
9 # CORRECT: Well-documented two-pass structure
```

Listing 436. C4 Implementation: GPT-5 (Score: 1.0)

6137
6138
6139
6140
6141

```
1 # DeepSeek R1: Two-pass but inefficient
2 with torch.no_grad():
3     coarse_densities = scene_representation.
        compute_density(points_flat)
```

```
4 # Later: compute radiance
5 colors_flat = scene_representation.compute_radiance(
    points_flat, rays_d_flat)
6 # ISSUE: Not clearly separated passes
```

Listing 437. C4 Implementation: DeepSeek R1 (Score: 0.8)

```
1 # Paper2Code: Single-pass rendering assumption
2 rendered_image = self.model.render(rays)
3 # MISSING: No two-pass implementation
```

Listing 438. C4 Implementation: Paper2Code (Score: 0.0)

```
1 # AutoP2C: No rendering implementation
2 self.rays.append((ray_origins, ray_directions))
3 # MISSING: No pass structure at all
```

Listing 439. C4 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY and GPT-5 correctly separate density and color evaluation. DeepSeek R1 has two passes but lacks clear separation. Paper2Code and AutoP2C miss the two-pass structure entirely.

Component C5: Lightweight 3-Layer CNN Denoiser - All Baselines

The paper specifies a 3-layer CNN with 8 output channels, ReLU activations, and 4 input channels (RGB+ α).

```
1 # \nerify\ : Complete configurable denoiser
2 class SimpleDenoiser(nn.Module):
3     def __init__(self, in_ch: int = 4, feat: int = 8):
4         :
5             self.net = nn.Sequential(
6                 nn.Conv2d(in_ch, feat, 3, padding=1),
7                 nn.ReLU(inplace=True),
8                 nn.Conv2d(feat, feat, 3, padding=1),
9                 nn.ReLU(inplace=True),
10                nn.Conv2d(feat, 3, 3, padding=1),
11            )
12 # CORRECT: Exact paper specification
```

Listing 440. C5 Implementation: NERFIFY (Score: 1.0)

```
1 # GPT-5: Proper 3-layer denoiser
2 class DenoiseNet(nn.Module):
3     def __init__(self):
4         self.fc1 = nn.Linear(3, hidden_dim)
5         self.fc2 = nn.Linear(hidden_dim, hidden_dim)
6         self.sigma_out = nn.Linear(hidden_dim, 1)
7 # CORRECT: 3-layer architecture with proper channels
```

Listing 441. C5 Implementation: GPT-5 (Score: 1.0)

```
1 # DeepSeek R1: Denoiser with feature computation
2 self.net = nn.Sequential(
3     nn.Conv2d(4, 8, kernel_size=3, padding=1),
4     nn.ReLU(), nn.Conv2d(8, 8, kernel_size=3, padding=
    1),
5     nn.ReLU(), nn.Conv2d(8, self.feature_dim + 2,
    kernel_size=3, padding=1))
6 # ISSUE: Extra complexity with feature branches
```

Listing 442. C5 Implementation: DeepSeek R1 (Score: 0.8)

```
1 # Paper2Code: Simplified denoiser attempt
2 class DenoiseNet(nn.Module):
3     def forward(self, noisy_image, alpha_channel):
4         x = self.activation(self.conv1(input_tensor))
5         x = self.activation(self.conv2(x))
```

6205

```
1 # ERROR: Missing proper initialization
```

Listing 443. C5 Implementation: Paper2Code (Score: 0.4)

6206
6207
6208
6209

```
1 # AutoP2C: Placeholder denoiser
2 denoised_output = F.conv2d(noisy_image, x, padding=1)
3 # CRITICAL ERROR: No actual network defined
```

Listing 444. C5 Implementation: AutoP2C (Score: 0.2)

Analysis: NERFIFY and GPT-5 implement the exact 3-layer specification. DeepSeek R1 adds unnecessary complexity. Paper2Code has structure but missing initialization. AutoP2C only has placeholder code.

Component C6: Combined MSE + SSIM Loss - All Baselines

The paper uses $\mathcal{L} = \mathcal{L}_{\text{MSE}} + \lambda \cdot \mathcal{L}_{\text{SSIM}}$ with $\lambda = 0.1$.

```
1 # \nerfify\ : Proper combined loss
2 total_loss = mse_loss + 0.1 * ssim_loss # lambda =
    0.1 as in paper
3 # CORRECT: Exact paper specification
```

Listing 445. C6 Implementation: NERFIFY (Score: 1.0)

6224
6225
6226
6227
6228
6229

```
1 # GPT-5: Correct SSIM implementation
2 loss_recon = F.mse_loss(denoised, gt_colors)
3 loss_ssim = ssim_loss(denoised, gt_colors)
4 loss = loss_recon + 0.1 * loss_ssim
5 # CORRECT: Proper weighting with lambda=0.1
```

Listing 446. C6 Implementation: GPT-5 (Score: 1.0)

6231
6232
6233
6234
6235
6236

```
1 # DeepSeek R1: SSIM but simplified
2 ssim_loss = 1 - self._ssim(denoised_images,
    target_images)
3 total_loss = mse_loss + 0.1 * ssim_loss
4 # ISSUE: Simplified SSIM computation
```

Listing 447. C6 Implementation: DeepSeek R1 (Score: 0.6)

6238
6239
6240
6241
6242

```
1 # Paper2Code: Has SSIM but no MC
2 ssim_val = compare_ssim(ground_truth, prediction,
    multichannel=True)
3 # CORRECT: SSIM computation but missing integration
```

Listing 448. C6 Implementation: Paper2Code (Score: 1.0)

6244
6245
6246

```
1 # AutoP2C: No loss implementation
2 # MISSING: No loss functions defined
```

Listing 449. C6 Implementation: AutoP2C (Score: 0.0)

Analysis: NERFIFY and GPT-5 implement the exact loss combination. DeepSeek R1 has correct weighting but simplified SSIM. Paper2Code has SSIM but not integrated. AutoP2C missing entirely.

Component C7: Per-Scene Denoiser Training - All Baselines

The paper trains a separate denoiser for each scene with 100k steps.

6256
6257
6258
6259
6260
6261
6262
6263
6264

```
1 # \nerfify\ : Full per-scene pipeline
2 # Per scene, render training views with MC sampling
3 # Optimize DenoiseNet with Adam, 100k steps, batch
    size 32
4 self.denoiser_optimizer = torch.optim.Adam(
5     model.denoiser.parameters(), lr=config.get('
        denoiser_lr', 1e-3))
6 # CORRECT: Complete per-scene training
```

Listing 450. C7 Implementation: NERFIFY (Score: 1.0)

6266
6267
6268
6269
6270
6271
6272

```
1 # GPT-5: Per-scene training present
2 def train_mcnerf(model, denoiser, optimizer, rays,
    gt_colors, epochs=200):
3     for ep in tqdm(range(epochs)):
4         denoised = denoiser(inputs).view(-1, 3)
5 # ISSUE: Simplified training loop
```

Listing 451. C7 Implementation: GPT-5 (Score: 0.6)

6274
6275
6276
6277
6278
6279

```
1 # DeepSeek R1: Has per-scene training
2 self.denoiser_optimizer = torch.optim.Adam(
3     model.denoiser.parameters(), lr=config.get('
        denoiser_lr', 1e-3))
4 # CORRECT: Per-scene optimization setup
```

Listing 452. C7 Implementation: DeepSeek R1 (Score: 0.8)

6281
6282
6283
6284
6285
6286
6287

```
1 # Paper2Code: Training loop exists
2 self.optimizer = optim.Adam(self.model.parameters(),
    lr=learning_rate)
3 for global_step in range(self.total_steps):
4     self.optimizer.step()
5 # ISSUE: Not scene-specific
```

Listing 453. C7 Implementation: Paper2Code (Score: 0.6)

6289
6290
6291
6292
6293
6294
6295

```
1 # AutoP2C: Generic training attempt
2 trainer = Trainer(model=mcnerf_model, optimizer=
    optimizer)
3 trainer.train(num_epochs=config['training']['
    num_epochs'])
4 # ERROR: No actual implementation
```

Listing 454. C7 Implementation: AutoP2C (Score: 0.4)

6297
6298
6299
6300
6301
6302

Analysis: NERFIFY has complete per-scene training. GPT-5 and DeepSeek R1 have training but simplified. Paper2Code has training but not scene-specific. AutoP2C only has placeholders.

Component C8: Real-time Performance Target - All Baselines

The paper targets 5 samples per pixel for real-time rendering at 20 fps.

6303
6304
6305
6306
6307
6308
6309
6310
6311
6312

```
1 # \nerfify\ : Optimized for real-time
2 mc_spp: int = 5 # Samples per pixel
3 # WebGL shader implementing two-pass MC importance
    sampling
4 # Achieves interactive frame rates on 2021 Apple M1
    laptop
5 # CORRECT: Full real-time optimization
```

Listing 455. C8 Implementation: NERFIFY (Score: 1.0)

6314
6315
6316
6317
6318

```
1 # GPT-5: Targets 5 spp but not optimized
2 M = 5 # use M=5 for real-time
3 # Basic implementation without GPU optimizations
4 # ISSUE: No real-time optimizations
```

Listing 456. C8 Implementation: GPT-5 (Score: 0.4)

6320
6321
6322
6323
6324

```
1 # DeepSeek R1: Has 5 spp setting
2 self.num_importance_samples = config.get('
    num_importance_samples', 5)
3 # ISSUE: No performance optimizations
```

Listing 457. C8 Implementation: DeepSeek R1 (Score: 0.4)

6326
6327
6328

```
1 # Paper2Code: No performance considerations
2 # MISSING: No real-time optimizations
```

Listing 458. C8 Implementation: Paper2Code (Score: 0.0)

6330
6331
6332

```
1 # AutoP2C: No performance implementation
2 # MISSING: No optimization for speed
```

Listing 459. C8 Implementation: AutoP2C (Score: 0.0)

Analysis: Only NERFIFY fully optimizes for real-time performance. GPT-5 and DeepSeek R1 set 5 spp but lack optimizations. Paper2Code and AutoP2C ignore performance entirely.

3.10.6. Scoring Analysis

Table 73. MCNeRF Implementation Scores

Method	C1	C2	C3	C4	C5	C6	C7	C8	Weighted Avg _{LLM}
NERFIFY	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.00
GPT-5	1.0	1.0	1.0	1.0	1.0	1.0	0.6	0.4	0.95
DeepSeek R1	0.8	0.8	0.6	0.8	0.8	0.6	0.8	0.4	0.74
Paper2Code	0.0	0.0	0.0	0.0	0.4	1.0	0.6	0.0	0.15
AutoP2C	0.0	0.0	0.0	0.0	0.2	0.0	0.4	0.0	0.08

3.10.7. Why Baselines Fail Despite Component Scores

6339

6340

6341

6342

6343

6344

6345

6346

6347

6348

6349

6350

6351

6352

6353

6354

6355

6356

6357

6358

6359

6360

6361

6362

6363

6364

6365

6366

6367

6368

6369

6370

6371

6372

6373

6374

6375

6376

6377

6378

6379

6380

6381

6382

6383

6384

6385

6386

6387

6388

6389

6390

6391

6392

6393

6394

6395

- Completely missing Monte Carlo sampling (core algorithm)
- No importance sampling or two-pass rendering
- Placeholder functions throughout

- **Result:** Fundamental misunderstanding of paper's core contribution

AutoP2C (Score: 0.08 components, 0% trainable)

- **Strengths:** Basic dataset loading structure

- **Fatal Issues:**

- No Monte Carlo implementation whatsoever
- Missing all core algorithms
- Placeholder denoiser with no actual network

- **Result:** Only implements data loading, misses entire algorithm

3.10.8. Hyperparameter Fidelity

Table 74. Hyperparameter Implementation Accuracy

Parameter	Paper	NERFIFY	GPT-5	R1	P2C	AutoP2C
Samples per pixel	5	✓	✓	✓	✗	✗
Coarse samples	128	✓	✓	✓	✗	✗
Denoiser layers	3	✓	✓	✓	~	✗
Denoiser channels	8	✓	✓	✓	~	✗
SSIM weight λ	0.1	✓	✓	✓	✓	✗
Stratified sampling	Yes	✓	✓	~	✗	✗
Learning rate	10^{-3}	✓	✓	✓	✓	✓
Batch size	32	✓	~	~	✓	✓
Training steps	100k	✓	~	~	✓	~
W Score	–	1.00	0.85	0.80	0.20	0.10

3.10.9. Conclusion

All baselines attempt to implement MCNeRF with varying degrees of sophistication. NERFIFY achieves perfect implementation (1.00 correct) as a complete Nerfstudio plugin with all components properly integrated. GPT-5 demonstrates exceptional algorithmic understanding (0.95 score) with correct Monte Carlo mathematics and stratified sampling, but lacks framework integration. DeepSeek R1 provides good coverage (0.74) with multiple backend support but has critical execution issues. Paper2Code (0.15) and AutoP2C (0.08) fundamentally fail by completely missing the core Monte Carlo sampling algorithm, implementing only peripheral components like metrics and data loading. The stark contrast between NERFIFY's 100% trainable implementation and the 0% trainable rate across all baselines underscores the critical importance of domain-specific synthesis with proper framework integration for reproducible research code generation.

6396

4. Implementation Details

6397

4.1. Context-Free Grammar

6398

6399

6400

6401

6402

6403

6404

6405

6406

6407

6408

6409

6410

6411

6412

6413

6414

6415

6416

6417

6418

6419

6420

6421

6422

6423

6424

6425

6426

6427

6428

6429

6430

6431

6432

6433

6434

6435

6436

6437

6438

6439

6440

6441

6442

6443

6444

6445

6446

6447

6448

6449

6450

6451

6452

6453

6454

6455

6456

6457

6458

6459

6460

6461

6462

6463

6464

6465

6466

6467

6468

6469

6470

6471

6472

6473

```

1
2 Context-Free Grammar
3 // Notation:
4 // - Terminal symbols are in double quotes.
5 // - Non-terminals are in angle brackets.
6 // - \epsilon denotes the empty string.
7 // - [] = optional, {} = zero or more repetitions.
8 // - <PY_EXPR>, <PY_STMT>, <MARKDOWN>, <TOML> are lexical non-terminals
9 // handled by a separate Python / Markdown / TOML grammar or left unconstrained.
10
11 // =====
12 // Top-level \nerify\ repository & multi-file protocol for the LLM
13 // =====
14
15 <NeRFifyRepo> ::= <ConfigFile> <PipelineFile> <ModelFile> <DataManagerFile>
16     [<FieldFile>]
17     {<AuxFile>}
18     <InitFile>
19     <PyProjectFile>
20     [<ReadmeFile>]
21
22 // Each file is wrapped in explicit markers that the LLM must output.
23 // In the paper we show them as terminals; in code they are literal strings.
24
25 <ConfigFile> ::= "@@FILE: method_template/template_config.py\n"
26     <ConfigModule>
27     "\n@@END_FILE\n"
28
29 <PipelineFile> ::= "@@FILE: method_template/template_pipeline.py\n"
30     <PipelineModule>
31     "\n@@END_FILE\n"
32
33 <ModelFile> ::= "@@FILE: method_template/template_model.py\n"
34     <ModelModule>
35     "\n@@END_FILE\n"
36
37 <DataManagerFile> ::= "@@FILE: method_template/template_datamanager.py\n"
38     <DataManagerModule>
39     "\n@@END_FILE\n"
40
41 <FieldFile> ::= "@@FILE: method_template/template_field.py\n"
42     <FieldModule>
43     "\n@@END_FILE\n"
44
45 // Optional auxiliary modules that the LLM is allowed to generate.
46 // They share the generic <AuxModuleBody> grammar (Python module).
47
48 <AuxFile> ::= "@@FILE: method_template/" <AuxFileName> ".py\n"
49     <AuxModuleBody>
50     "\n@@END_FILE\n"
51
52 <AuxFileName> ::= "utils" | "losses" | "encoders" | "render_utils"
53     | "metrics" | "diffusion" | "samplers" | "scripts"
54
55 // Package initialiser and entry point in pyproject.toml
56
57 <InitFile> ::= "@@FILE: method_template/__init__.py\n"
58     [<ModuleDocstring>]
59     "from .template_config import " <MethodVar> "\n"
60     "__all__ = [\"" <MethodVar> "\"]\n"
61     "@@END_FILE\n"
62
63 <PyProjectFile> ::= "@@FILE: pyproject.toml\n"
64     "[project.entry_points.'nerfstudio.method_configs']\n"
65     <MethodName> " = \"method_template.template_config: " <MethodVar> "\"\n"
66     {<TomlLine>}
67     "@@END_FILE\n"
68
69 <ReadmeFile> ::= "@@FILE: README.md\n"
70     <MARKDOWN>
71     "@@END_FILE\n"
72
73 <TomlLine> ::= <TOML> // any additional TOML lines
74

```

```

76 // Method identity used consistently across files
77 <MethodName> ::= <LowerIdent>           // snake_case method name
78 <MethodVar>  ::= <LowerIdent>           // Python variable (e.g. method_template)
79
80 // =====
81 // Configuration module: MethodSpecification + TrainerConfig wiring
82 // =====
83
84
85 <ConfigModule> ::= [<ModuleDocstring>
86   <ConfigImports>
87   <MethodSpecDef>
88
89 <ModuleDocstring> ::= '"""' {<Char>} '"""' "\n"
90
91 <ConfigImports> ::= <RequiredConfigImports> {<ExtraImport>}
92
93 <RequiredConfigImports> ::= 
94   "from __future__ import annotations\n"
95   "from method_template.template_datamanager import TemplateDataManagerConfig\n"
96   "from method_template.template_model import TemplateModelConfig\n"
97   "from method_template.template_pipeline import TemplatePipelineConfig\n"
98   "from nerfstudio.configs.base_config import ViewerConfig\n"
99   "from nerfstudio.data.dataparsers.nerfstudio_dataparser import NerfstudioDataParserConfig\n"
100  "from nerfstudio.engine.optimizers import AdamOptimizerConfig\n"
101  "from nerfstudio.engine.schedulers import ExponentialDecaySchedulerConfig\n"
102  "from nerfstudio.engine.trainer import TrainerConfig\n"
103  "from nerfstudio.plugins.types import MethodSpecification\n"
104
105 <ExtraImport> ::= 
106   "from" <ModulePath> "import" <IdentList> "\n"
107 | "import" <ModulePath> "\n"
108
109 <ModulePath> ::= <LowerIdent> {"." <LowerIdent>}
110 <IdentList>  ::= <Ident> {"." <Ident>}
111 <Ident>      ::= <LowerIdent> | <UpperIdent>
112
113 <MethodSpecDef> ::= <MethodVar> " = MethodSpecification("
114   <TrainerConfig> ","
115   <DescriptionArg>
116 ") \n"
117
118 <TrainerConfig> ::= "config=TrainerConfig(" <TrainerKwargs> ")"
119
120 <TrainerKwargs> ::= <TrainerKwarg> {"." <TrainerKwarg>}
121
122 <TrainerKwarg> ::= "method_name=" <String>
123   | "pipeline=" <PipelineConfig>
124   | "datamanager=" <DataManagerConfig>
125   | "model=" <ModelConfig>
126   | "viewer=" <ViewerConfigExpr>
127   | "optimizers=" <OptimizerDict>
128   | "vis=" <String>
129   | <Ident> "=" <PY_EXPR>    // catch-all extra kwargs
130
131 <PipelineConfig> ::= "TemplatePipelineConfig(" <PyArgs> ")"
132 <DataManagerConfig> ::= "TemplateDataManagerConfig(" <PyArgs> ")"
133 <ModelConfig>  ::= "TemplateModelConfig(" <PyArgs> ")"
134
135 <your entire grammar goes here>
136
137 <ViewerConfigExpr> ::= "ViewerConfig(" <PyArgs> ")"
138
139 <OptimizerDict> ::= "{" <OptimizerEntry> {"." <OptimizerEntry>} "}"
140
141 <OptimizerEntry> ::= <String> ":" {
142   "optimizer": <OptimizerConfig> ","
143   "scheduler": <SchedulerConfig>
144 }
145
146 <OptimizerConfig> ::= <Ident> "(" <PyArgs> ")"
147 <SchedulerConfig> ::= <Ident> "(" <PyArgs> ")"
148
149 <PyArgs> ::= [<PY_EXPR> {"." <PY_EXPR>}]
150
151 <DescriptionArg> ::= "description=" <String>
152
153 <String> ::= '"' {<Char>} "'"
154

```

```

6553 // =====
6554 // Pipeline module: VanillaPipelineConfig + VanillaPipeline subclass
6555 // =====
6556
6557 <PipelineModule> ::= [<ModuleDocstring>
6558     <PipelineImports>
6559     <PipelineConfigClass>
6560     <PipelineClass>
6561
6562 <PipelineImports> ::= {<ImportStmt>}
6563
6564 <ImportStmt> ::= 
6565     "import" <ModulePath> "\n"
6566     | "from" <ModulePath> "import" <IdentList> "\n"
6567
6568 <PipelineConfigClass> ::=
6569     "@dataclass\n"
6570     "class " <PipelineConfigName> "(VanillaPipelineConfig):" <IndentedSuite>
6571
6572 <PipelineClass> ::=
6573     "class " <PipelineName> "(VanillaPipeline):" <PipelineSuite>
6574
6575 <PipelineSuite> ::= <InitMethod> {<PipelineMethod>}
6576
6577 <InitMethod> ::=
6578     "def __init__(self, " <PipelineParams> ")" <IndentedSuite>
6579     // Inside the suite we require core wiring patterns:
6580     // - super().__init__(...)
6581     // - datamanager setup
6582     // - model setup
6583     // - optional DDP setup
6584
6585 <PipelineParams> ::= <Param> {" ", " <Param> }
6586 <Param> ::= <Ident> [": " <Type>] [" = " <PY_EXPR>]
6587
6588 <Type> ::= "int" | "float" | "bool" | "Literal" | "Optional" | "Type" | "DataManager" | "Model" | <Ident>
6589
6590 <PipelineMethod> ::= <DefHeader> <IndentedSuite>
6591
6592 <DefHeader> ::= "def " <Ident> "(" [<Param> {" ", " <Param>}] ")"
6593
6594 <IndentedSuite> ::= "\n" <Indent> {<PY_STMT>} <Dedent>
6595
6596 // We do not expand <PY_STMT> further; it's a lexical non-terminal for valid Python code.
6597
6598 // =====
6599 // Model module: NerfactoModelConfig + NerfactoModel subclass
6600 // =====
6601
6602 <ModelModule> ::= [<ModuleDocstring>
6603     <ModelImports>
6604     <ModelConfigClass>
6605     <ModelClass>
6606
6607 <ModelImports> ::= {<ImportStmt>}
6608
6609 <ModelConfigClass> ::=
6610     "@dataclass\n"
6611     "class " <ModelConfigName> "(" <BaseModelConfig> ")" <IndentedSuite>
6612
6613 <BaseModelConfig> ::= "NerfactoModelConfig" | "ModelConfig" | <Ident>
6614
6615 <ModelClass> ::=
6616     "class " <modelName> "(" <BaseModel> ")" <ModelSuite>
6617
6618 <BaseModel> ::= "NerfactoModel" | "Model" | <Ident>
6619
6620 <ModelSuite> ::= <PopulateModules>
6621     [<GetLossDict>]
6622     {<ModelMethod>}
6623
6624 <PopulateModules> ::=
6625     "def populate_modules(self):" <IndentedSuite>
6626
6627 <GetLossDict> ::=
6628     "def get_loss_dict(self, outputs, batch, metrics_dict=None):" <IndentedSuite>
6629
6630 <ModelMethod> ::= <DefHeader> <IndentedSuite>
6631

```

```

234 // =====
235 // DataManager module: VanillaDataManagerConfig + VanillaDataManager subclass
236 // =====
237
238 <DataManagerModule> ::= [<ModuleDocstring>
239     <DataManagerImports>
240     <DataManagerConfigClass>
241     <DataManagerClass>
242
243 <DataManagerImports> ::= {<ImportStmt>}
244
245 <DataManagerConfigClass> ::= 6632
246     "@dataclass\n"
247     "class " <DataManagerConfigName> "(VanillaDataManagerConfig):" <IndentedSuite> 6633
248
249 <DataManagerClass> ::= 6634
250     "class " <DataManagerName> "(VanillaDataManager):" <DataManagerSuite> 6635
251
252 <DataManagerSuite> ::= <DataManagerInit> 6636
253     [<NextTrainMethod>] 6637
254     [<NextEvalMethod>] 6638
255     {<DataManagerMethod>} 6639
256
257 <DataManagerInit> ::= 6640
258     "def __init__(self, " <Param> {" , " <Param>} "):" <IndentedSuite> 6641
259
260 <NextTrainMethod> ::= 6642
261     "def next_train(self, step: int) -> Tuple[RayBundle, Dict]:"
262         <IndentedSuite> 6643
263
264 <NextEvalMethod> ::= 6644
265     "def next_eval(self, step: int) -> Tuple[RayBundle, Dict]:"
266         <IndentedSuite> 6645
267
268 <DataManagerMethod> ::= <DefHeader> <IndentedSuite> 6646
269
270 // =====
271 // Field module (optional): NerfactoField subclass
272 // =====
273
274 <FieldModule> ::= [<ModuleDocstring>
275     <FieldImports>
276     <FieldClass>
277
278 <FieldImports> ::= {<ImportStmt>} 6647
279
280 <FieldClass> ::= 6648
281     "class " <FieldName> "(" <BaseField> ")" <FieldSuite> 6649
282
283 <BaseField> ::= "NerfactoField" | "Field" | <Ident> 6650
284
285 <FieldSuite> ::= <FieldInit> 6651
286     [<ForwardMethod>] 6652
287     {<FieldMethod>} 6653
288
289 <FieldInit> ::= 6654
290     "def __init__(self, " <Param> {" , " <Param>} "):" <IndentedSuite> 6655
291
292 <ForwardMethod> ::= 6656
293     "def forward(self, ray_samples: RaySamples) -> Dict[str, Tensor]:"
294         <IndentedSuite> 6657
295
296 <FieldMethod> ::= <DefHeader> <IndentedSuite> 6658
297
298 // =====
299 // Generic auxiliary module body (for utils.py, diffusion.py, etc.)
300 // =====
301
302 <AuxModuleBody> ::= [<ModuleDocstring>
303     {<ImportStmt>}
304     {<TopLevelDef>}
305
306 <TopLevelDef> ::= <DefHeader> <IndentedSuite> 6659
307     | "@dataclass\n" "class " <UpperIdent> "(" [<Ident>] ")" <IndentedSuite> 6660
308     | "class " <UpperIdent> "(" [<Ident>] ")" <IndentedSuite> 6661
309
310 // =====
311 <LowerIdent> ::= <LowerLetter> {<Letter> | "_" | <Digit>} 6662
312 <UpperIdent> ::= <UpperLetter> {<Letter> | "_" | <Digit>} 6663

```

```
6711  
6712 313 <LowerLetter> ::= "a" | "b" | ... | "z"  
6713 314 <UpperLetter> ::= "A" | "B" | ... | "Z"  
6714 315 <Letter> ::= <LowerLetter> | <UpperLetter>  
6715 316 <Digit> ::= "0" | "1" | ... | "9"  
6716 317 <Char> ::= any printable character except '"' and newline
```

6718 4.1.1. Code Generation Templates

6719 4.2. Multi-Agent Architecture and Agent Specifications

6720 The system is built on LangChain for agent orchestration and LangGraph for workflow management. We use DeepAgent as
6721 the base agent framework and GPT-5 API for code generation. Paper parsing employs MinerU [19] to extract text, equations,
6722 and figures from PDFs. External knowledge retrieval uses Tavily for web search when agents encounter unfamiliar techniques
6723 or require implementation details absent from the paper. For visual feedback, NeRFify uses Qwen3-VL 8B model.

6724 4.3. LLM Prompts Used in NeRFify

6725 This section lists the main prompts used by the agents in our NERFIFY framework. For brevity, we show the system message
6726 and the static part of the user template for each agent. Placeholders such as {template_tree} or {image_paths} are
6727 filled programmatically at runtime.

6728 4.3.1. Prompt 1: Markdown Cleaning Agent (agentic_clean.py)

6729 System prompt

```
6730  
6731 1 You are a meticulous Markdown cleaner for research papers.  
6732 2 Apply the following strictly. Output ONLY valid GitHub-flavored  
6733 3 Markdown no commentary, no code fences, no extra prose.  
6734 4  
6735 5 Rules:  
6736 6 1) Text hygiene: de-hyphenate wrapped words, fix OCR ligatures,  
6737 7 normalize quotes/dashes/units; remove duplicate lines/sections;  
6738 8 fix spacing/formatting.  
6739 9  
6740 10 2) Equations: NEVER delete or alter equations. Preserve inline  
6741 11 $...$ and display $$...$$ math exactly as authored, including  
6742 12 numbering/labels if present.  
6743 13  
6744 14 3) Scope pruning: remove generic narrative sections  
6745 15 (Introduction/background, Related Work surveys,  
6746 16 qualitative/marketing-style Results). Keep ONLY content needed  
6747 17 to implement and reproduce experiments: problem setup,  
6748 18 assumptions, notation, model architecture, objectives/losses,  
6749 19 algorithms/pseudocode, training schedule, datasets,  
6750 20 preprocessing, hyperparameters, ablations that affect  
6751 21 implementation, and evaluation protocol/metrics definitions.  
6752 22  
6753 23 4) Tables: delete benchmark/comparison tables. If a single clear  
6754 24 takeaway is obvious, replace with a one-line textual takeaway.  
6755 25 Retain implementation-critical tables (hyperparameters, layer  
6756 26 configs, dataset splits) but convert them into concise bullet  
6757 27 lists; no raw HTML or Markdown tables.  
6758 28  
6759 29 5) Figures: remove images/captions/links/placeholders. If nearby  
6760 30 text contains implementation-relevant details, keep that text  
6761 31 as plain prose. Do NOT invent content.  
6762 32  
6763 33 6) Strip raw HTML artifacts entirely (<td>, <tr>, <table>,  
6764 34 inline styles). Keep only pure Markdown and math.  
6765 35  
6766 36 7) Citations: remove unnecessary/dangling citations and citation
```

```

37     dumps. Keep citations only when required to identify
38     datasets/codebases/definitions essential for reproduction.
39
40 8) Summaries: when paragraphs are verbose or narrative, compress
41     to 3 6 bullets emphasizing actionable implementation details
42     and experimental setup (no fixed word threshold).
43
44 9) Final check: output must parse as clean, minimal Markdown with
45     intact math, no images, no HTML table tags, no benchmark
46     tables, and no broken anchors.
47
48 10) Brevity: make the output as short as possible while
49     preserving all information required for implementation and
50     experiments.

```

6767
6768
6769
6770
6771
6772
6773
6774
6775
6776
6777
6778
6779
6780

User prompt template

```

1 Clean the following Markdown. Output ONLY the cleaned Markdown.
2
3 <MARKDOWN CONTENT HERE>

```

6782
6783
6784
6785
6786

4.3.2. Prompt 2: Citation Search Agent (agentic_citation_recovery.py)

6788

System prompt (SEARCH_AGENT_SYSTEM_PROMPT)

6789

```

1 You are a NeRF and 3D vision research assistant focused on
2 citation discovery and dependency analysis.
3
4 For each user query:
5 - Use 'internet_search' to find the canonical page(s) for the
6   requested NeRF-related paper (arXiv, project page, or main
7   publication page).
8 - Skim titles, abstracts, and key metadata from the results.
9 - Produce a concise textual bundle that will later be parsed by
10 another LLM.
11
12 Your answer MUST:
13 - Clearly state the best-guess canonical title, arXiv ID (if any),
14   year, venue (if obvious), and 3 6 bullet points summarizing the
15   method.
16 - List 3 8 likely upstream dependency / baseline papers, each with
17   a one-line reason describing what is reused or extended.
18 - Include the strongest URLs (arXiv / project page / PDF) inline.
19
20 Keep the answer compact but information-dense. Do not write
21 generic advice; only report what you found in the search results.

```

6790
6791
6792
6793
6794
6795
6796
6797
6798
6799
6800
6801
6802
6803
6804
6805
6806
6807
6808
6809
6810
6811

User prompt template for a single query

6813

```

1 Search for the NeRF-related paper and its canonical page or arXiv:
2 {query}
3
4 Use the 'internet_search' tool as needed. Return a concise
5 textual bundle with:
6 - Canonical title and year
7 - arXiv ID (if any)
8 - Venue (if obvious)
9 - 3 6 bullet points summarizing the method

```

6814
6815
6816
6817
6818
6819
6820
6821
6822
6823

6824 10 - 3 8 likely upstream dependency papers **with** one-line reasons each
6825 11 - Best URLs (arXiv / project / PDF)

6827 4.3.3. Prompt 3: Dependency Extraction Prompt

6828 System prompt (DEPENDENCY_EXTRACTION_SYSTEM)

```
6829 1 You are a research assistant building a citation dependency graph
6830 2 for NeRF papers.
6831 3
6832 4 Given text that may include scraped search results, references,
6833 abstracts, or markdown, identify:
6834 5 (a) the paper's canonical title,
6835 6 (b) arXiv id if any,
6836 7 (c) year,
6837 8 (d) core components it introduces or uses as architectural modules,
6838 9 loss functions, training protocols, or datasets,
6839 10 (e) explicit dependencies on prior NeRF-style methods.
6840 11
6841 12
6842 13 Return a single JSON object with keys:
6843 14 - title: string
6844 15 - arxiv_id: string or null
6845 16 - year: integer or null
6846 17 - venue: string or null
6847 18 - url: string or null
6848 19 - summary: short string
6849 20 - components: {modules:[], losses:[], protocols:{}}
6850 21 - dependencies: [{title_or_id, reason, components_borrowed}]
6851 22 - citations: [{title_or_id, url?}]
6852 23
6853 24 Return STRICT JSON only, with no extra commentary.
```

6855 User prompt template

```
6856 1 Use the following bundle to extract paper facts and dependencies.
6857 2
6858 3 <BUNDLE AS JSON OR TEXT HERE>
6859 4
6860 5 Return strict JSON only.
```

6863 4.3.4. Prompt 4: Target Analysis Prompt

6864 System prompt (TARGET_ANALYSIS_SYSTEM)

```
6865 1 You extract novelties, formulas, and implementation details from
6866 2 NeRF papers and their dependencies.
6867 3
6868 4 Given a JSON citation/dependency graph plus metadata for a target
6869 paper, produce a compact analysis describing:
6870 5 - core novelties of the target method (architecture, losses,
6871 training/inference protocols),
6872 6 - key equations (with informal names, LaTeX-like forms, and short
6873 descriptions),
6874 7 - a high-level but implementation-ready plan of components.
6875 8
6876 9
6877 10 Return STRICT JSON with:
6878 11 - novelties: [{aspect, description}]
6879 12 - formulas: [{name, equation, description}]
```

```

15 - implementation_plan: {
16     modules:[...],
17     losses:[...],
18     protocols:[...],
19     data_requirements:[...],
20     pseudo_steps:[...]
21 }
22
23 No extra commentary outside this JSON.

```

6880
6881
6882
6883
6884
6885
6886
6887
6888
6889

User prompt template

6890
6891
6892
6893
6894
6895
6896
6897
6898

```

1 Analyze the following target paper text to extract novelties,
2 formulas, and an implementation plan.
3
4 <TARGET PAPER TEXT OR DERIVED MARKDOWN HERE>
5
6 Return strict JSON only.

```

6899

6900

6901

6902

6903

6904

6905

6906

6907

6908

6909

6910

6911

6912

6913

6914

6915

6916

6917

6918

6919

6920

6921

6922

6923

6924

6925

6926

6927

6928

6929

6930

6931

6932

6933

6934

6935

6936

4.3.5. Prompt 5: DAG + File Plan with Repo Snapshot (agentic_dag.py)**User prompt (DAG + file-plan generator)**

```

1 You are a senior Nerfstudio engineer.
2
3 TASK A      IMPORT DAG:
4 - Build a file-level DAG where each node is a relative file path
5   (e.g., "method_template/template_model.py").
6 - Add a directed edge for each internal Python import:
7   { "from": "<file>", "to": "<file>", "relation": "imports" }.
8 - The result must be a valid DAG (no cycles). If cycles are
9   detected, break them by removing the least number of edges.
10 - Use the REPO SNAPSHOT and OBSERVED IMPORTS as ground truth.
11 - If there is disagreement, prefer OBSERVED IMPORTS.
12
13 TASK B      FILE PLAN (optional for codegen):
14 - Propose the minimal set of files under "method_template/"
15   (subset of TEMPLATE TREE) that must be implemented for this
16   paper.
17 - For each file, provide:
18   path,
19   purpose,
20   depends_on (relative file paths),
21   key_classes,
22   key_functions.
23
24 Return STRICT JSON of the form:
25 {
26   "nodes": [
27     {"id": "<file>", "label": "<short description>"}
28   ],
29   "edges": [
30     {"from": "<src_file>", "to": "<dst_file>", "relation": "imports"}
31   ],
32   "files": [
33     {
34       "path": "method_template/template_model.py",
35       "purpose": "<short description>",

```

```
6937     "depends_on": [
6938         "method_template/template_field.py",
6939         ...
6940     ],
6941     "key_classes": ["METHOD_Model"],
6942     "key_functions": ["forward"]
6943 },
6944 ...
6945 ]
6946 }
6947
6948 Constraints:
6949 - Only include paths that exist in the TEMPLATE TREE or REPO
6950   SNAPSHOT.
6951 - Keep depends_on lists short and meaningful.
6952 - Do NOT invent new directories or top-level packages beyond what
6953   you see.
6954 - JSON must be syntactically valid and parseable.
6955
6956 REPO SNAPSHOT (selected files with excerpts):
6957 <snapshot_text>
6958
6959 OBSERVED INTERNAL IMPORTS (static analysis):
6960 <imports_text>
```

6962 System prompt used for this call

```
6963
6964 1 You are a NeRF planning agent. Return JSON only.
```

6966 4.3.6. Prompt 6: DAG + File Plan from Template Tree Only (build_dag.py)

6967 User prompt

```
6968
6969 1 You are a senior Nerfstudio engineer. Using ONLY the provided
6970   TEMPLATE TREE and TEMPLATE FILE CONTENTS, design a complete DAG
6971   of the method and a file-generation plan.
6972   - Do NOT invent files outside TEMPLATE TREE.
6973   - Do NOT include any in-context examples.
6974   - Be concise but complete.
6975
6976   8 Return strict JSON with:
6977     - "nodes": [
6978       {"id":"snake_case_id",
6979        "label":"Short title",
6980        "methods":["method1","method2"]}
6981     ]
6982     - "edges": [
6983       {"from":"node_id",
6984        "to":"node_id",
6985        "relation":"feeds|queries|supervises|produces|writes"}
6986     ]
6987     - "files": [
6988       {
6989         "path":"relative/path.py or README.md",
6990         "purpose":"short description",
6991         "depends_on":["other/file.py", ...],
6992         "key_classes":["ClassA"],
6993         "key_functions":["fn_a", "fn_b"]}
```

```
26     }
27 ]
28
29 TEMPLATE TREE (authoritative list of allowed files):
30 {template_tree}
31
32 TEMPLATE FILE CONTENTS (use to match APIs/imports):
33 {template_files}
```

6994
6995
6996
6997
6998
6999
7000
7001

System prompt

```
1 You are a NeRF planning agent. Return JSON only.
```

7003
7004
7005

4.3.7. Prompt 7: NeRFify Code Generation Prompt (`agentic_coding.py`)

7007

Base full-code prompt (`_build_full_prompt`)

```
1 You are a senior NeRFStudio engineer. Synthesize a complete
2 implementation for the TARGET PAPER by learning patterns from
3 in-context examples and by adhering EXACTLY to the provided
4 NeRFStudio method template.
5
6 ### WHAT YOU GET
7 A) TEMPLATE TREE (authoritative list of output files & their
8 relative paths under method_template):
9 {template_tree}
10
11 B) TEMPLATE FILE CONTENTS (current stubs you must overwrite; use
12 for APIs/imports naming):
13 {template_files}
14
15 C) IN-CONTEXT EXAMPLES (order matters; learn structure, naming,
16 losses, schedulers, data flows):
17 {examples}
18
19 D) TARGET PAPER      FINAL to IMPLEMENT
20 {final_paper}
21 {optional_file_plan_block}
22
23 ### IMPLEMENTATION REQUIREMENTS
24 1) OUTPUT FILES:
25 Generate code ONLY for the files that exist in TEMPLATE TREE
26 under method_template. Do NOT invent new files or paths
27 beyond this list.
28
29 2) METHOD NAME:
30 Infer a concise snake_case METHOD_NAME from the paper title
31 (e.g., "mip_nerf", "seathru_nerf"). Use this consistently
32 across files (config, model, field, datamanager, pipeline).
33
34 3) COMPATIBILITY & STYLE:
35 - Match Nerfstudio APIs and signatures shown in TEMPLATE
36 FILE CONTENTS.
37 - Use type hints and relative imports.
38 - Use defaults from the TARGET PAPER; encode losses exactly.
39 - Put data requirements/transforms in template_datamanager.py.
40 - Wire metrics in template_pipeline.py; compute them in model/
41 field as needed.
```

7008
7009
7010
7011
7012
7013
7014
7015
7016
7017
7018
7019
7020
7021
7022
7023
7024
7025
7026
7027
7028
7029
7030
7031
7032
7033
7034
7035
7036
7037
7038
7039
7040
7041
7042
7043
7044
7045
7046
7047
7048
7049
7050

```
7051    42      - Avoid unresolved imports or circular dependencies.  
7052    43  
7053    44 4) SELF-CONSISTENCY:  
7054    45      - The project should install with `pip install -e .`.  
7055    46      - Imports like  
7056    47          'from method_template.template_model import METHOD_NAME_Model'  
7057    48          must succeed.  
7058    49      - All files should form a coherent, runnable Nerfstudio method.  
7059    50  
7060    51 5) OUTPUT FORMAT:  
7061    52      - Return your answer as a sequence of file blocks:  
7062    53          @@FILE: relative/path.py  
7063    54          <file contents>  
7064    55          @@END_FILE  
7065    56      - One block per file. No extra commentary, no Markdown fences,  
7066    57          no explanations outside the file blocks.
```

7068 Incremental / consolidation variants

7069 The incremental and consolidation prompts (`_build_incremental_prompt` and `_build_consolidation_prompt`) reuse the same
7070 structure as above, but:

- 7071 • Restrict the allowed output paths to a subset `{allowed_files}` (plus optional `{extra_allowed_files}`).
- 7072 • Add an extra section E) CURRENT IMPLEMENTATION (YOU MAY MODIFY AS NEEDED) containing the current
7073 generated files:

```
7074    1 E) CURRENT IMPLEMENTATION (YOU MAY MODIFY AS NEEDED)  
7075    2 The following files were generated incrementally. You may change  
7076    3 ANY of them to ensure correctness and consistency:  
7077    4 {current_text}
```

7080 4.3.8. Prompt 8: ReAct Wrapper for Code Generation (`agentic_coding.py`)

7081 System prompt for ReAct agent

```
7082    1 You are a senior NeRFStudio engineer. You may use tools  
7083    2 (web_search, http_get) to verify facts. Your FINAL answer must  
7084    3 be ONLY the multi-file output in the exact format specified by  
7085    4 the user using @@FILE blocks and @@END_FILE, with NO extra  
7086    5 commentary, NO code fences, NO prefixes, and NO explanations.
```

7089 The human message for this agent is simply:

```
7090    1 {input}
```

7093 where `{input}` is the full-code prompt from Prompt 7 (or its incremental variant).

7094 4.3.9. Prompt 9: NeRF Artifact Detection Prompt (`nerf_qwen_vl.artifacts.py`)

7095 Image-level Qwen3-VL prompt (`build_prompt`)

```
7096    1 You are an expert in Neural Radiance Field (NeRF) rendering  
7097    2 diagnostics.  
7098    3 The image you see is a NeRF render named '{image_name}' with  
7099    4 resolution {width}x{height} pixels.  
7100    5  
7101    6 Your tasks:  
7102    7 1. Inspect the image for visible NeRF-specific rendering defects.  
7103    8 2. For each defect, provide a tight bounding box and classify it.  
7104    9 3. For each defect, propose concrete code-level changes to a  
7105    10 typical NeRF implementation that would likely reduce or
```

```
11     remove this artifact.  
12  
13 Consider the following defect types and use their keys exactly as  
14 written:  
15 {defect_list_text}  
16  
17 Examples of code-level suggestions you may use (adapt as  
18 appropriate):  
19 - Increase samples per ray (coarse and/or fine); increase  
20 proposal network steps.  
21 - Enable mip-NeRF-style cone tracing / anti-aliasing to reduce  
22 aliasing.  
23 - Increase distortion loss weight or density regularization to  
24 remove floaters.  
25 - Adjust near/far bounds or background handling to fix background  
26 leakage.  
27 - Add depth/normal smoothness loss to fix tearing or missing  
28 geometry.  
29 - Improve pose estimation (re-run COLMAP with higher quality) if  
30 multi-view misalignment is visible.  
31 - Increase training iterations or reduce learning rate if  
32 underfit / noisy.  
33 - Use higher input image resolution or improve dataloader  
34 resizing.  
35  
36 IMPORTANT:  
37 - Set width to {width} and height to {height} in the JSON.  
38 - Use as few defects as needed but up to 10 per image.  
39 - Be specific in code_suggestions: reference knobs like samples  
40 per ray, proposal networks, loss weights, occupancy grids,  
41 background color handling, pose optimization, etc.  
42  
43 Return ONLY a single valid JSON object, with this structure:  
44  
45 {  
46     "image_name": "<exact image file name>",  
47     "width": <integer width in pixels>,  
48     "height": <integer height in pixels>,  
49     "defects": [  
50         {  
51             "type": "<one of the defect type keys above>",  
52             "bbox": [x_min, y_min, x_max, y_max],  
53             "description": "<short natural-language description>",  
54             "codeSuggestions": [  
55                 "<one concrete change to the NeRF code or hyperparameters>",  
56                 "<another concrete suggestion if useful>"  
57             ]  
58         }  
59     ]  
60 }  
61  
62 Constraints:  
63 - bbox coordinates must be integers in pixel space, with:  
64     0 <= x_min < x_max <= width,  
65     0 <= y_min < y_max <= height.  
66 - If no defects are present, return "defects": [] but keep the  
67 rest of the JSON.  
68 - Do NOT output any extra commentary, markdown, or text outside
```

7107
7108
7109
7110
7111
7112
7113
7114
7115
7116
7117
7118
7119
7120
7121
7122
7123
7124
7125
7126
7127
7128
7129
7130
7131
7132
7133
7134
7135
7136
7137
7138
7139
7140
7141
7142
7143
7144
7145
7146
7147
7148
7149
7150
7151
7152
7153
7154
7155
7156
7157
7158
7159
7160
7161
7162
7163
7164

7166 69 the JSON.

7167 4.3.10. Prompt 10: ReAct Aggregation Prompt for NeRF Artifacts

7168 System prompt for LangChain ReAct agent

```
7169 1 You are orchestrating an automated inspection pipeline for NeRF
7170 2 renderings.
7171 3 You have access to a tool called 'inspect_nerf_render' which,
7172 4 given an image path string, returns a JSON string describing
7173 5 defects in that image.
7174 6
7175 7 Your job:
7176 8 1. For each image path provided in the user input, call
7177 9 'inspect_nerf_render'.
7178 10 2. Parse each JSON string returned by the tool.
7179 11 3. Aggregate them into a single JSON object with this structure:
7180 12
7181 13 {
7182 14     "images": [
7183 15         {
7184 16             "image_name": "<file name>",
7185 17             "width": <integer>,
7186 18             "height": <integer>,
7187 19             "defects": [
7188 20                 {
7189 21                     "type": "<defect type>",
7190 22                     "bbox": [x_min, y_min, x_max, y_max],
7191 23                     "description": "<description>",
7192 24                     "code_suggestions": ["<suggestion 1>", "<suggestion 2>", "..."]
7193 25                 }
7194 26             ]
7195 27         }
7196 28     ]
7197 29 }
7198 30
7199 31 Rules:
7200 32 - Call 'inspect_nerf_render' on EVERY image path and do not skip
7201 33 any.
7202 34 - Preserve the order of images as given.
7203 35 - If the tool reports an error or returns no defects, still
7204 36 include the image with "defects": [] (you may carry over an
7205 37 error field if present).
7206 38 - Your FINAL answer must be ONLY that JSON object, with no extra
7207 39 text.
```

7210 User prompt template

```
7211 1 You are given the following list of NeRF render image paths:
7212 2
7213 3 {JSON-ENCODED LIST OF IMAGE PATHS}
7214 4
7215 5 For EACH path in this list, you must:
7216 6 - Call the tool 'inspect_nerf_render' with the exact path string.
7217 7 - Afterwards, aggregate all tool outputs into the single JSON
7218 8 object described in the system message. Output ONLY that JSON
7219 9 object.
```

7222 =

7223

5. NeRFifyBench Dataset



Figure 3. Categorization of NeRF Papers by Integrability in Nerfstudio. Papers are grouped by implementation difficulty within the Nerfstudio framework. **Category 1:** Directly integrable methods modifying architecture, rendering, or losses without external dependencies. **Category 2:** Methods requiring pretrained models (CLIP, diffusion, SLAM) with substantial engineering effort. **Category 3:** Out-of-scope works where NeRF is used in the pipeline but serves different objectives (GANs, downstream applications, etc). The categorization illustrates alignment with Nerfstudio’s modular design and implementation feasibility using NERFIFY .

7224

5.1. Paper Selection Criteria

7225 NERFIFY-BENCH comprises 30 carefully curated NeRF research papers selected to provide comprehensive coverage of 7226 different implementation challenges. Our selection process 7227 prioritized diversity across architectural innovations, training 7228 strategies, and integration complexity to ensure robust 7229 evaluation of paper-to-code synthesis systems. 7230

7231 The benchmark is organized into four distinct set chosen 7232 from three categories as shown in figure 3, each designed 7233 to test specific capabilities of automated code generation 7234 systems.

7235 5.1.1. Category 1: Never-Implemented Papers (10 pa- 7236 pers)

7237 This category includes papers without any publicly available 7238 source code. To ensure evaluation integrity and avoid potential 7239 training data contamination, we specifically selected 7240 papers where no implementation exists in the public domain. 7241 This guarantees that large language models used in 7242 our agents could not have encountered corresponding code 7243 during pretraining, enabling unbiased assessment of purely 7244 paper-driven synthesis capabilities.

7245 For these papers, we commissioned expert reimplementations 7246 from graduate students with extensive NeRF research 7247 experience. Each implementation underwent rigorous veri- 7248 fication against paper-reported metrics on standard benchmarks. 7249 These expert implementations serve as ground truth 7250 references for quantitative evaluation.

7251 Papers in Category 1:

- 7252 • KeyNeRF [21]: Informative ray selection for few-shot 7253 neural radiance fields
- 7254 • mi-MLP NeRF [34]: Minimal MLP approach for few-shot 7255 view synthesis
- 7256 • LiNeRF [4]: Rethinking directional integration in neural 7257 radiance fields

- Anisotropic Neural Representation: Adaptive resolution for complex geometries 7258
- Efficient Ray Sampling [24]: Optimized sampling strategies for radiance field reconstruction 7259
- HybNeRF [28]: Hybrid multiresolution encoding for neural radiance fields 7260
- TVNeRF [33]: Total variation maximization for few-view 7261 neural volume rendering
- Surface Sampling [32]: Near-surface sampling with point 7262 cloud generation
- NeRF-ID [1]: Learning to sample for view synthesis 7263
- AR-NeRF [30]: Few-shot NeRF by Adaptive Rendering 7264 Loss Regularization 7265

5.1.2. Category 2: Non-Nerfstudio Papers (5 papers) 7266

These papers have existing public implementations but are not integrated into the Nerfstudio framework. This category enables direct comparison between our synthesized Nerfstudio plugins and original author implementations, evaluating both functional correctness and framework integration quality.

Papers in Category 2:

- Deblur-NeRF [15]: Neural radiance fields from blurry images 7267
- InfoNeRF [11]: Ray entropy minimization for few-shot 7268 neural volume rendering
- L0-Sampler [14]: Adaptive ray sampling for neural radiance 7269 fields
- NerfingMVS [29]: NerfingMVS: Guided Optimization of 7270 Neural Radiance Fields for Indoor Multi-view Stereo 7271

5.1.3. Category 3: Nerfstudio-Integrated Papers (5 pa- 7272 pers)

Papers already integrated into Nerfstudio serve as gold-standard references. These enable evaluation of synthesis quality and API compliance against production-quality implementations maintained by the research community.

Papers in Category 3:

- SeaThru-NeRF [13]: Neural radiance fields for underwater scenes 7273
- BioNeRF [22]: Bioneerf: Biologically Plausible Neural Radiance Fields for View Synthesis 7274
- Vanilla-NeRF [16]: Original neural radiance fields imple- 7275 mentation
- Instant-NGP [17]: Instant neural graphics primitives with 7276 multiresolution hash encoding
- Nerfacto [25]: Nerfstudio: A modular framework for neural 7277 radiance field development

5.1.4. Category 4: Novelty-Coverage Papers (10 papers) 7278

Selected for their distinct technical contributions, these papers evaluate how well synthesis systems capture and implement key research innovations. Each introduces novel loss

7308 functions, architectural components, or training strategies
 7309 that require deep understanding for correct implementation.

7310 Papers in Category 4:

- 7311 • Mip-NeRF [2]: Multiscale representation for anti-aliasing
 7312 neural radiance fields
- 7313 • BioNeRF [22]: Biologically plausible neural radiance
 7314 fields for view synthesis
- 7315 • PyNeRF [26]: Pyramidal neural radiance fields
- 7316 • TensoRF [3]: Tensorial radiance fields
- 7317 • Tetra-NeRF [12]: Representing neural radiance fields us-
 7318 ing tetrahedra
- 7319 • E-NeRF: Efficient neural radiance fields with learned em-
 7320 beddings
- 7321 • StyleNeRF [7]: Style-based 3D-aware generator for high-
 7322 resolution image synthesis
- 7323 • iNeRF [31]: Inverting neural radiance fields for pose esti-
 7324 mation
- 7325 • SigNeRF [5]: Scene integrated generation for neural radi-
 7326 ance fields
- 7327 • MCNeRF [8]: Monte Carlo rendering and denoising for
 7328 real-time NeRFs

7329 For each paper in NERIFY-BENCH we provide com-
 7330 prehensive resources to enable reproducible evaluation:

7331 **Frozen PDF Versions** The original paper PDF files are
 7332 included to ensure consistency across evaluations. This pre-
 7333 vents version drift and guarantees that all systems process
 7334 identical input documents.

7335 **Dual Markdown Representations** We provide two mark-
 7336 down versions for each paper:

- 7337 • *Raw extraction*: Direct conversion from PDF preserving
 7338 all text, equations, and structure
- 7339 • *Cleaned version*: Processed through our paper extraction
 7340 pipeline, with artifacts removed, mathematical notation
 7341 standardized, and formatting normalized for optimal pars-
 7342 ing

7343 **Tex source code** We provide LaTex source code of each
 7344 paper.

7345 **Ground Truth Repositories** Where available, we include
 7346 either the original author implementation or expert reimple-
 7347 mentations verified for correctness. These serve as reference
 7348 implementations for quantitative metrics including PSNR,
 7349 SSIM, and LPIPS on standard test scenes.

7350 5.1.5. Categorization Methodology

7351 5.2. Evaluation Protocol

7352 Our benchmark evaluates synthesized code across following
 7353 dimensions:

7354 Executability Metrics

- 7355 • *Build success rate*: Percentage of repositories that compile
 7356 without errors
- 7357 • *Import resolution*: Whether all dependencies resolve cor-
 7358 rectly

- *Training stability*: Convergence without NaN losses, gra-
 7359 dient explosions, or memory errors
- 7360

7361 **Rendering Quality Metrics** For successfully trained
 7362 models, we measure:

- *PSNR*: Peak signal-to-noise ratio on held-out test views
- 7363 • *SSIM*: Structural similarity index
- 7364 • *LPIPS*: Learned perceptual image patch similarity
- 7365 Results are averaged across multiple scenes from standard
 7366 benchmarks to ensure statistical significance.
- 7367

7368 **Novelty Fidelity** We evaluate coverage of paper-specific
 7369 innovations:

- *Correct (C)*: Percentage of novel components implemented
 7370 exactly as specified
- *Incorrect (I)*: Percentage with partial matches or flawed
 7372 logic
- *Missing (M)*: Percentage absent from generated code
- 7374 • *Weight match (W)*: Percentage of hyperparameters within
 7375 10% of paper specifications

7377 Runtime Efficiency

- *Training throughput*: Rays processed per second
- 7378 • *Memory usage*: Peak GPU memory consumption
- 7379 • *Convergence speed*: Iterations to reach target quality

7380 Code Quality

- *Nerfstudio conventions*: Adherence to framework patterns
 7382 and APIs
- *Documentation*: Presence of docstrings and inline com-
 7384 ments
- *Modularity*: Proper separation of concerns and reusable
 7386 components
- 7387

7388 5.2.1. Evaluation Dataset Details

7389 Standard datasets used for evaluation:

7390 Papers in Category 1:

- **KeyNeRF** [21]: Informative ray selection for few-shot
 7391 neural radiance fields
 - *Dataset*: NeRF-Synthetic
 - *Scene*: Average across all scenes
- **mi-MLP NeRF** [34]: Minimal MLP approach for few-
 7395 shot view synthesis
 - *Dataset*: NeRF-Synthetic
 - *Scene*: Average across all scenes
- **LiNeRF** [4]: Rethinking directional integration in neural
 7399 radiance fields
 - *Dataset*: NeRF-Synthetic
 - *Scene*: Drums
- **Anisotropic Neural Representation**: Adaptive resolution
 7403 for complex geometries
 - *Dataset*: NeRF-Synthetic
 - *Scene*: Average across all scenes
- **Efficient Ray Sampling** [24]: Optimized sampling strate-
 7407 gies for radiance field reconstruction
 - *Dataset*: DTU
 - *Scene*: Average on DTU test set

- 7411 • **HybNeRF** [28]: Hybrid multiresolution encoding for neural radiance fields
7412 – *Dataset*: NeRF-Synthetic
7413 – *Scene*: Average across all scenes
 - 7414 • **TVNeRF** [33]: Total variation maximization for few-view
7415 neural volume rendering
7416 – *Dataset*: NeRF-Synthetic
7417 – *Scene*: Hotdog
 - 7418 • **Surface Sampling** [32]: Near-surface sampling with point
7419 cloud generation
7420 – *Dataset*: NeRF-Synthetic
7421 – *Scene*: Ship (Table 2)
 - 7422 • **NeRF-ID** [1]: Learning to sample for view synthesis
7423 – *Dataset*: NeRF-Synthetic
7424 – *Scene*: Drums (Number reported in appendix)
 - 7425 • **AR-NeRF** [30]: Few-shot NeRF by Adaptive Rendering
7426 Loss Regularization
7427 – *Dataset*: DTU
7428 – *Scene*: 3-view setting
- Papers in Category 2:**
- 7431 • **Deblur-NeRF** [15]: Neural radiance fields from blurry
7432 images
7433 – *Dataset*: Deblur-NeRF Dataset
7434 – *Scene*: Cozyroom Camera Motion
 - 7435 • **InfoNeRF** [11]: Ray entropy minimization for few-shot
7436 neural volume rendering
7437 – *Dataset*: NeRF-Synthetic
7438 – *Scene*: Ship
 - 7439 • **L0-Sampler** [14]: Adaptive ray sampling for neural radi-
7440 ance fields
7441 – *Dataset*: NeRF-Synthetic
7442 – *Scene*: Ficus
 - 7443 • **NerfingMVS** [29]: Guided optimization of neural radi-
7444 ance fields for indoor multi-view stereo
7445 – *Dataset*: ScanNet
7446 – *Scene*: Scene 0653 (processed data)
- Papers in Category 3:**
- 7447 • **SeaThru-NeRF** [13]: Neural radiance fields for underwa-
7448 ter scenes
7449 – *Dataset*: SeaThru-NeRF Dataset
7450 – *Scene*: Panama
 - 7451 • **BioNeRF** [22]: Biologically plausible neural radiance
7452 fields for view synthesis
7453 – *Dataset*: NeRF-Synthetic
7454 – *Scene*: Drums
 - 7455 • **Vanilla-NeRF** [16]: Original neural radiance fields imple-
7456 mentation
7457 – *Dataset*: NeRF-Synthetic
7458 – *Scene*: Chair
 - 7459 • **Instant-NGP** [17]: Instant neural graphics primitives with
7460 multiresolution hash encoding
7461 – *Dataset*: NeRF-Synthetic
7462 – *Scene*: Lego

- 7463 • **Nerfacto** [25]: A modular framework for neural radiance
7464 field development
7465 – *Dataset*: Nerfstudio Dataset
7466 – *Scene*: Poster

6. Note on Figure Rendering

A rendering artifact in Figure 2 of the main paper omits two arrows in Stage 4: from rendered views to Critique Agent, and from Critique Agent to JSON box. These connections complete the visual feedback loop central to our refinement process.

References

- [1] Relja Arandjelović and Andrew Zisserman. Nerf in detail: Learning to sample for view synthesis. *arXiv preprint arXiv:2106.05264*, 2021. 1, 77, 79
- [2] Jonathan T Barron, Ben Mildenhall, Matthew Tancik, Peter Hedman, Ricardo Martin-Brualla, and Pratul P Srinivasan. Mip-nerf: A multiscale representation for anti-aliasing neural radiance fields. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 5855–5864, 2021. 18, 78
- [3] Anpei Chen, Zexiang Xu, Andreas Geiger, Jingyi Yu, and Hao Su. Tensorrf: Tensorial radiance fields. In *European conference on computer vision*, pages 333–350. Springer, 2022. 27, 78
- [4] Congyue Deng, Jiawei Yang, Leonidas Guibas, and Yue Wang. Rethinking directional integration in neural radiance fields. *arXiv preprint arXiv:2311.16504*, 2023. 1, 77, 78
- [5] Jan-Niklas Dihlmann, Andreas Engelhardt, and Hendrik Lensch. Signerf: Scene integrated generation for neural radiance fields. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 6679–6688, 2024. 78
- [6] Stephan J Garbin, Marek Kowalski, Matthew Johnson, Jamie Shotton, and Julien Valentin. Fastnerf: High-fidelity neural rendering at 200fps. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 14346–14355, 2021. 1
- [7] Jiatao Gu, Lingjie Liu, Peng Wang, and Christian Theobalt. Stylenarf: A style-based 3d-aware generator for high-resolution image synthesis. *arXiv preprint arXiv:2110.08985*, 2021. 43, 78
- [8] Kunal Gupta, Milos Hasan, Zexiang Xu, Fujun Luan, Kalyan Sunkavalli, Xin Sun, Manmohan Chandraker, and Sai Bi. Mcnerf: Monte carlo rendering and denoising for real-time nerfs. In *SIGGRAPH Asia 2023 Conference Papers*, pages 1–11, 2023. 59, 78
- [9] HKUDS. Deepcode: Open agentic coding. <https://github.com/HKUDS/DeepCode>, 2025. GitHub repository. 1
- [10] Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu Zheng, Yuheng Cheng, Jinlin Wang, Ceyao Zhang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, et al. Metagpt: Meta programming for a multi-agent collaborative framework. In *The Twelfth*

- 7517 International Conference on Learning Representations, 2023.
7518 1
- 7519 [11] Mijeong Kim, Seonguk Seo, and Bohyung Han. Infonerf: Ray
7520 entropy minimization for few-shot neural volume rendering.
7521 In *Proceedings of the IEEE/CVF Conference on Computer*
7522 *Vision and Pattern Recognition*, pages 12912–12921, 2022.
7523 1, 77, 79
- 7524 [12] Jonas Kulhanek and Torsten Sattler. Tetra-nerf: Representing
7525 neural radiance fields using tetrahedra. In *Proceedings of the*
7526 *IEEE/CVF international conference on computer vision*,
7527 pages 18458–18469, 2023. 32, 78
- 7528 [13] Deborah Levy, Amit Peleg, Naama Pearl, Dan Rosenbaum,
7529 Derya Akkaynak, Simon Korman, and Tali Treibitz. Seathru-
7530 nerf: Neural radiance fields in scattering media. In *Proceed-*
7531 *ings of the IEEE/CVF conference on computer vision and*
7532 *pattern recognition*, pages 56–65, 2023. 77, 79
- 7533 [14] Liangchen Li and Juyong Zhang. L0-sampler: An l0 model
7534 guided volume sampling for nerf. In *Proceedings of the*
7535 *IEEE/CVF Conference on Computer Vision and Pattern*
7536 *Recognition*, pages 21390–21400, 2024. 1, 77, 79
- 7537 [15] Li Ma, Xiaoyu Li, Jing Liao, Qi Zhang, Xuan Wang, Jue
7538 Wang, and Pedro V Sander. Deblur-nerf: Neural radiance
7539 fields from blurry images. In *Proceedings of the IEEE/CVF*
7540 *conference on computer vision and pattern recognition*, pages
7541 12861–12870, 2022. 1, 77, 79
- 7542 [16] Ben Mildenhall, Pratul P. Srinivasan, Matthew Tancik,
7543 Jonathan T. Barron, Ravi Ramamoorthi, and Ren Ng. Nerf:
7544 Representing scenes as neural radiance fields for view syn-
7545 thesis. In *European Conference on Computer Vision (ECCV)*,
7546 pages 405–421, 2020. 1, 77, 79
- 7547 [17] Thomas Müller, Alex Evans, Christoph Schied, and Alexan-
7548 der Keller. Instant neural graphics primitives with a multires-
7549 olution hash encoding. *ACM Transactions on Graphics*, 41
7550 (4):1–15, 2022. 77, 79
- 7551 [18] Thomas Neff, Pascal Stadlbauer, Mathias Parger, Andreas
7552 Kurz, Joerg H Mueller, Chakravarty R Alla Chaitanya, Anton
7553 Kaplanyan, and Markus Steinberger. Donerf: Towards real-
7554 time rendering of compact neural radiance fields using depth
7555 oracle networks. In *Computer Graphics Forum*, pages 45–59.
7556 Wiley Online Library, 2021. 1
- 7557 [19] Junbo Niu, Zheng Liu, Zhuangcheng Gu, Bin Wang, Linke
7558 Ouyang, Zhiyuan Zhao, Tao Chu, Tianyao He, Fan Wu, Qin-
7559 tong Zhang, et al. Mineru2.5: A decoupled vision-language
7560 model for efficient high-resolution document parsing. *arXiv*
7561 preprint arXiv:2509.22186, 2025. 68
- 7562 [20] OpenAI. Gpt-5 technical report. <https://openai.com>,
7563 2025. 1
- 7564 [21] Marco Orsingher, Anthony Dell’Eva, Paolo Zani, Paolo
7565 Medici, and Massimo Bertozzi. Informative rays selec-
7566 tion for few-shot neural radiance fields. *arXiv preprint*
7567 arXiv:2312.17561, 2023. 1, 77, 78
- 7568 [22] Leandro Aparecido Passos, Danilo Samuel Jodas, Ahsan
7569 Adeel, João P Papa, Douglas Rodrigues, and Kelton Augusto
7570 Pontara da Costa. Bionerf: Biologically plausible neural radi-
7571 ance fields for view synthesis. Available at SSRN 5384186,
7572 2024. 77, 78, 79
- 7573 [23] Chen Qian, Xin Cai, Cheng Liu, Yang Liu, Juyuan Dang, Lin
7574 Wang, et al. Chatdev: Communicative agents for software
7575 development. *arXiv preprint arXiv:2307.07924*, 2023. 1
- 7576 [24] Shilei Sun, Ming Liu, Zhongyi Fan, Qingliang Jiao, Yuxue
7577 Liu, Liquan Dong, and Lingqin Kong. Efficient ray sampling
7578 for radiance fields reconstruction. *Computers & Graphics*,
7579 118:48–59, 2024. 1, 77, 78
- 7580 [25] Matthew Tancik, Ethan Weber, Eevonne Ng, Ruilong Li, Brent
7581 Yi, Terrance Wang, Alexander Kristoffersen, Jake Austin,
7582 Kamyr Salahi, Abhik Ahuja, et al. Nerfstudio: A modular
7583 framework for neural radiance field development. In *ACM*
7584 *SIGGRAPH Conference Proceedings*, 2023. 1, 77, 79
- 7585 [26] Haithem Turki, Michael Zollhöfer, Christian Richardt, and
7586 Deva Ramanan. Pynerf: Pyramidal neural radiance fields.
7587 *Advances in neural information processing systems*, 36:37670–
7588 37681, 2023. 78
- 7589 [27] Yifan Wang, Jun Xu, Y Zeng, and Y Gong. Anisotropic neu-
7590 ral representation learning for high-quality neural rendering.
7591 *arXiv preprint arXiv:2311.18311*, 2023. 1
- 7592 [28] Yifan Wang, Yi Gong, and Yuan Zeng. Hyb-nerf: A multires-
7593 olution hybrid encoding for neural radiance fields. In *2024*
7594 *IEEE/CVF Winter Conference on Applications of Computer*
7595 *Vision (WACV)*, pages 3677–3686. IEEE, 2024. 1, 77, 79
- 7596 [29] Yi Wei, Shaohui Liu, Yongming Rao, Wang Zhao, Jiwen Lu,
7597 and Jie Zhou. Nerfingmvs: Guided optimization of neural
7598 radiance fields for indoor multi-view stereo, 2021. 77, 79
- 7599 [30] Qingshan Xu, Xuanyu Yi, Jianyao Xu, Wenbing Tao, Yew-
7600 Soon Ong, and Hanwang Zhang. Few-shot nerf by adaptive
7601 rendering loss regularization. In *European Conference on*
7602 *Computer Vision*, pages 125–142. Springer, 2024. 77, 79
- 7603 [31] Lin Yen-Chen, Pete Florence, Jonathan T Barron, Alberto
7604 Rodriguez, Phillip Isola, and Tsung-Yi Lin. inferf: Inverting
7605 neural radiance fields for pose estimation. In *2021 IEEE/RSJ*
7606 *International Conference on Intelligent Robots and Systems*
7607 (*IROS*), pages 1323–1330. IEEE, 2021. 48, 78
- 7608 [32] Hye Bin Yoo, Hyun Min Han, Sung Soo Hwang, and Il Yong
7609 Chun. Improving neural radiance fields using near-surface
7610 sampling with point cloud generation. *Neural Processing*
7611 *Leters*, 56(4):214, 2024. 1, 77, 79
- 7612 [33] Yao Zhang, Jiangshu Wei, Bei Zhou, Fang Li, Yuxin Xie,
7613 and Jiajun Liu. Tvnarf: Improving few-view neural volume
7614 rendering with total variation maximization. *Knowledge-
7615 Based Systems*, 301:112273, 2024. 1, 77, 79
- 7616 [34] Hanxin Zhu, Tianyu He, Xin Li, Bingchen Li, and Zhibo
7617 Chen. Is vanilla mlp in neural radiance field enough for
7618 few-shot view synthesis? In *Proceedings of the IEEE/CVF*
7619 *Conference on Computer Vision and Pattern Recognition*,
7620 pages 20288–20298, 2024. 1, 77, 78